

Splunk 101/201 Hands-on Workshops



splunk>

Forward- looking statements

This presentation may contain forward-looking statements regarding future events, plans or the expected financial performance of our company, including our expectations regarding our products, technology, strategy, customers, markets, acquisitions and investments. These statements reflect management's current expectations, estimates and assumptions based on the information currently available to us. These forward-looking statements are not guarantees of future performance and involve significant risks, uncertainties and other factors that may cause our actual results, performance or achievements to be materially different from results, performance or achievements expressed or implied by the forward-looking statements contained in this presentation.

For additional information about factors that could cause actual results to differ materially from those described in the forward-looking statements made in this presentation, please refer to our periodic reports and other filings with the SEC, including the risk factors identified in our most recent quarterly reports on Form 10-Q and annual reports on Form 10-K, copies of which may be obtained by visiting the Splunk Investor Relations website at www.investors.splunk.com or the SEC's website at www.sec.gov. The forward-looking statements made in this presentation are made as of the time and date of this presentation. If reviewed after the initial presentation, even if made available by us, on our website or otherwise, it may not contain current or accurate information. We disclaim any obligation to update or revise any forward-looking statement based on new information, future events or otherwise, except as required by applicable law.

In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. We undertake no obligation either to develop the features or functionalities described, in beta or in preview (used interchangeably), or to include any such feature or functionality in a future release.

Splunk, Splunk> and Turn Data Into Doing are trademarks and registered trademarks of Splunk Inc. in the United States and other countries. All other brand names, product names or trademarks belong to their respective owners.
© 2024 Splunk Inc. All rights reserved.



Workshop Logistics

See Zoom chat for Splunk workshop instance details

Workshop flow:

- **LEARN:** New capability and outcomes on each slide
- **PRACTICE:** Apply techniques in a demo environment after each slide
- **WORK:** Apply techniques in your environment with your Splunk team after the session

Ask questions! What outcomes are you trying to achieve?

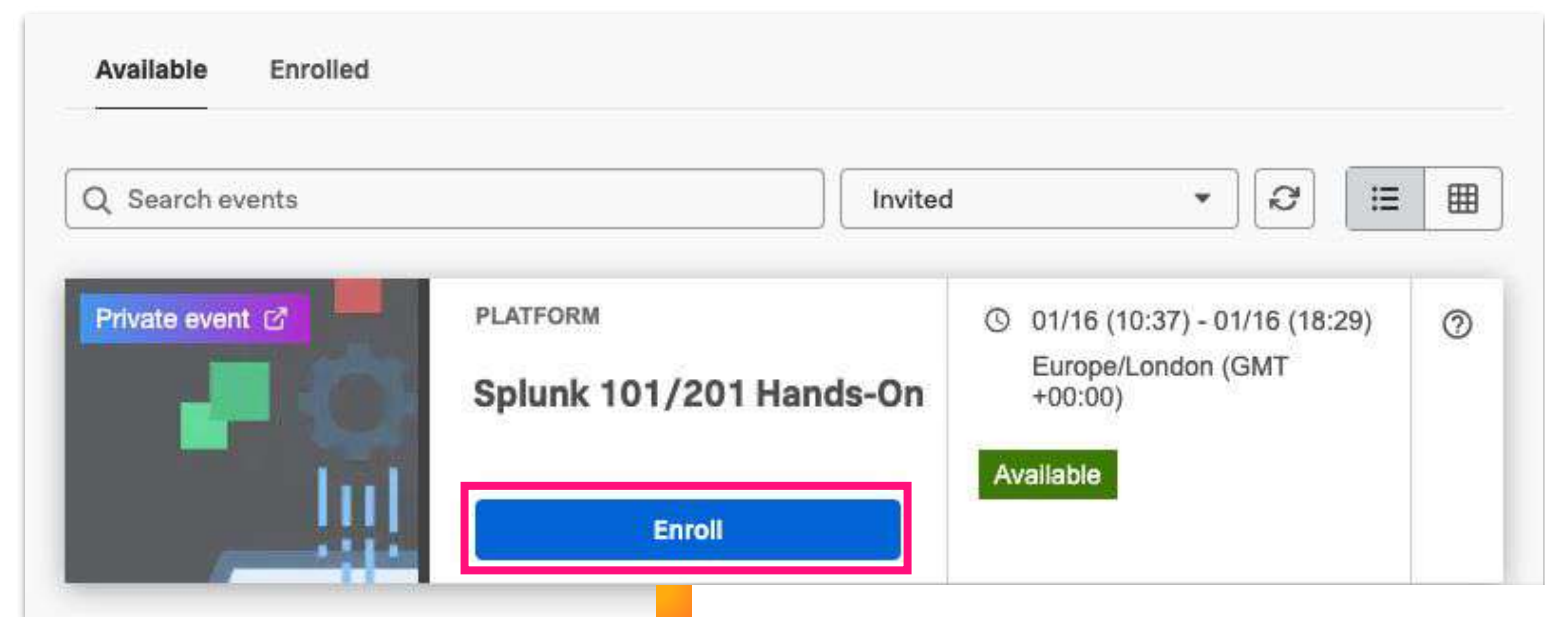


Enroll in Today's Workshop

Tasks

1. Get a splunk.com account if you don't have one yet:
<https://splk.it/SignUp>
2. Enroll in the Splunk Show workshop event:
<https://show.splunk.com/event/<eventID>>
3. Download a copy of today's slide deck:
<https://splk.it/101-201-Attendee>

Goal



Enroll in today's event

Introduction



Splunk 101 Session Overview

Session Goals:

- Enable your teams to search, investigate, analyze & report on data
- Drive more value and better outcomes from machine data
- Increase teams' productivity and efficiency in resolving issues

Splunk 101 Agenda (for users):

- Search UI
- Anatomy of a Search
- Search Best Practices
- Open chevron, event details, actions, etc.
- Customized time spans
- search commands, Booleans, save as dashboard
- stats, timechart, eval & where functions
- rename, dedup, fields/table, head/rare/top functions
- Common use cases in Security, IT, and Observability

Splunk 201 Session Overview

Session Goals:

- Enable your teams to search, investigate, analyze & report on data
- Drive more value and better outcomes from machine data
- Increase teams' productivity and efficiency in resolving issues

Splunk 201 Agenda (for power users):

- Transactional analysis (stats list/values, transaction, histograms)
- Advanced statistics (stats stdev/variance, eventstats & streamstats)
- Advanced Dashboarding (form inputs, tokens, post-process search, Dashboard Studio)
- Schema-on-read and Field Extraction (rex, Interactive Field Extractor, eventtypes/tags/macros, Common Information Model)

Splunk 301 for Security and IT/Observability

We have lots of 301-level workshops available!

Security 301 Workshops include:

- Security Operations Suite Hands-On Workshop
- Enterprise Security Hands-On Workshop

IT/Observability 301 Workshops include:

- Business Service Insights Workshop
- ITSI Event Analytics Workshop
- Splunk4Ninjas - IT Service Intelligence (ITSI)

For more options, reach out to your Splunk account team and we can facilitate based on your use cases

Additional Resources

Splunk Docs (with detailed reference for all commands):

<https://docs.splunk.com/>

FREE Splunk Fundamentalstraining:

https://www.splunk.com/en_us/training/free-courses/overview.html

Splunk Lantern (best practices & use cases):

<https://lantern.splunk.com>

Splunk Community (answers to common & uncommon questions):

<https://community.splunk.com/>

.conf Online (incl. past keynotes & customer presentations):

<https://conf.splunk.com/>

Splunk Quick Reference Guide:

<https://www.splunk.com/pdfs/solution-guides/splunk-quick-reference-guide.pdf>



Splunk Use Cases & Customer Examples

IT Operations



Predict service-level degradation before it occurs

Application Observability



Improve app performance and reliability using app logs and infrastructure

Security & Compliance



Speed up security investigations and reduce the impact of insider threats

Business Analytics



Drive more orders with marketing campaigns across website and mobile

Internet of Things



Monitor and resolve problems from 10,000s of sensors in real time

Our Investigative Approach


Adaptable | Real-Time | Fast To Value | Massive Scale

Send

unstructured data from all systems, devices, and people

React

quickly to changing circumstances by asking questions immediately



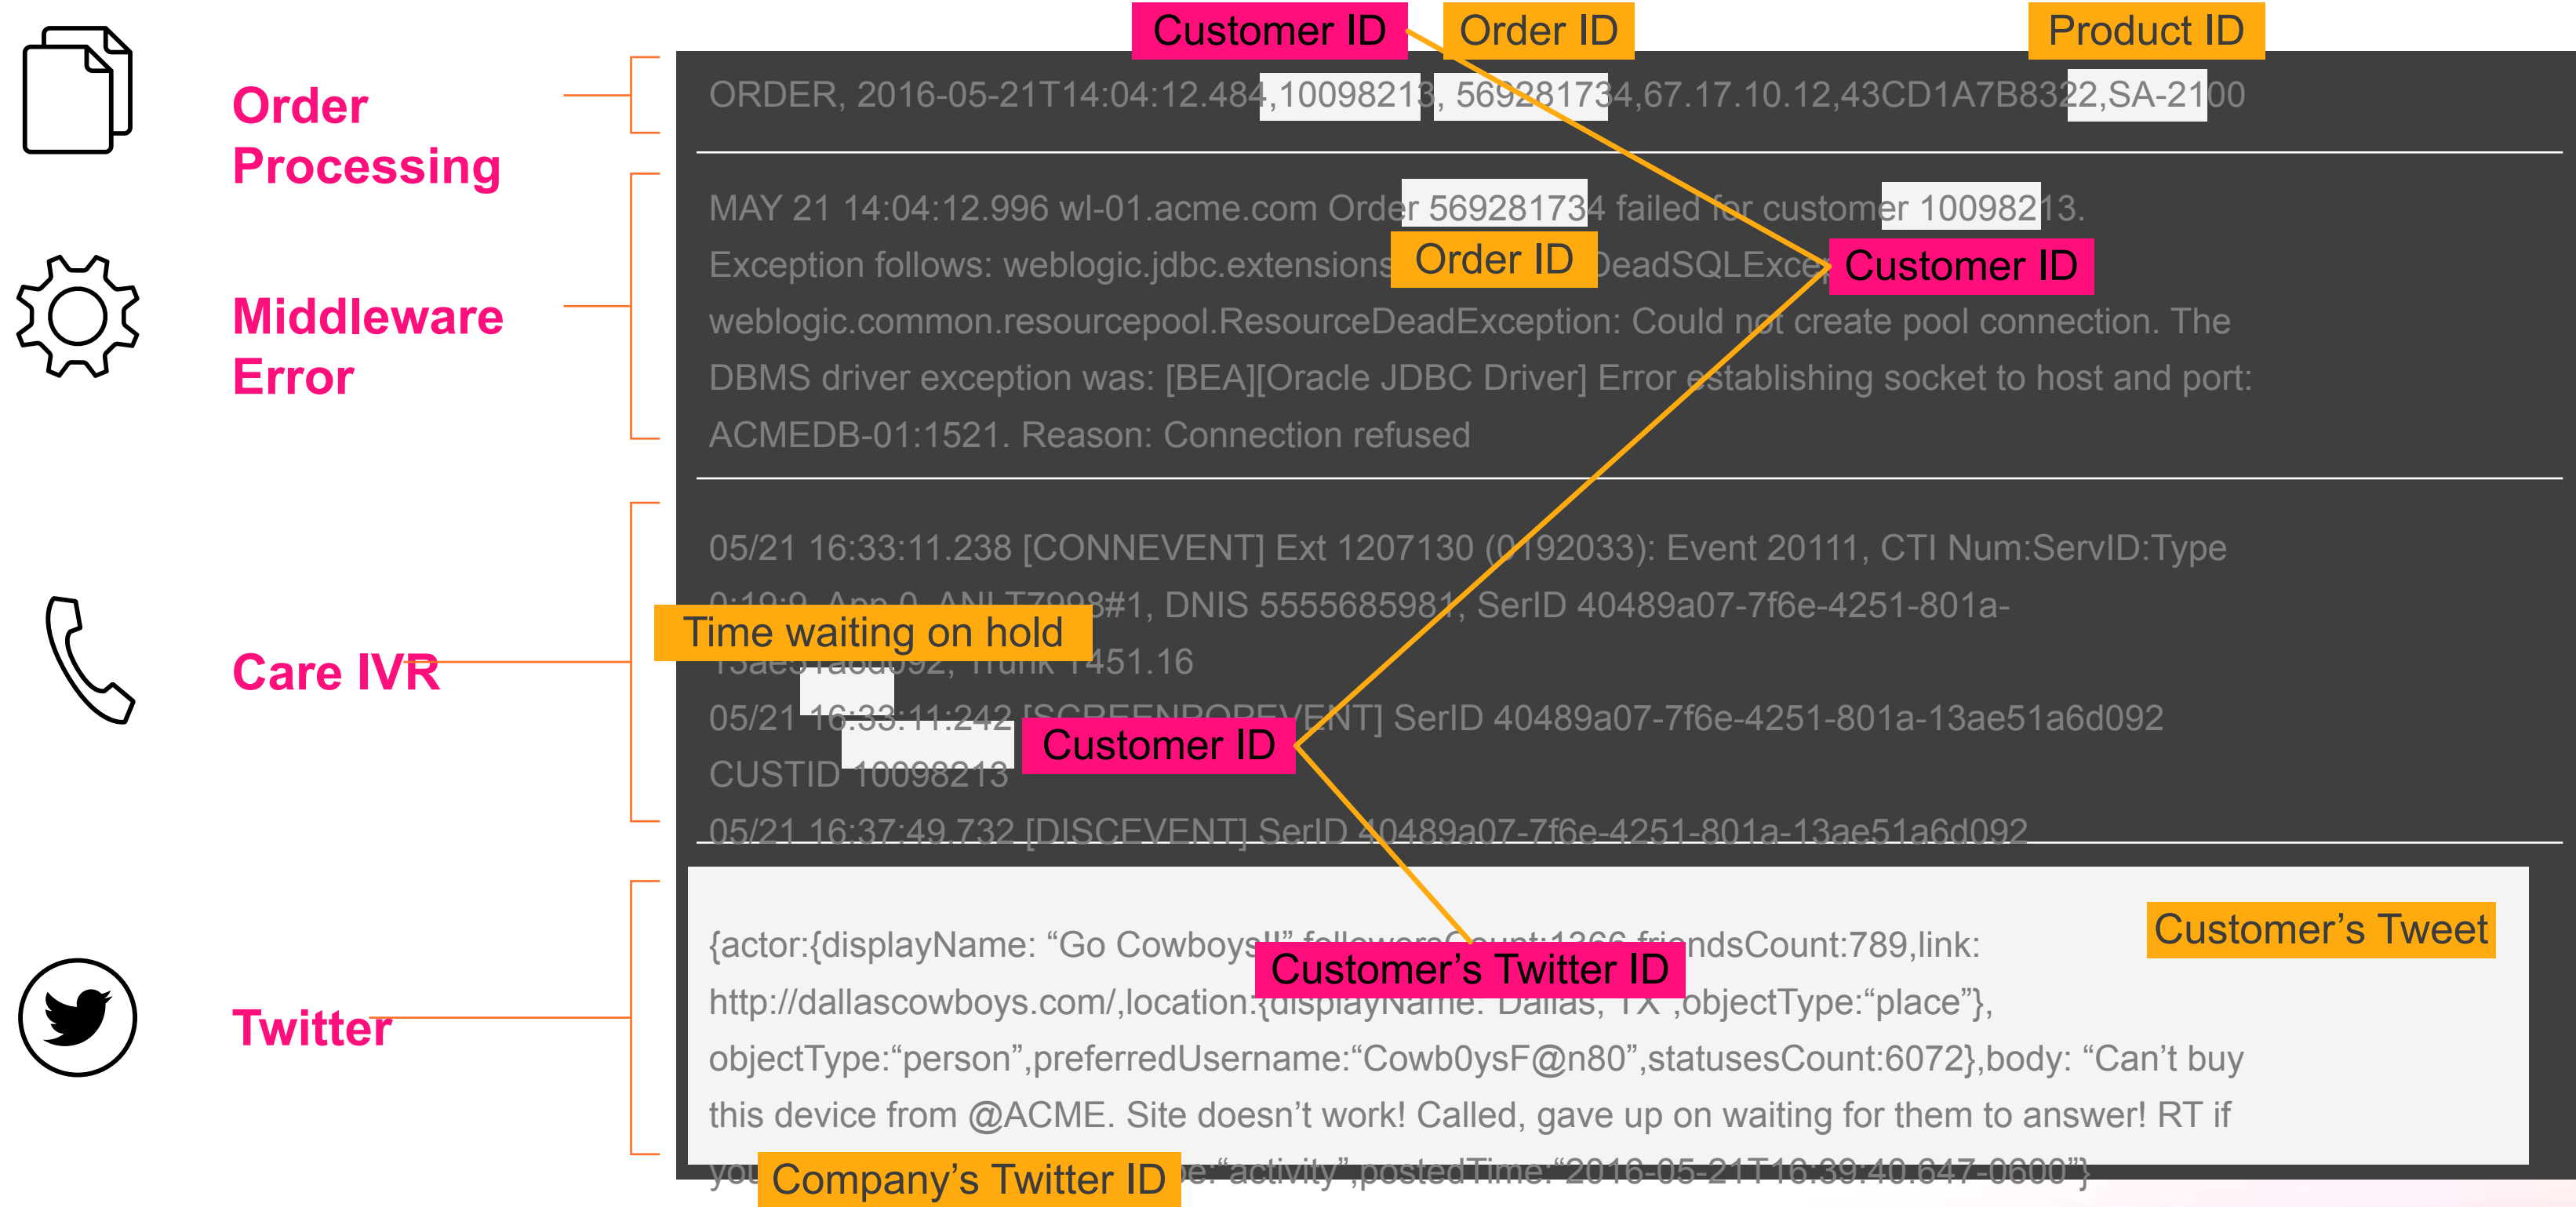
Don't Structure
your data until you are ready

Splunk 101



Reduce impact by observing end-to-end transactions

SOURCES



Search UI

Search & investigate logs for errors:

- `sourcetype=* error`

Quick reports:

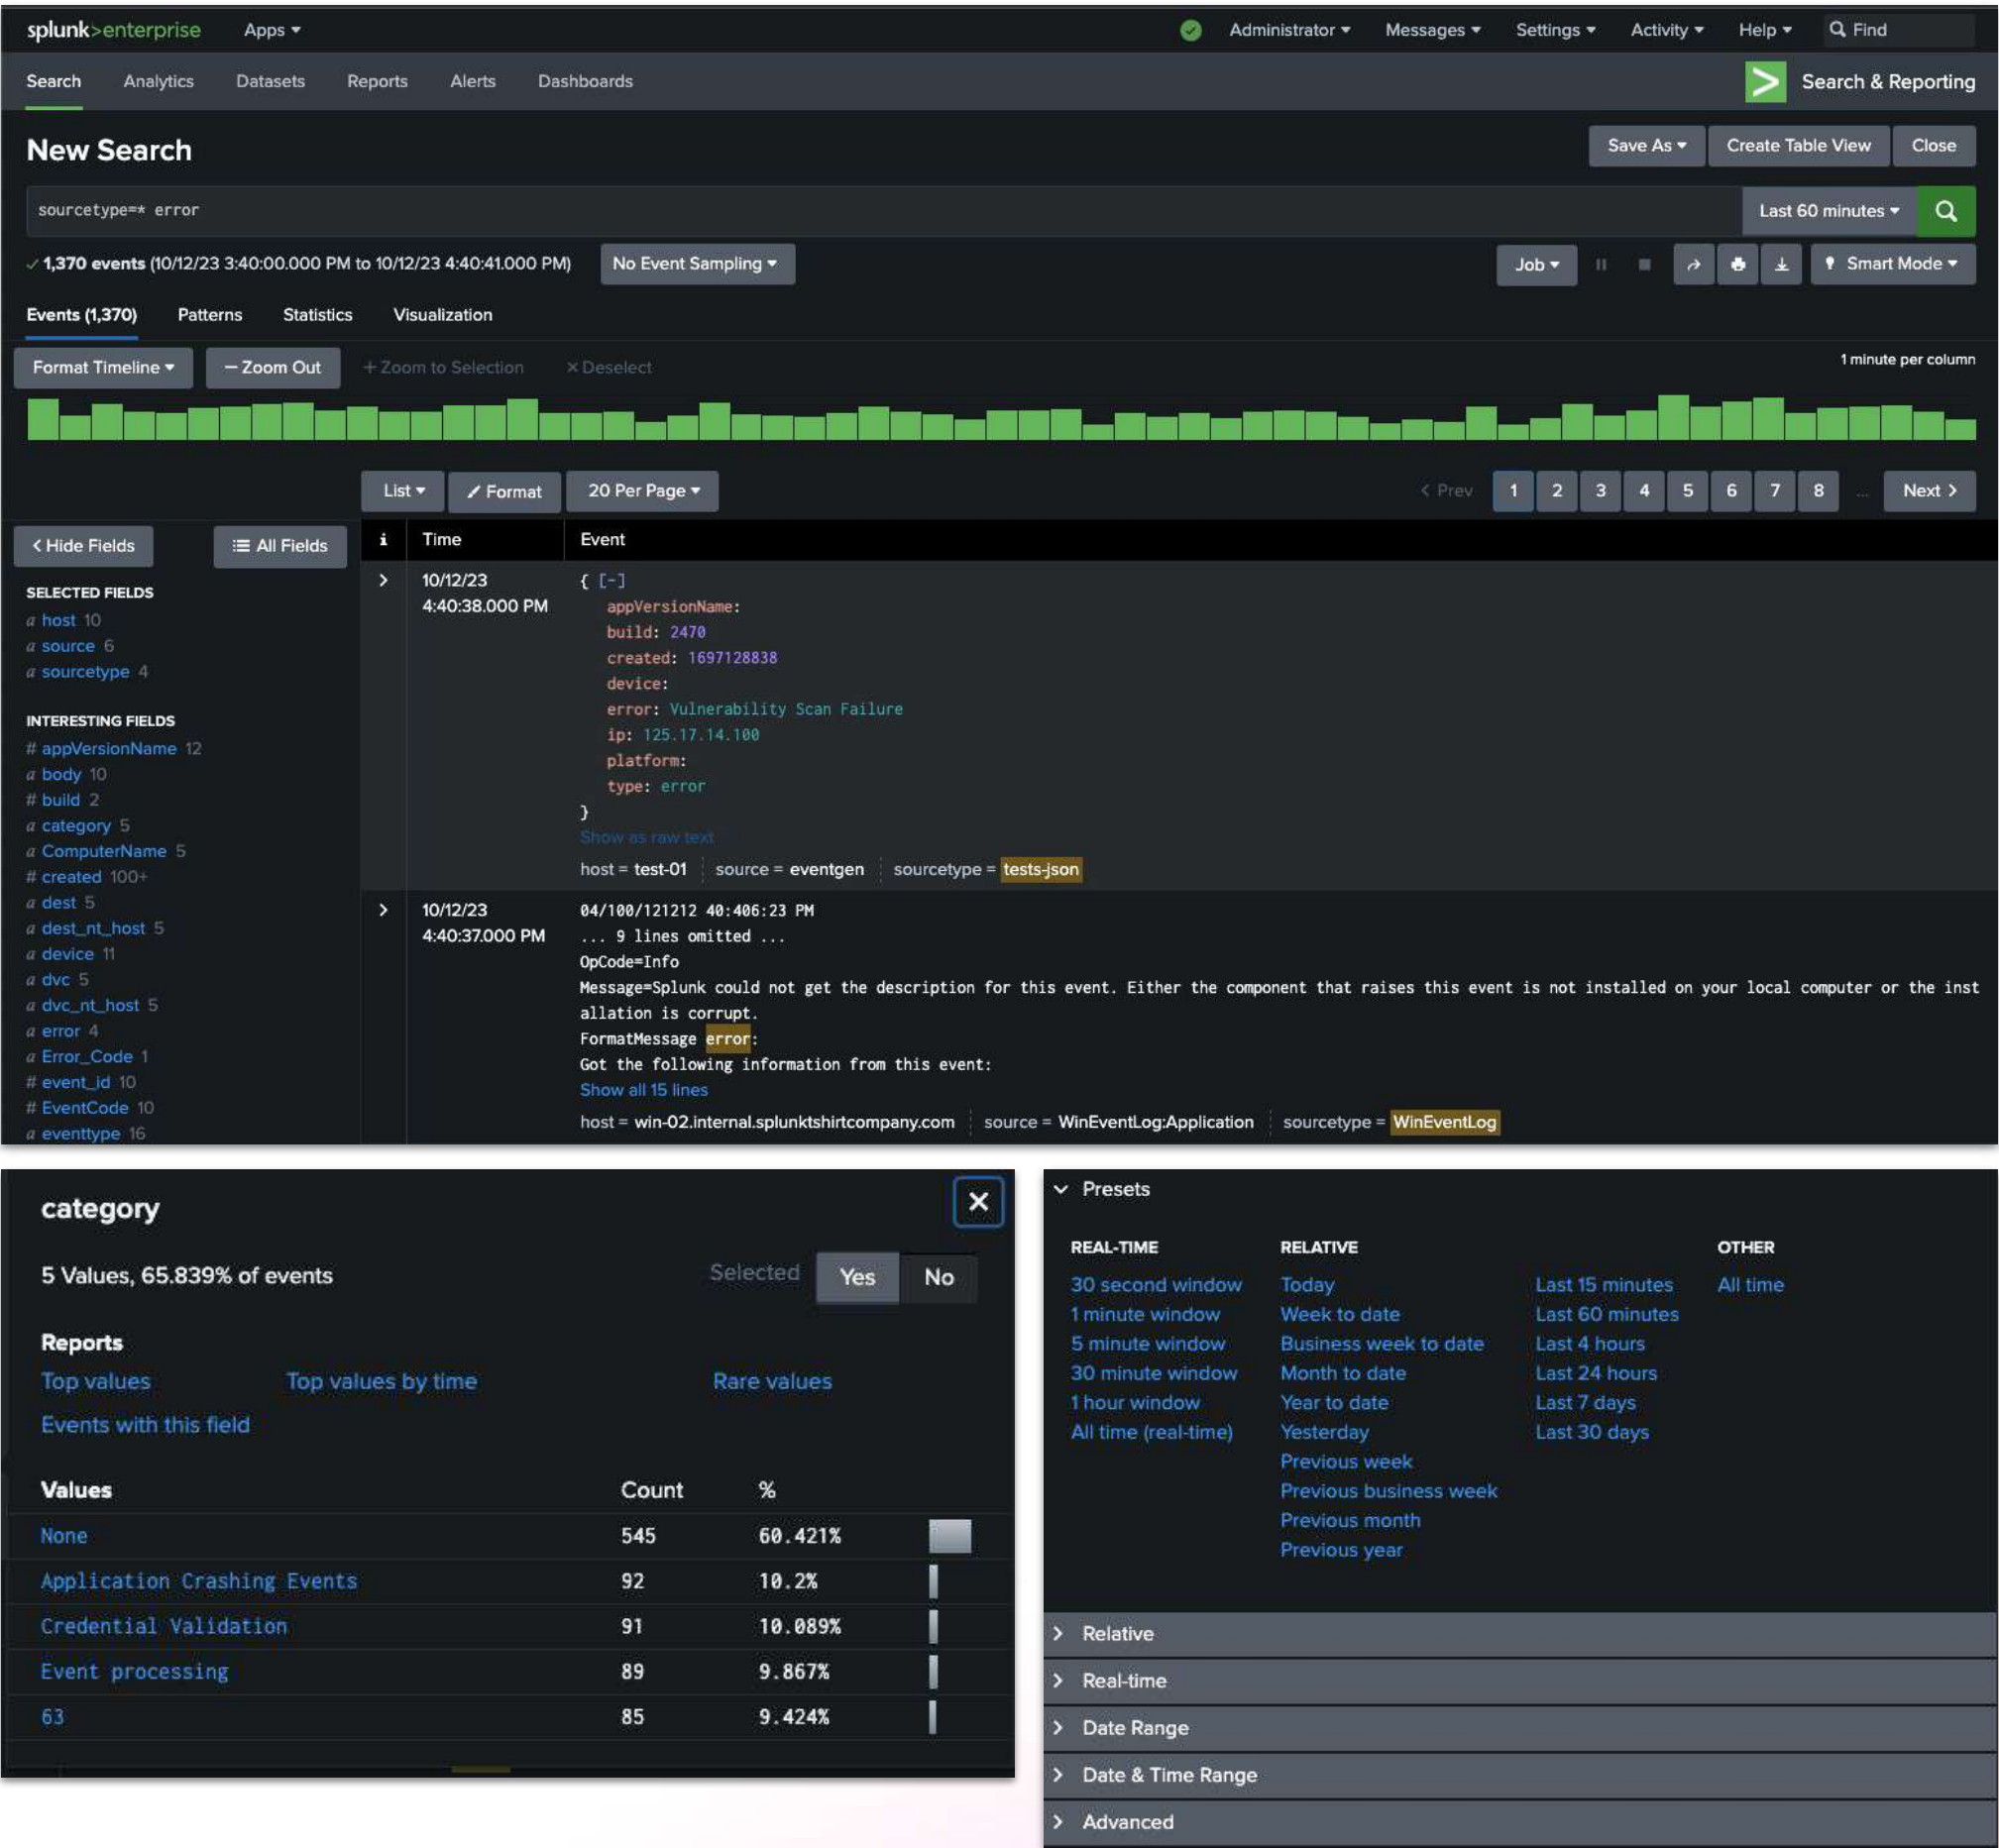
- Top values (by time), rare values
- Events with this field
- Average/Maximum/Minimum over time

Customize time spans

- Relative, Real-time, Date & Time Range

Outcomes & value:

- Find needle-in-haystack problems or threats
- Reduce Mean Time To Resolve (MTTR) issues



Search UI (cont.)

Open chevron:

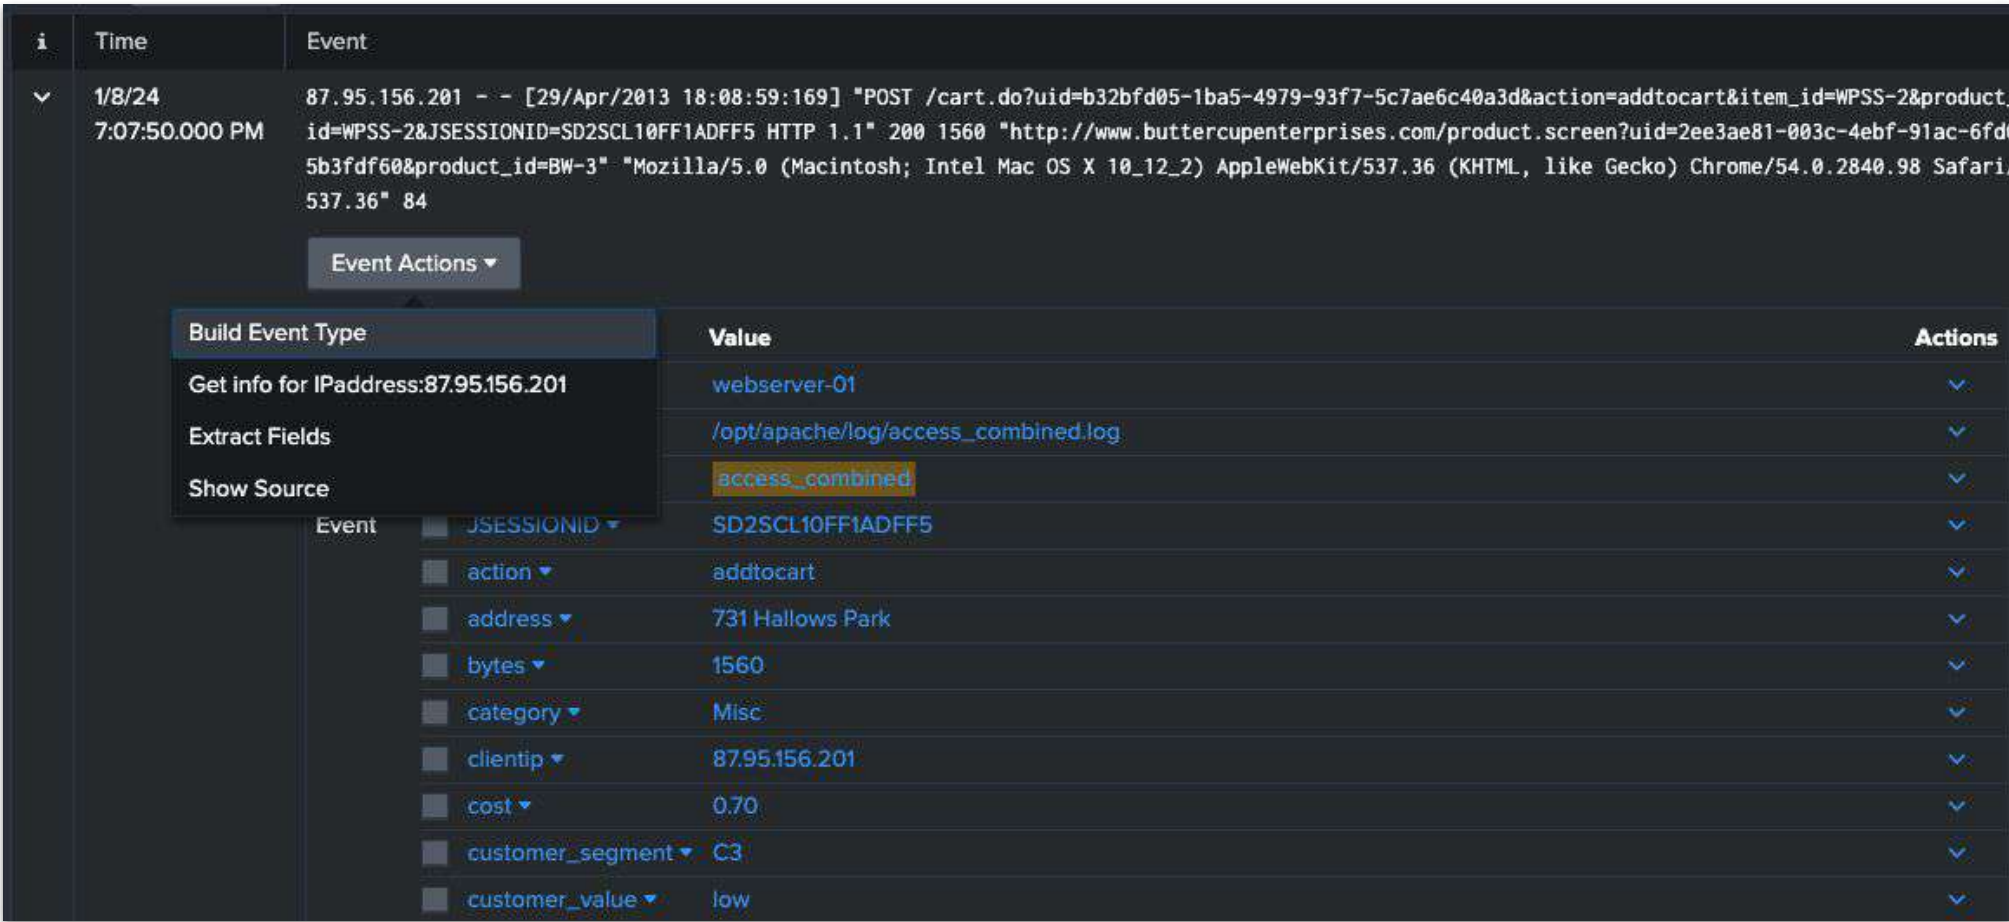
- Get event details
- Add new tags
- Add/exclude or new search from terms
- Trigger event actions & automation
(Work with power users & Splunk admin team)

Search nearby events:

- Click timestamp
- Events before or after
- Modify search within new time range

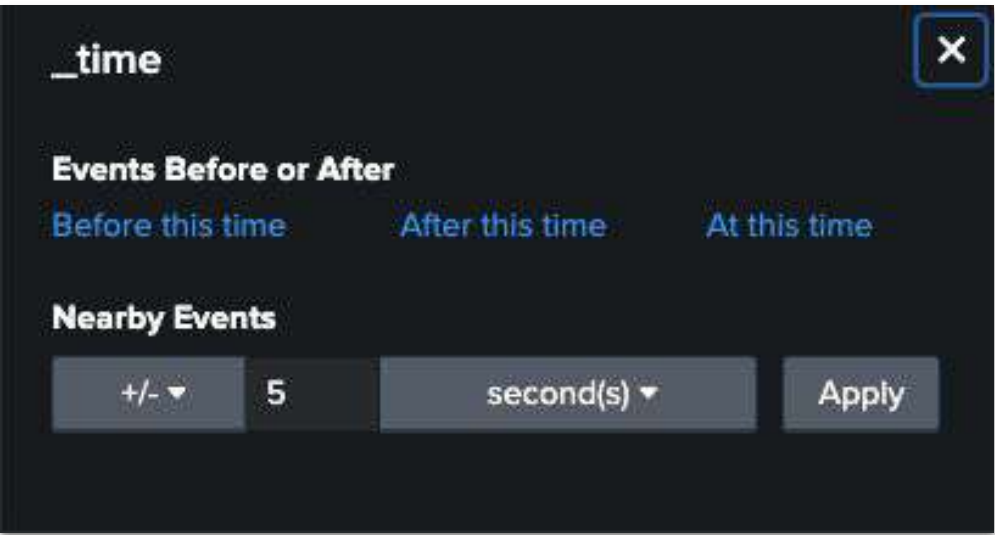
Outcomes & value:

- Reduce MTTR by surfacing related events via metadata, keywords, and time

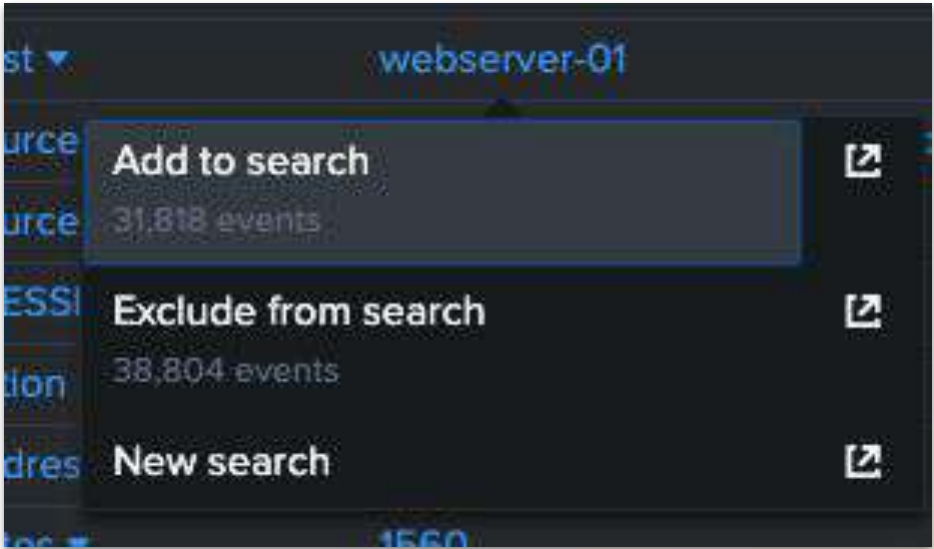


The screenshot shows a Splunk search results table with columns 'i', 'Time', and 'Event'. A specific event is expanded, showing a dropdown menu for 'Event Actions'. The menu options are 'Build Event Type', 'Get info for IPAddress:87.95.156.201', 'Extract Fields', and 'Show Source'. Below the menu, a table lists event fields and their values:

Field	Value	Actions
webserver-01		⌵
/opt/apache/log/access_combined.log		⌵
access_combined		⌵
JSESSIONID	SD2SCL10FF1ADFF5	⌵
action	addtocart	⌵
address	731 Hallows Park	⌵
bytes	1560	⌵
category	Misc	⌵
clientip	87.95.156.201	⌵
cost	0.70	⌵
customer_segment	C3	⌵
customer_value	low	⌵



The screenshot shows the 'Nearby Events' panel. It has a title bar with a close button. Below the title, there are three tabs: 'Events Before or After', 'Before this time', 'After this time', and 'At this time'. The 'Events Before or After' tab is selected. Below the tabs, there is a section for 'Nearby Events' with a range selector set to '+/- 5 second(s)' and an 'Apply' button.



The screenshot shows a panel with three options: 'Add to search', 'Exclude from search', and 'New search'. Each option has a corresponding icon to its right.

<https://docs.splunk.com/Documentation/Splunk/latest/Search/Usertimeaccelerators>

Search with Booleans

Combine events with Booleans (**OR**, **AND**, **NOT**)

Search combined app log + web log data:

- `sourcetype=tests-json OR sourcetype=access_combined`

Search app log data for host X **AND** platform Y:

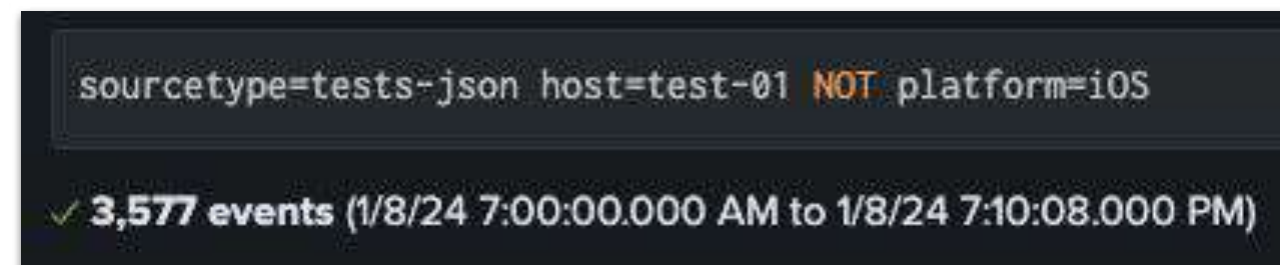

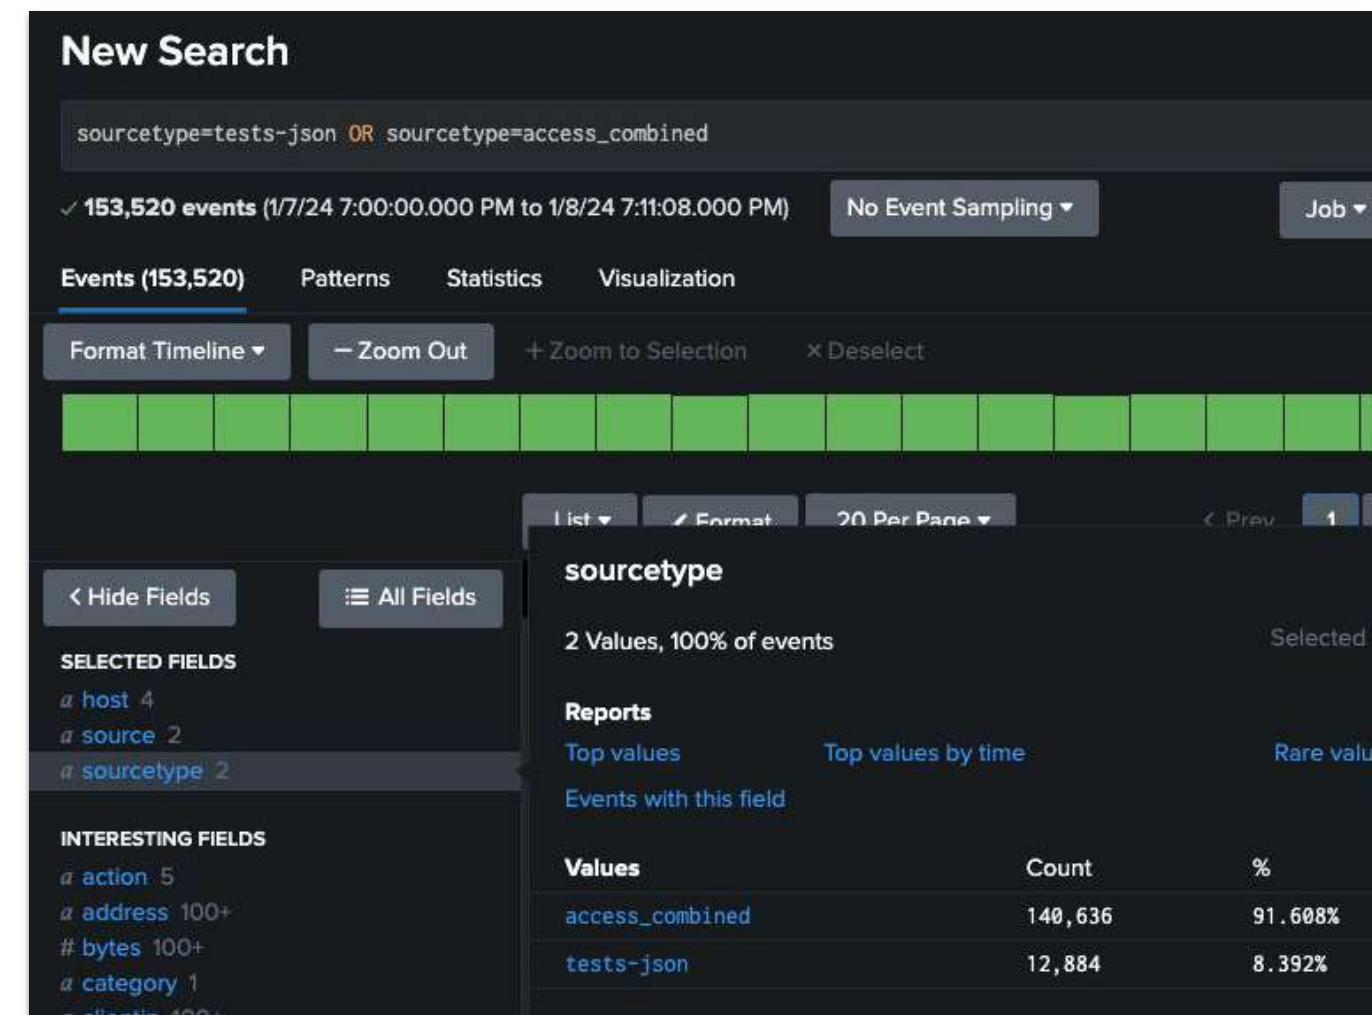
- `sourcetype=tests-json host=test-01 platform=iOS`

Search app log data for host X **AND NOT** platform Y:

- `sourcetype=tests-json host=test-01 NOT platform=iOS`

Outcomes & value:

- Detect and resolve issues by narrowing in on exact events (needles in haystacks)




Anatomy of a Search

Search Query Structure

```
index=waldo error OR fail* | eval loc=long+lat+alt | geoip loc
```

retrieve events



filter/transform/operate/map



Search Best Practices

Search for just what you need:

- Small time ranges (e.g., 15 minutes, not all time)
- Exact indexes and sourcetypes:
 - `index=oidemo (sourcetype=access_combined OR sourcetype=tests-json)`
- Use eventtypes and tags for consistent and repeatable searches:
 - `tag=network`
 - `eventtype=nix-all-logs`

Reduce load on system with optimal exact techniques instead of easy ones:

- More efficient and optimal with “stats list” or “stats value” commands:
 - `index=oidemo sourcetype=access_combined | stats list(action) by JSESSIONID`
- Less efficient but easy with “transaction” command:
 - `index=oidemo sourcetype=access_combined | transaction JSESSIONID`

Outcomes & value:

- Reduce MTTR by pulling back exact data and running faster searches
- Reduce impact to other users and deliver more efficient Splunk service

Search & stats

Use `stats` command to turn data into aggregate tables

Example: Track errors vs successes and load times by platform:

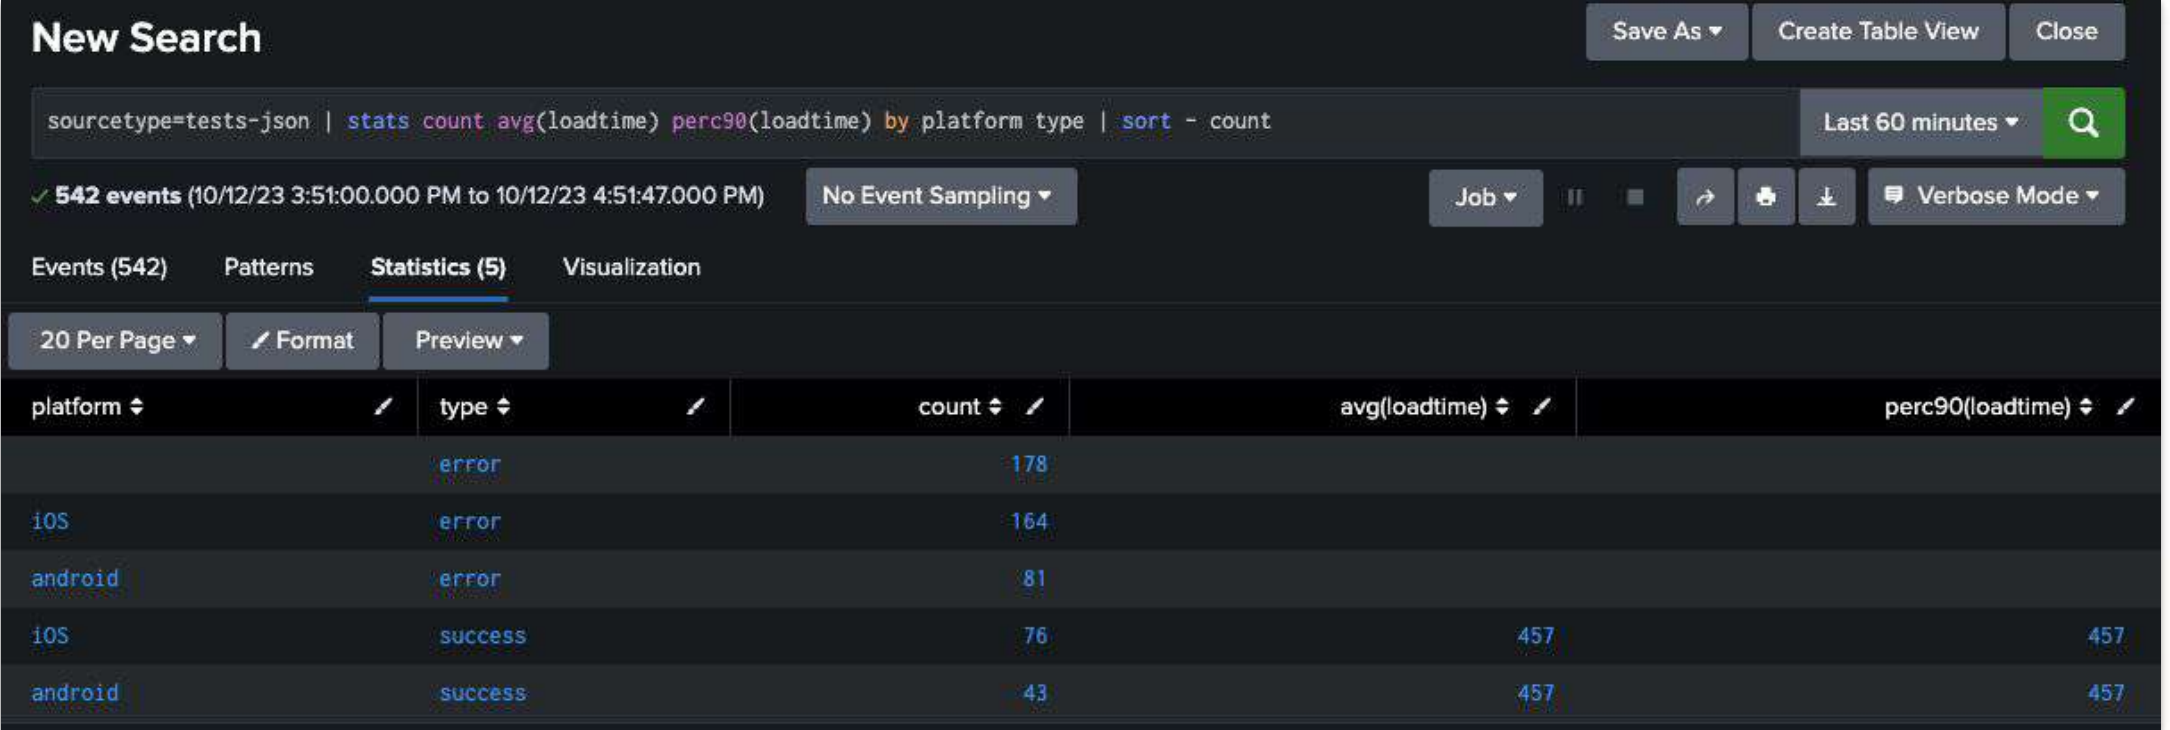
- `sourcetype=tests-json | stats count avg(loadtime) perc90(loadtime) by platform type | sort - count`

Statistical functions:

- `c / count / sum / distinct_count / dc`
- `min / max`
- `avg / stdev`
- `median / perc90 perc<int1-100>`
- `list / values`
- `earliest / latest`

Outcomes & value:

- Group data to find patterns to detect future problems, report on current situation & context



The screenshot shows the Splunk Search interface with a search bar containing the command: `sourcetype=tests-json | stats count avg(loadtime) perc90(loadtime) by platform type | sort - count`. The search results are displayed in a table with 5 columns: platform, type, count, avg(loadtime), and perc90(loadtime). The table shows 5 rows of data, grouped by platform and type.

platform	type	count	avg(loadtime)	perc90(loadtime)
	error	178		
iOS	error	164		
android	error	81		
iOS	success	76	457	457
android	success	43	457	457

<https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Stats>

Search & timechart

Use `timechart` command to turn data into timecharts

Example: Chart median loadtimes across platforms:

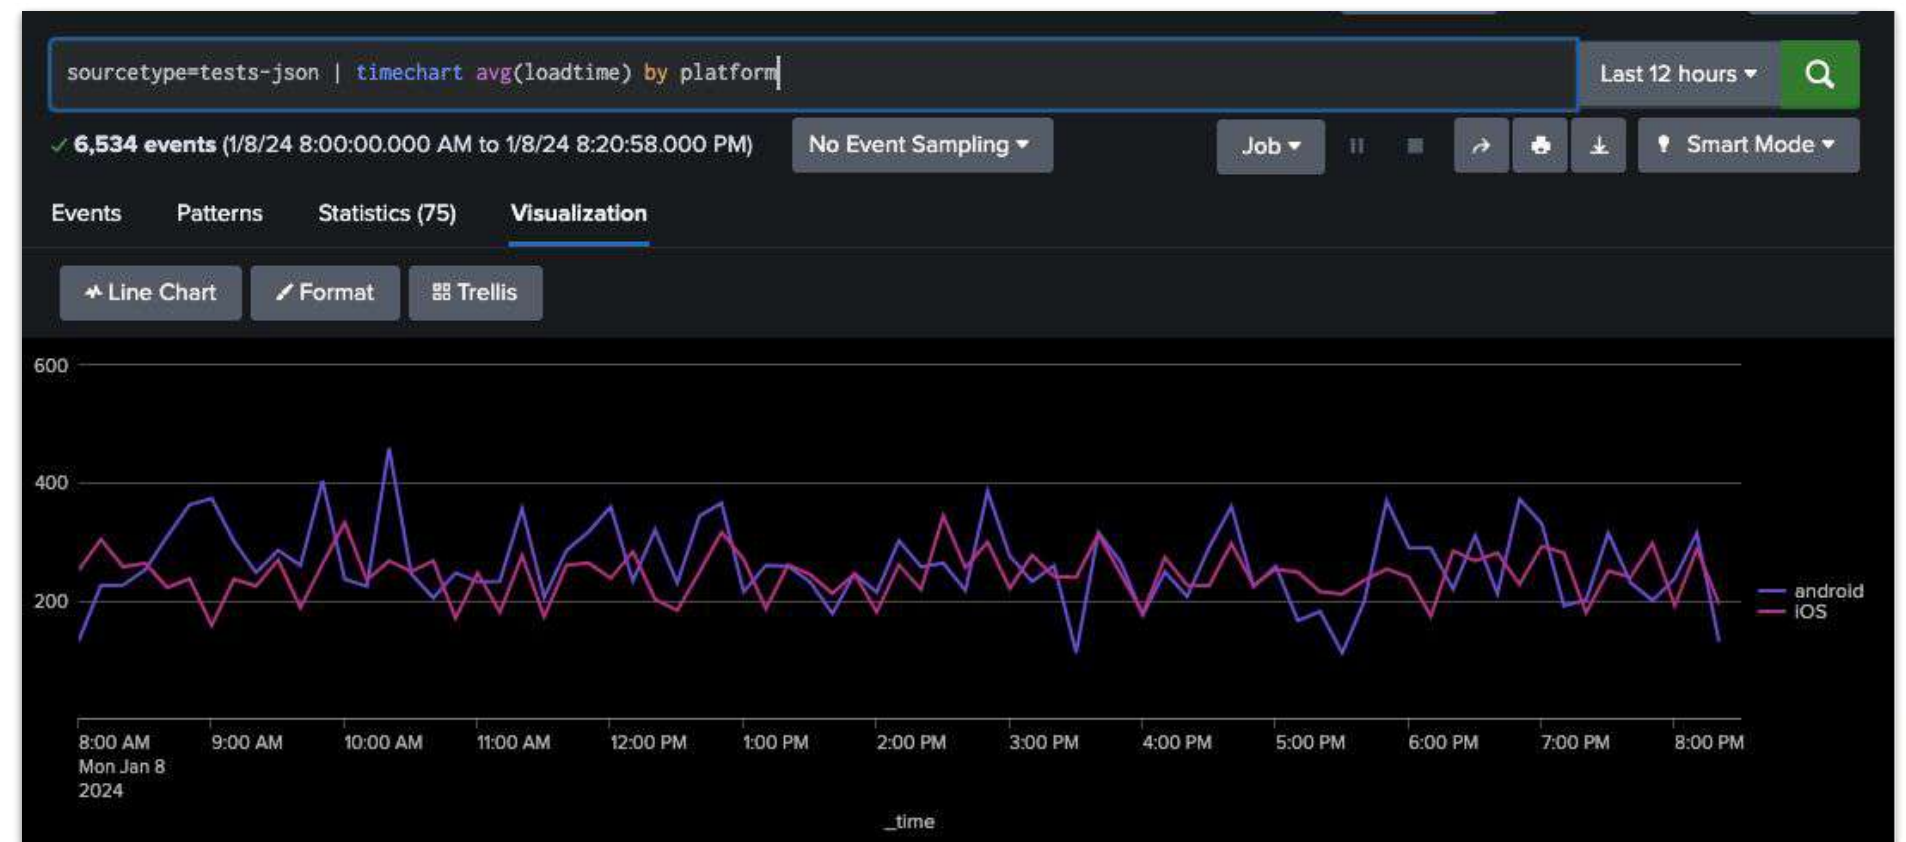
- `sourcetype=tests-json | timechart avg(loadtime) by platform`

Statistical functions:

- `count / sum`
- `min / max`
- `avg / stdev`
- `median / perc90`
- `earliest / latest`
- `distinct_count / dc / estdc`

Outcomes & value:

- Report and analyze patterns and trends, communicate anomalies to management



<https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Timechart>

Save As Report, Private Dashboard or Alert

Save any search as report, private dashboard, or alert

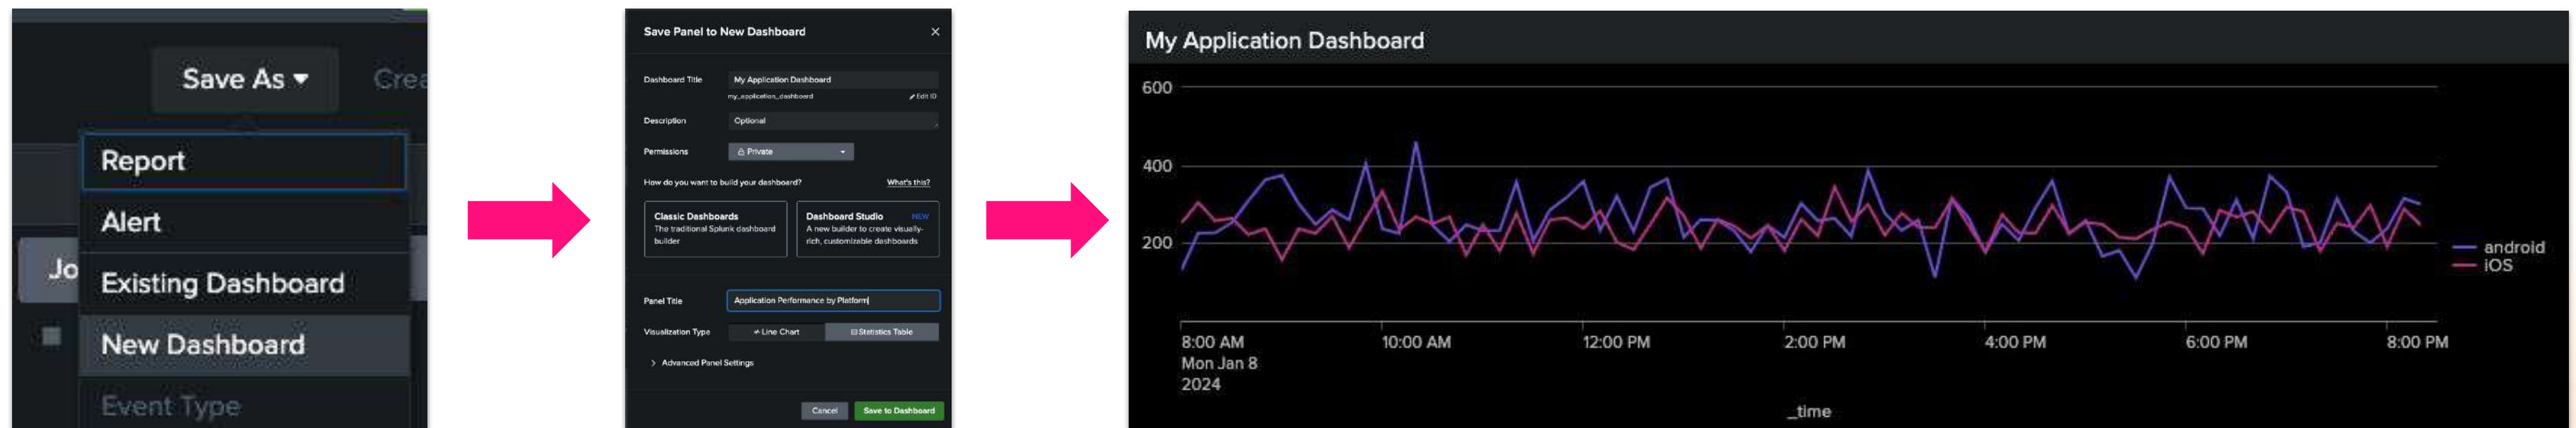
- Click "Save As" then "Report" or "Dashboard Panel" or "Alert"

Edit visualizations, add form inputs, adjust to your needs

Work with power users & Splunk admin team to share & publish content (dashboards, alerts)

Outcomes & value:

- Improve productivity by collaborating with team members using dashboards & reports
- Improve MTTD using alerts



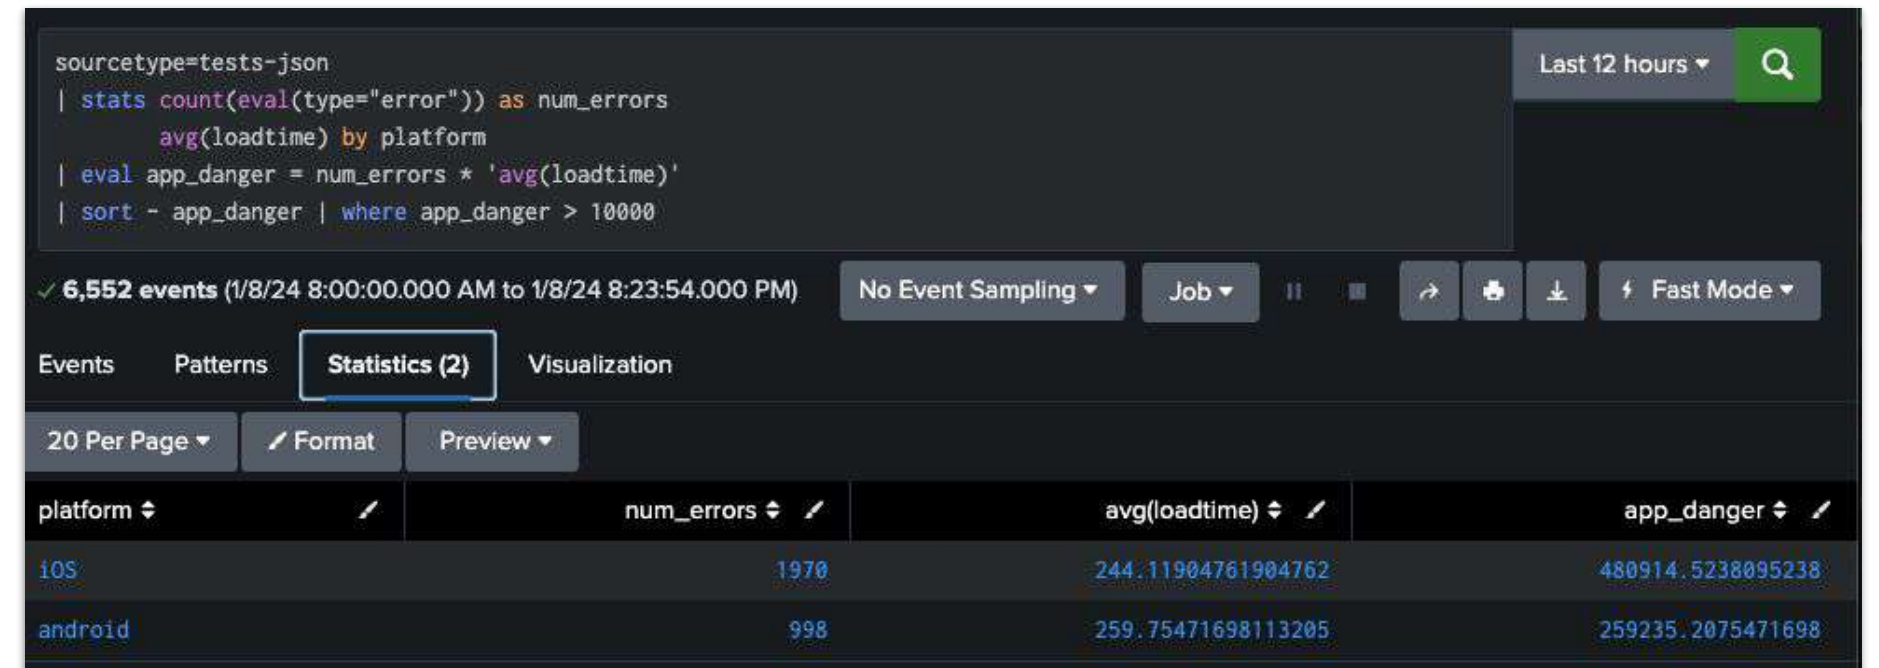
eval & where Commands

Use `eval` command to calculate expressions and create new fields

Use `where` command to filter based on Booleans

Example: Find apps with high errors or load times:

- `sourcetype=tests-json`
| `stats count(eval(type="error")) as num_errors`
| `avg(loadtime) by platform`
| `eval app_danger = num_errors * 'avg(loadtime)'`
| `sort - app_danger | where app_danger > 10000`



The screenshot shows a Splunk search interface with the following search query:

```
sourcetype=tests-json
| stats count(eval(type="error")) as num_errors
| avg(loadtime) by platform
| eval app_danger = num_errors * 'avg(loadtime)'
| sort - app_danger | where app_danger > 10000
```

The results show 6,552 events from 1/8/24 8:00:00.000 AM to 1/8/24 8:23:54.000 PM. The search is in Fast Mode. The results are displayed in a table with the following columns: platform, num_errors, avg(loadtime), and app_danger.

platform	num_errors	avg(loadtime)	app_danger
iOS	1970	244.11904761904762	480914.5238095238
android	998	259.75471698113205	259235.2075471698

`eval` / `where` functions:

- `+`, `-`, `*`, `/`
- `'fieldnames'` / `"strings"`
- `if` / `case` / `coalesce`
- `isnull` / `isnotnull`

Outcomes & value:

- Improve collaboration and detections with common field aliases and new combinations of fields

Eval command: <https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Eval>

Where command: <https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Where>

More Common Commands

Rename fields with rename:

- `sourcetype=tests-json | rename error as Error_Description`

Deduplicate repeated events with dedup:

- `sourcetype=tests-json | dedup ip`

Show only key fields in events or table view with fields/table:

- `sourcetype=tests-json | fields _time platform type error`
- `sourcetype=tests-json | table _time platform type error`

Get most recent/rarest/most common events with head/rare/top:

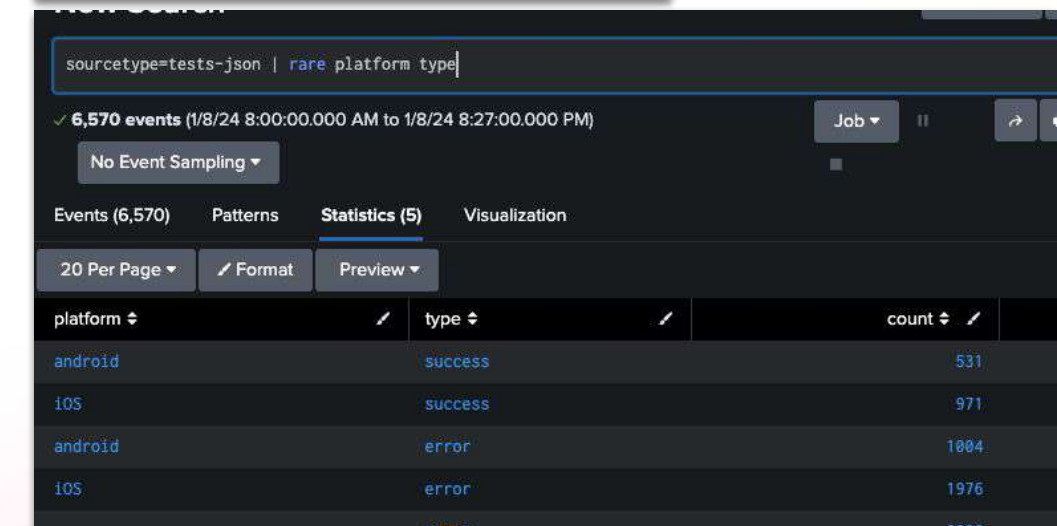

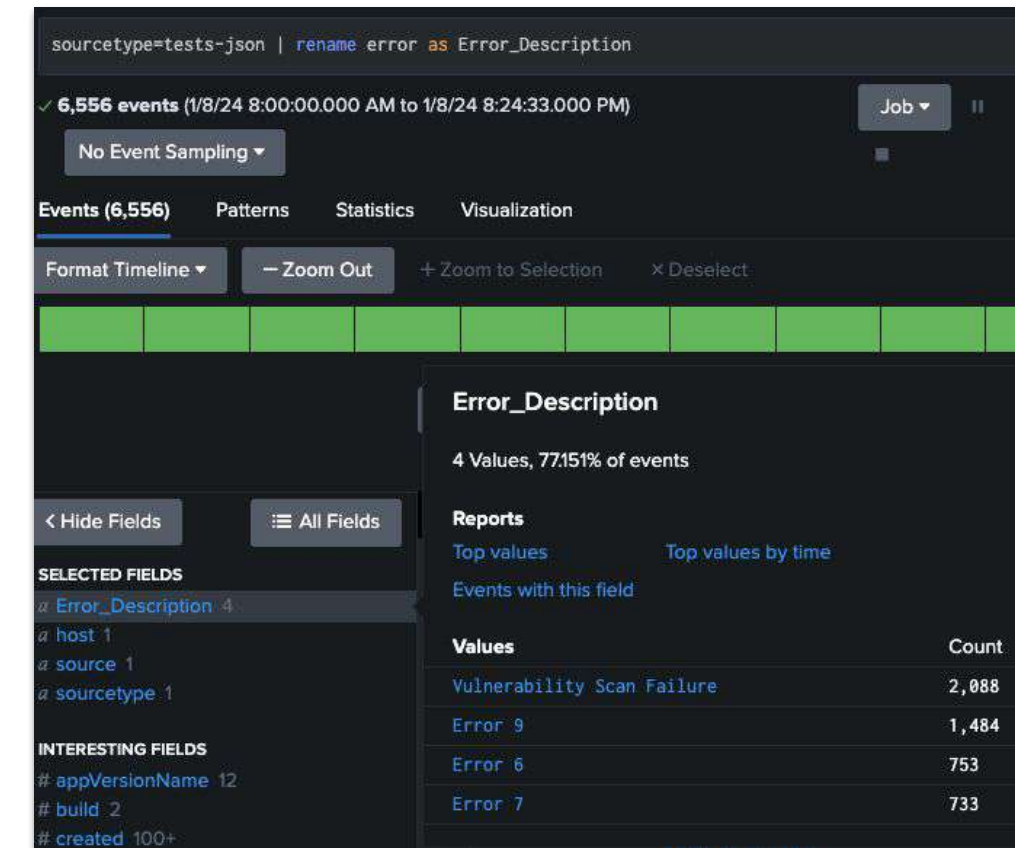
- `sourcetype=tests-json | rare platform type`

Outcomes & value:

- Detect and report on important information in your data

Plus 100+ more commands:

- <https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/ListOfSearchCommands>



Common Use Cases in Security, IT & Observability



Splunkbase

Splunkbase has 3100+ apps

Enable new use cases and extend your teams' Splunk capabilities:

<https://splunkbase.splunk.com>

Use cases include:

- IT Ops, Security, Observability, Business Analytics, IoT & Industrial Data
- Financial Services, Retail, Telecom, Healthcare, Energy

Technology vendors include:

- Cisco, Dell EMC, Amazon Web Services, Palo Alto Networks, Microsoft, Google

Please work with Splunk Admin team to ensure that you can install



Splunk Security Essentials

Reduce time to identify and resolve threats with out-of-the-box searches

Pre-built detections and data recommendations

- 1300+ pre-built detections from the Splunk Threat Research Team and data recommendations to stay ahead of existing and emerging threats

Operationalize security with industry frameworks

- Measure coverage and identify gaps in your defenses
- Maps your environment to frameworks like MITRE ATT&CK and cyber kill chain

Outcomes & value:

- Improved security detections
- More efficient Security Operations teams

Get Security Essentials for free:

<https://splunkbase.splunk.com/app/3435>

The screenshot displays the Splunk Security Essentials web interface. The top navigation bar includes links for Introduction, Security Content, Security Data Journey, Data Source Check, Documentation, and Advanced. The main content area is titled 'Security Content / Concentration of Attacker Tools by Filename'. It features a 'Description' section explaining the search's purpose, a 'Use Case' section for Advanced Threat Detection, and a 'Category' section for Endpoint Compromise. The 'Security Impact' section describes the search's effectiveness in detecting suspicious activity. The 'Alert Volume' is shown as Low, and the 'SPL Difficulty' is Basic. On the right, there are tabs for 'View', 'Demo Data', and 'Live Data'. Below the main content, there are expandable sections for 'Related Splunk Capabilities', 'How to Implement', 'Known False Positives', 'How To Respond', 'Show Search', and 'Help'. At the bottom, a table shows the search results: 1 Outlier(s) and 11,047 Raw Event(s).

Splunk IT Essentials Learn & Work

Reduce time to identify & resolve infrastructure problems with out-of-the-box searches

Learn how to deploy common IT use cases and apply those use cases in your environment

Includes 50+ use cases across:

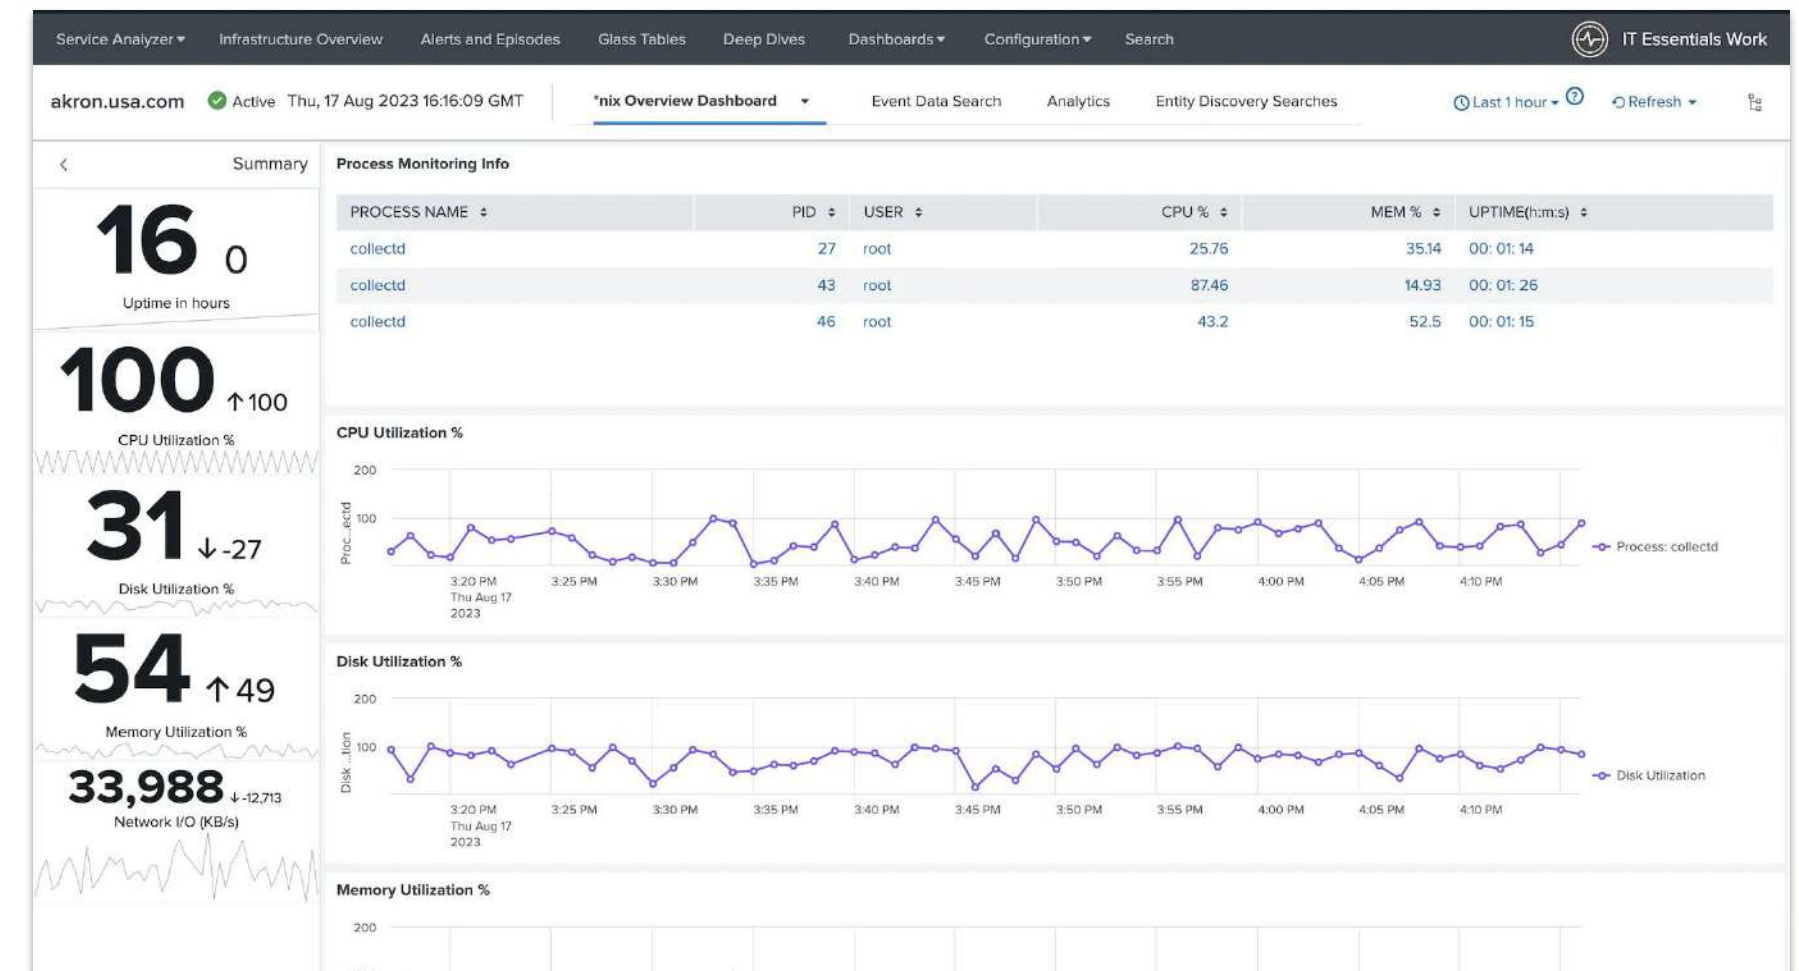
- Infrastructure Troubleshooting
- Container & Server Monitoring

Outcomes & value:

- Track infrastructure issues in one place
- More efficient IT Operations teams

Get Splunk IT Essentials Learn & Work for free:

<https://splunkbase.splunk.com/app/5390/>
<https://splunkbase.splunk.com/app/5403/>



Splunk Infrastructure Monitoring

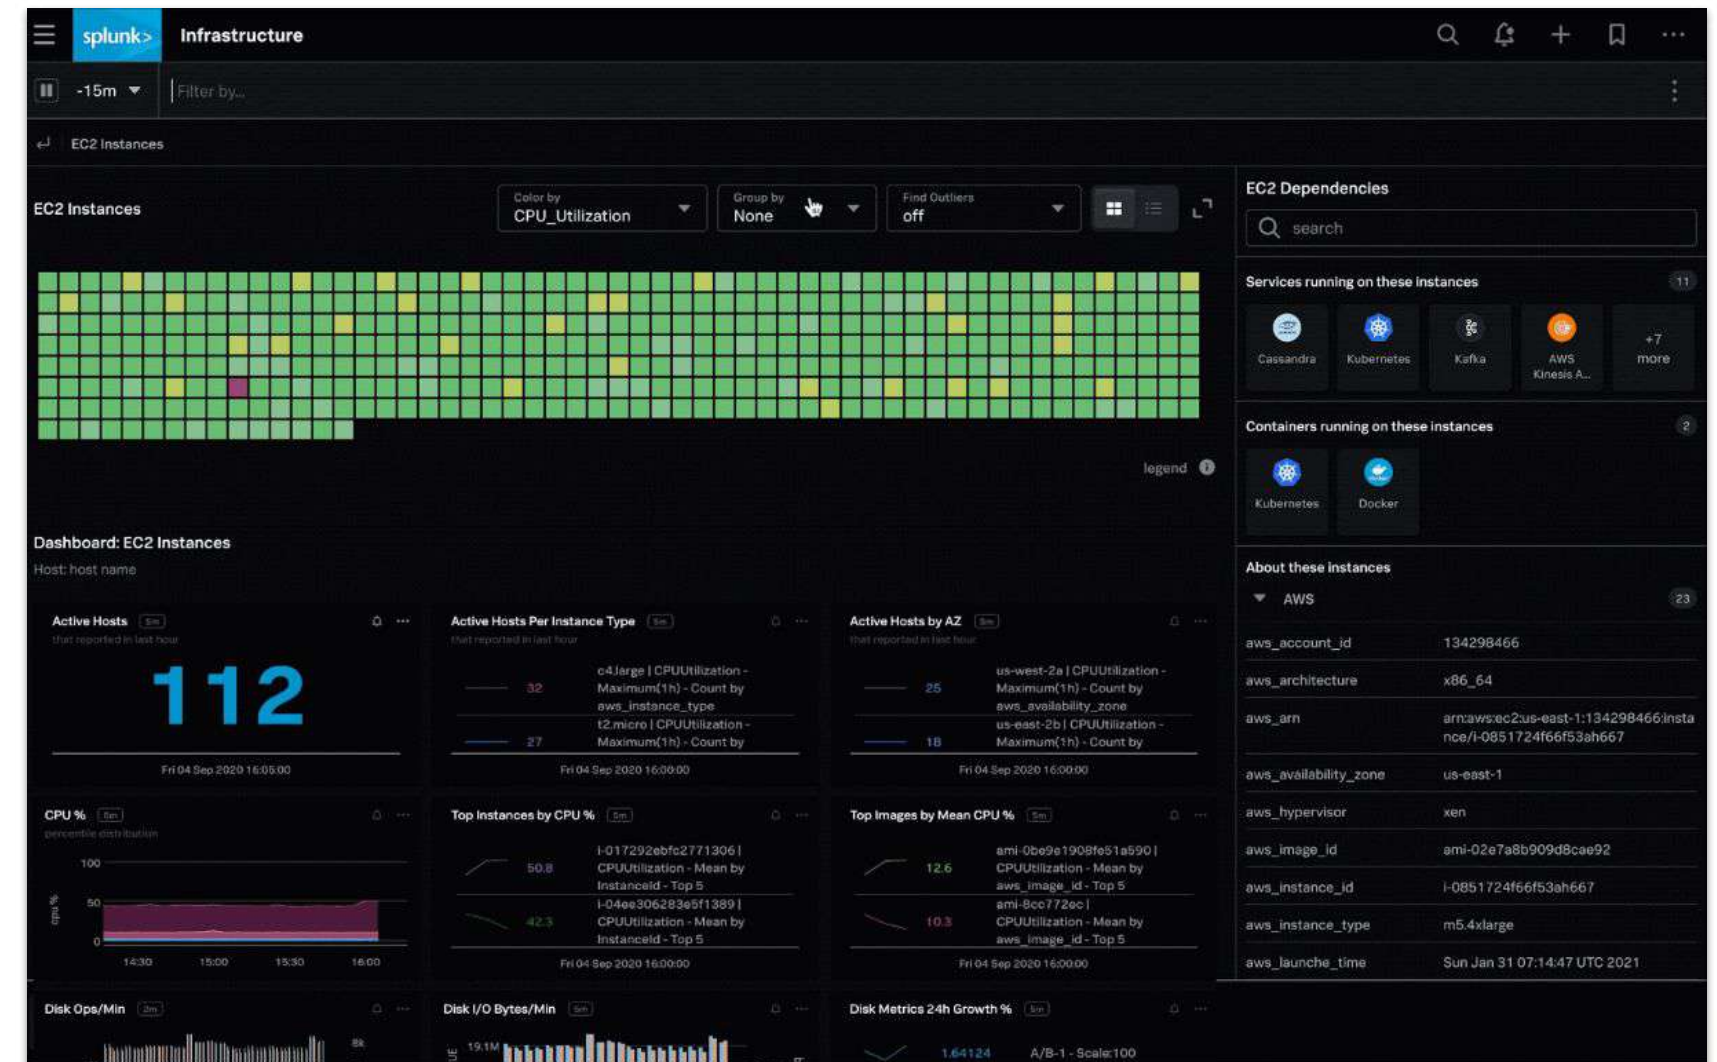
Built on a real-time streaming metrics platform

On-prem to multi cloud monitoring

- Linux, Windows, VMware
- AWS, Google Cloud, Azure
- Kubernetes & Serverless

Automatic service discovery

- 200+ pre-built integrations
- Analytics-driven alerting



Start monitoring your infrastructure with Splunk IM:

https://www.splunk.com/en_us/download/infrastructure-monitoring.html

Other Splunk Apps on Splunkbase

Technology Apps

- ▶ Amazon Web Services
 - ▶ Cisco ASA/IPS/ISE
 - ▶ DNS Analytics
 - ▶ Microsoft Exchange
 - ▶ F5 BIG-IP
 - ▶ Google Cloud Platform
 - ▶ Infoblox
 - ▶ McAfee
 - ▶ Microsoft Azure
 - ▶ Microsoft Active Directory
 - ▶ Microsoft IIS
 - ▶ Microsoft SQL Server
 - ▶ Microsoft Windows
 - ▶ Microsoft Windows DNS
 - ▶ New Relic
 - ▶ Unix & Linux
 - ▶ Oracle
 - ▶ Palo Alto Networks
 - ▶ Qualys
 - ▶ ServiceNow
 - ▶ Splunk Connect for Kubernetes
 - ▶ Splunk Connect for Syslog
 - ▶ VMWare
- + 3000 more!!

Thank you

Any questions?



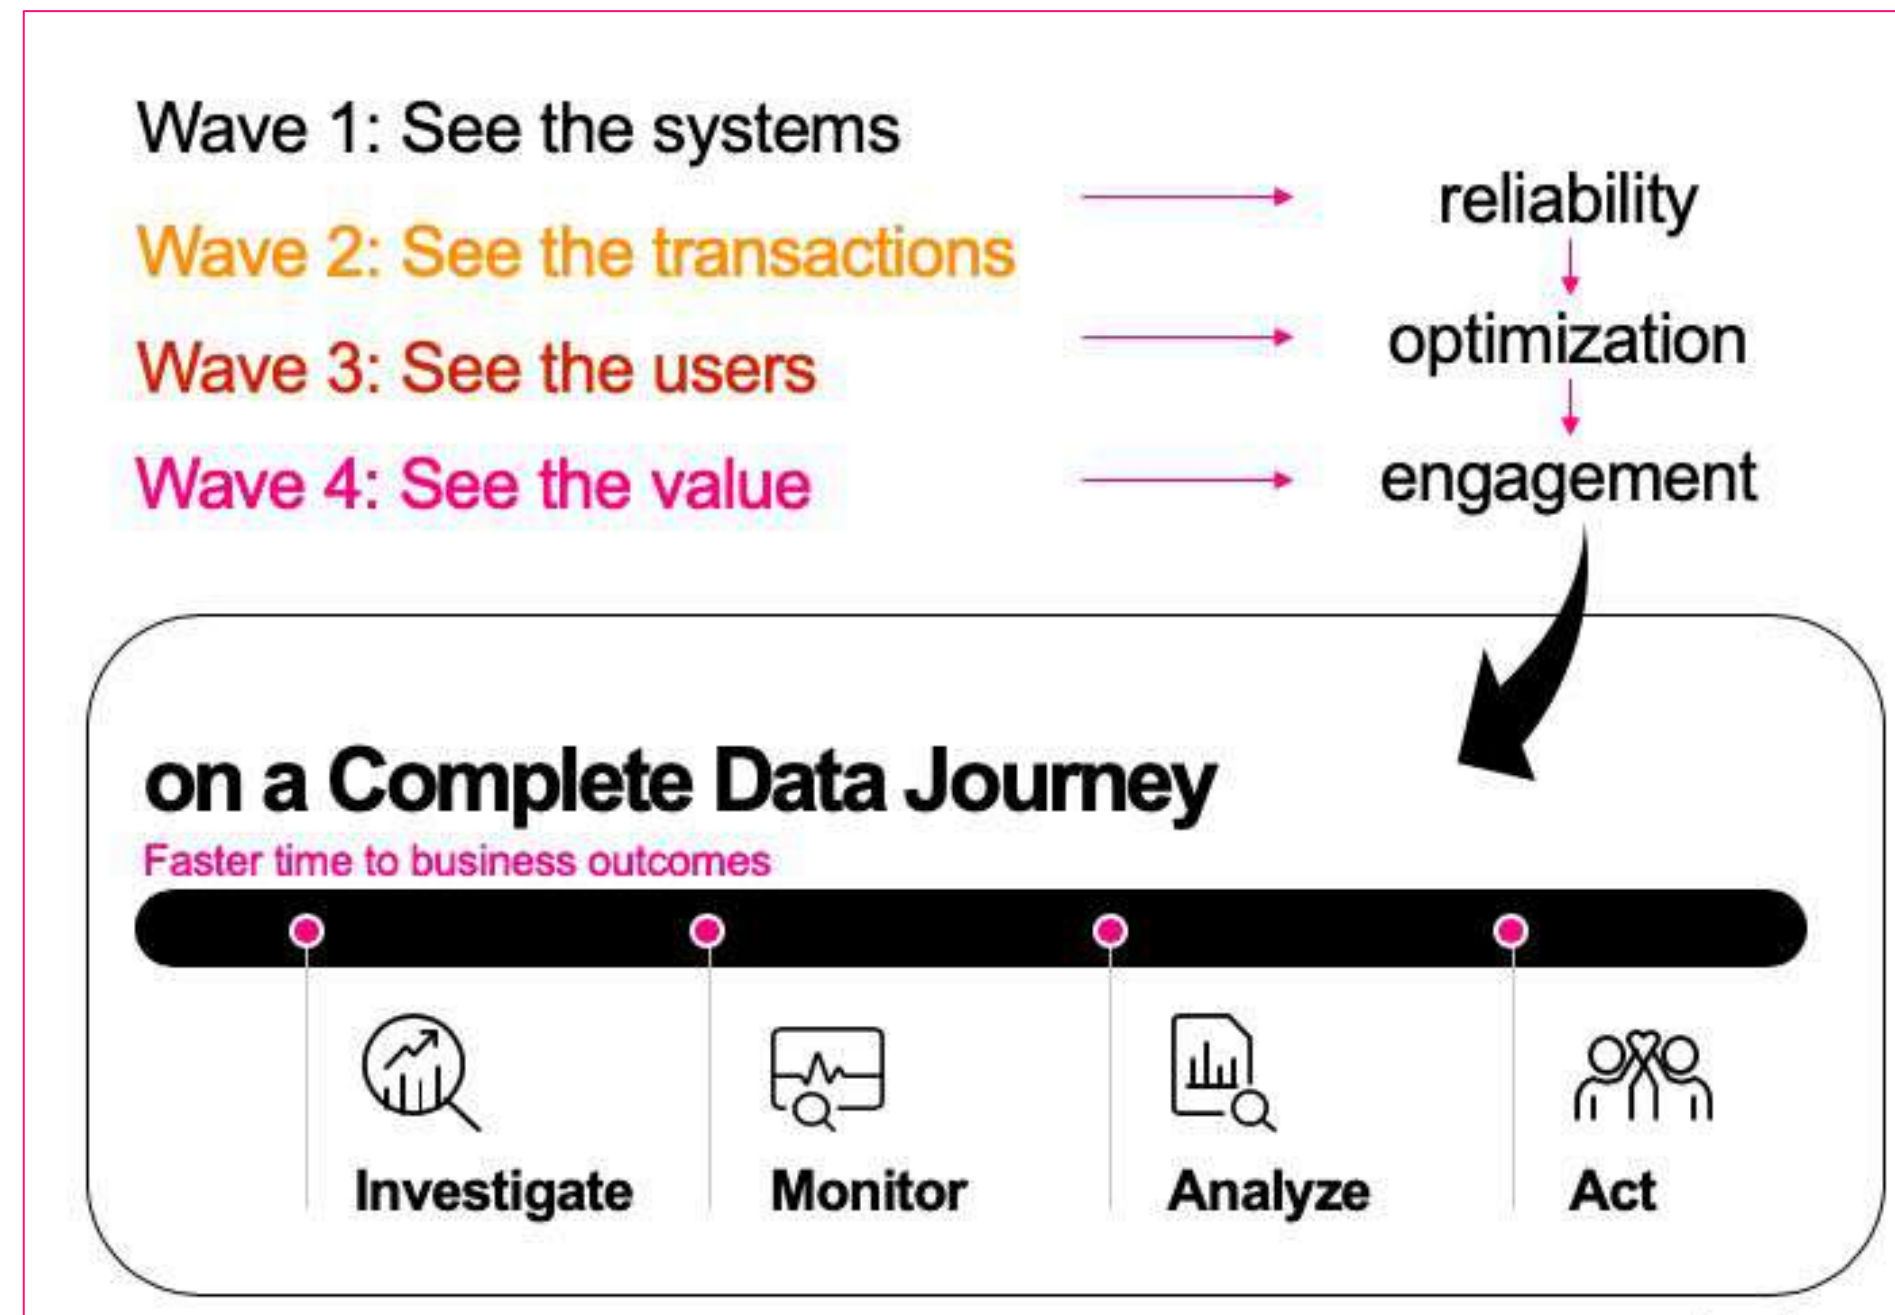
Splunk 201



Transactional Analysis



WHY transactional analysis is important

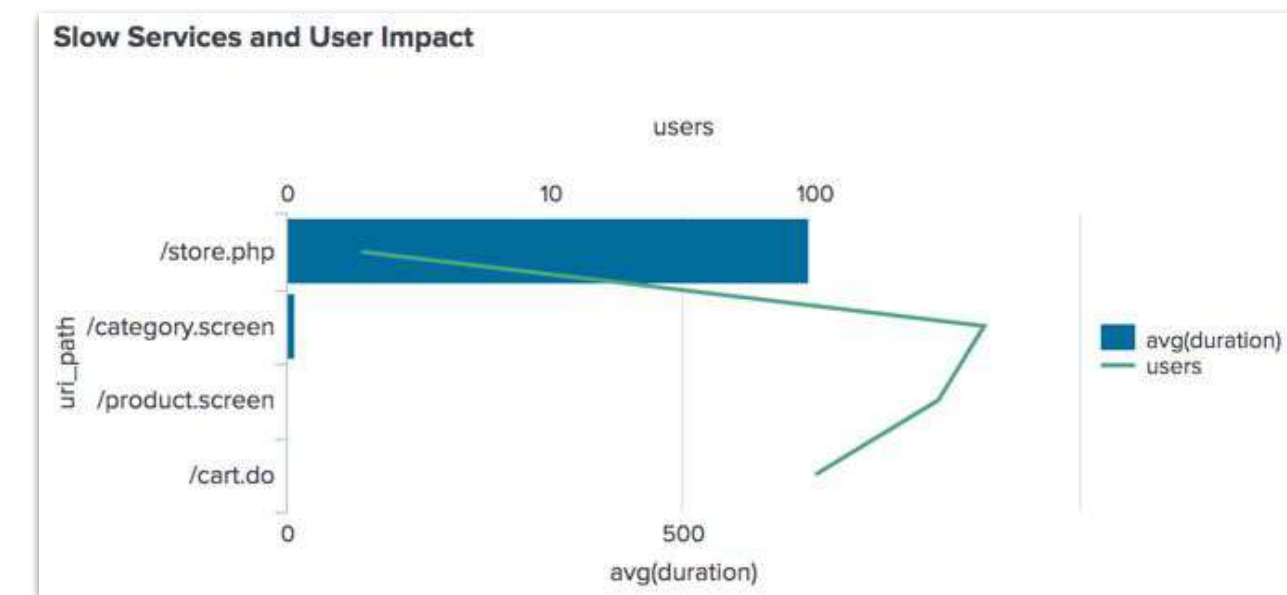


Use cases & value for transactional analysis

Transactional analysis helps you identify higher-level patterns

Use cases:

- Find impacted systems: improve system reliability
- Find impacted transactions: improve % of successful transactions & app performance/reliability
- Find impacted users: improve user engagement and value
- Find impacted process: improve revenue and efficiency



Data and events are the “atoms” of analysis

Use statistical groupings and the transaction command to build system, transactions and user profiles out of granular machine data

Correlate application logs and Application Performance Management (APM) metrics

Build transactions with stats list/values

More efficient and optimal with “stats list” or “stats value” commands

Get all details for slowest transactions using stats list or stats values and group by common ID (e.g., username, src, JSESSIONID):

- `range(_time)` AS duration - compute duration of transaction
- `list(X)` - list all values in sequential order
- `values(X)` - list deduplicated values in alphabetical order

Example:

```
sourcetype=access_combined
| fields - address, category_id
| stats range(_time) as duration
  values(action) as actions
  list(_raw) as _raw by JSESSIONID
| sort -duration
| head 10
```

20 Per Page ▾		/ Format	Preview ▾
JSESSIONID ▴ ▾	duration ▴ ▾	actions ▴ ▾	_raw ▴ ▾
SD9SAL4FF3ADFF7	3567	purchase remove view	166.110.50.46 -- [29/Apr/2013 18:09:04:101] "GET /product.screen?uid=7bf94858-e4e7-42d8-88f7-3c06e511812a276b&product_id=BS-2" Mozilla/5.0 (iPhone; CPU iPhone OS 10_0_2 like Mac OS X) AppleWebKit/119.241.19.122 -- [29/Apr/2013 18:09:04:162] "GET /category.screen?uid=0bd397e1-ccb5-4d1e-bc81-fc2329455d0da0&action=purchase&item_id=ZSG-2&product_id=MCB-5" Mozilla/5.0 (Linux; Android 6.0.1; Nexus 5X Build/MMB29P) AppleWebKit/188.157.157.176 -- [29/Apr/2013 18:09:00:146] "GET /cart.do?uid=239093d4-d1c1-4bf1-89eb-6f4755932ab2d-dfb16bef34d6&action=purchase&item_id=MCF-3&product_id=MCF-3" Mozilla/5.0 (Windows; WOW64) AppleWebKit/184.132.8.30 -- [29/Apr/2013 18:09:04:173] "POST /category.screen?uid=1341c65e-537a-4e12-8fb3-78b59c75ebcd82f&product_id=CM-1" Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_1) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.0 (iPhone; CPU iPhone OS 10_2 like Mac OS X) AppleWebKit/602.2.1.202.249.116.52 -- [29/Apr/2013 18:08:59:169] "POST /cart.do?uid=d48b3196-605c-41d9-a11f-711c5bf5c444d-a701-2f32cf7813be&product_id=DFS-2" Mozilla/5.0 (Windows; WOW64) AppleWebKit/537.36 (KHTML, 244.221.135.140 -- [29/Apr/2013 18:09:05:132] "GET /category.screen?uid=a1124351-79e5-44e5-9650-f7140835d184b&category=Misc" Mozilla/5.

Build transactions with transaction command

Less efficient but easy with “transaction” command

Get all details for slowest transactions using transaction command

```
sourcetype=access_combined
| transaction JSESSIONID username
| sort -duration
| head 10
| table JSESSIONID duration username uri_path *
```

Warning: the “transaction” command is not efficient!

- Use the “stats list” or “stats value” commands if you can
- Talk to your Splunk admin team
- Read more: <https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Transaction>

JSESSIONID	duration	username	uri_path	action	address	bytes	category	category_id	clientip	clo
SD1SAL1FF6ADFF4	2841	edavisdj	/cart.do	purchase view	8878 Oakridge Plaza	259 2868	Misc		201.36.233.151 54.115.1.118	
SD3SL1FF7ADFF8	2829	wcampbellnz	/cart.do /category.screen	purchase	37566 Logan Plaza	1829 3538	Misc	Misc	10.2.1.33 219.135.155.193	
SD8SL9FF9ADFF8	2658	wmorganqa	/cart.do	changequantity	122 Buses	1274 308	Misc		210.204.91.255 60.87.149.135	

stats list/values vs transaction command

stats list and stats values commands are good in most circumstances:

- Efficient transactions
- Transactions with common correlation fields or transaction ID
- stats list - all values, sequential
- stats values - deduplicated values, alphabetical

transaction command can be used in a pinch:

- Quick prototypes
- Complex transactions (startswith, endswith, maxspan)

Warning: the “transaction” command is not efficient!

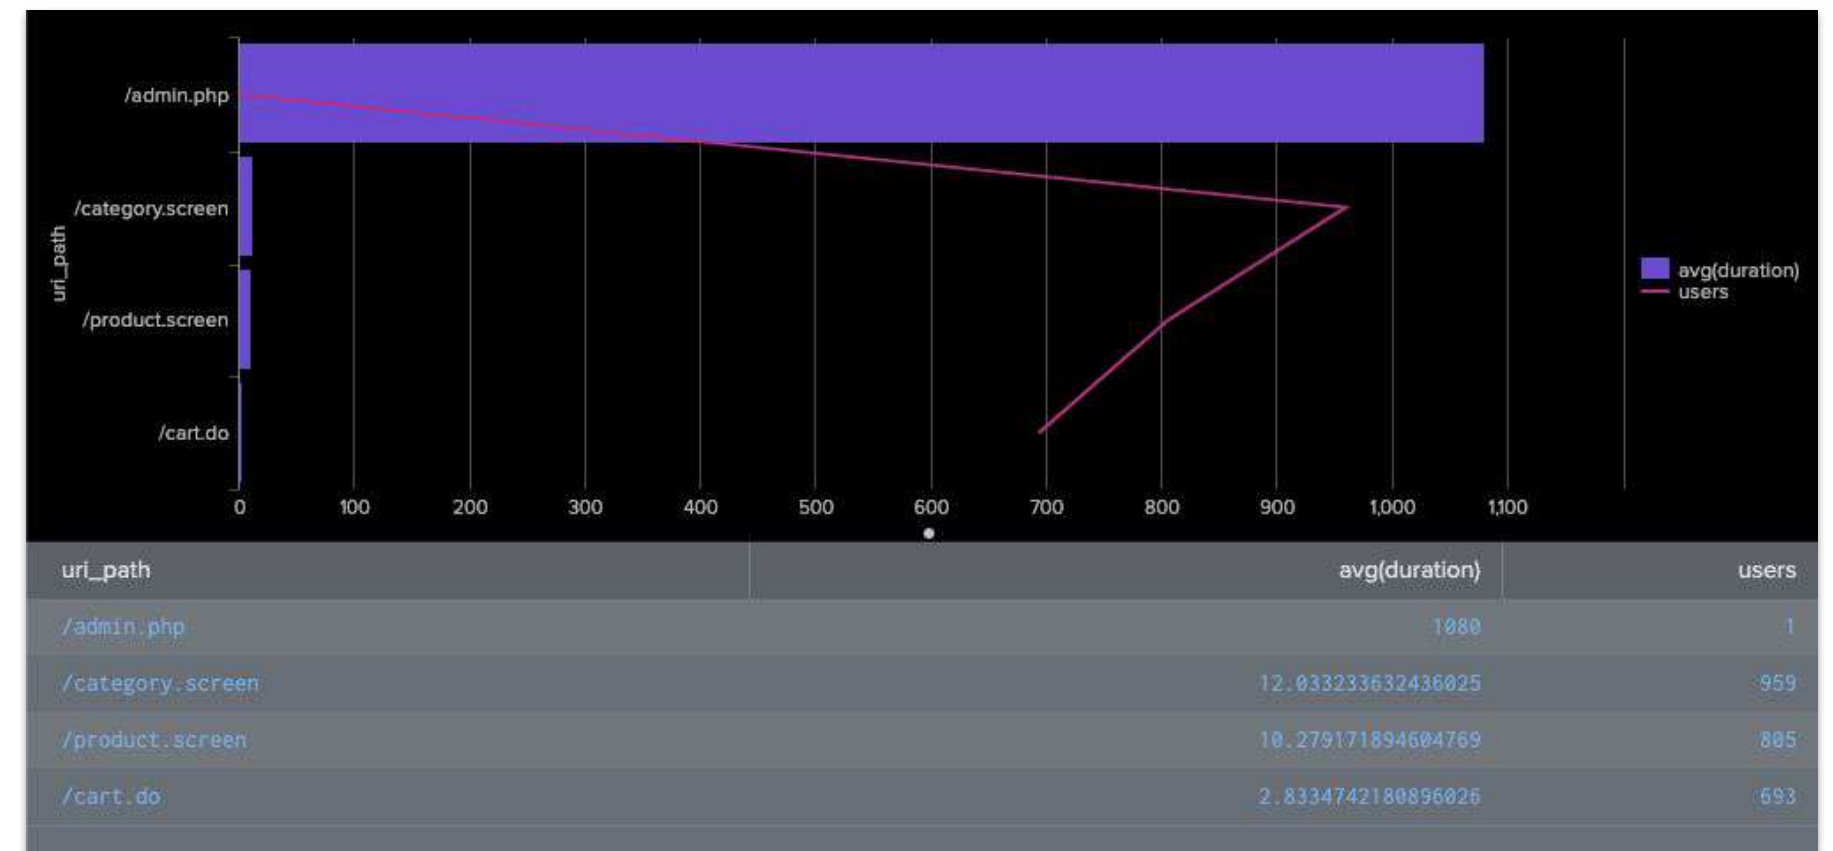
- Use the “stats list” or “stats value” commands if you can
- Talk to your Splunk admin team

Analyze transactions with additional stats commands

Analyze transactions with additional aggregations:

- count
- sum
- avg / stdev
- median / perc
- dc / estdc

Example: Find slow services and user impact by aggregating transactions:



```
sourcetype=access_combined
| stats values(action) as actions range(_time) as duration
  by JSESSIONID username src uri_path
| stats avg(duration) dc(username) as users by uri_path
| sort - "avg(duration)"
```

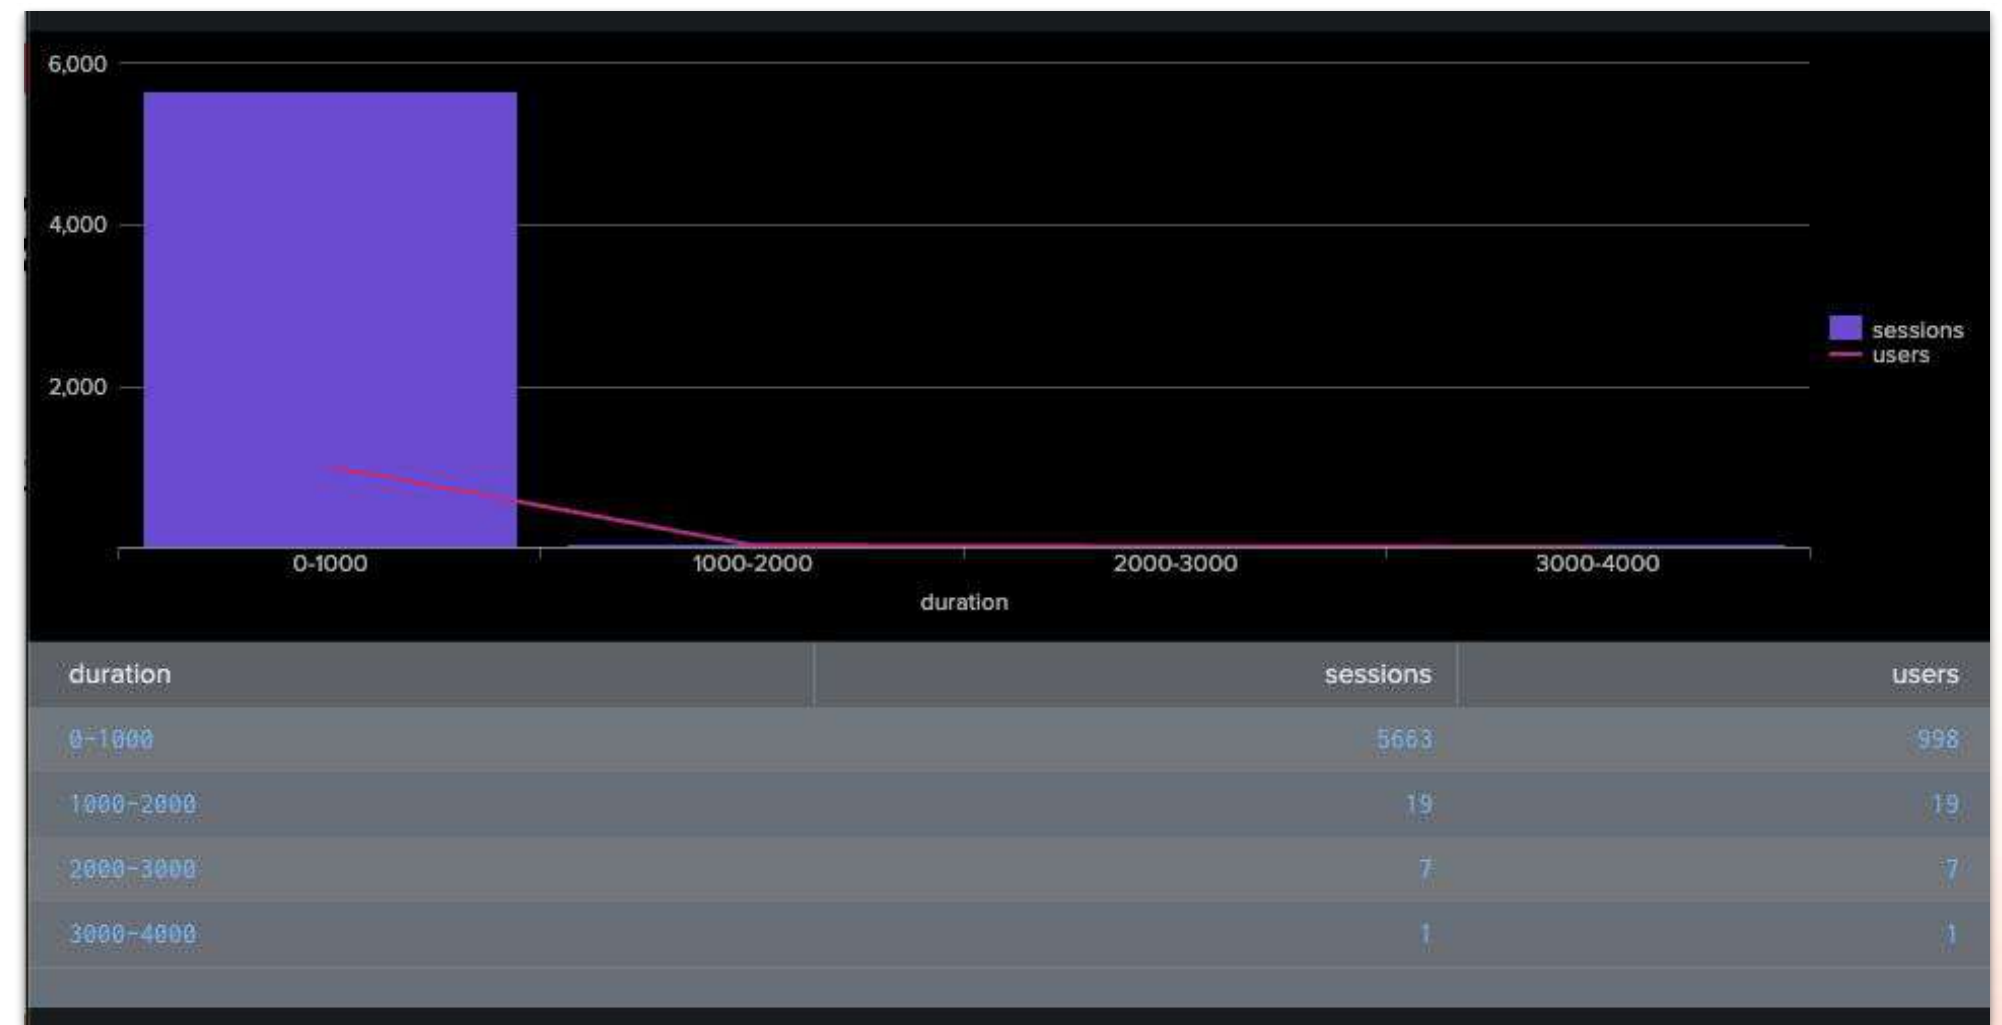
Monitor systems with histograms & chart command

Compute histograms for monitoring real-time transactions

Use `chart` command with `bins` parameter, or standalone `bin` command

Monitor transaction speeds and user impact:

```
sourcetype=access_combined
| stats range(_time) as duration
  by JSESSIONID username
| chart count as sessions
  dc(username) as users
  by duration bins=10
```



Advanced Statistics




Comparing statistics over time

Track the “typical” value using averages and medians. Track the fluctuations using stdev and percentiles. If there are problems, likely to modify either typical values or fluctuations.

Compare different transaction speeds over time:

```
sourcetype=access_combined
| stats earliest(_time) as _time range(_time) as duration by JSESSIONID username
| timechart avg(duration) stdev(duration) perc25(duration) median(duration) perc75(duration)
```



Advanced Statistics vs Machine Learning

Splunk makes it easy to do multiple types of statistics:

- **(101) Statistics:** aggregate and analyze numerical data split by groups (counts, sums, avgs & medians)
- **(201) Transactional Statistics:** analyze higher-level entities and transactions from raw data
- **(201) Advanced Statistics:** incorporate variability (stdev/percentiles), apply eventstats & streamstats
- **(301) Machine Learning:** use statistical models to detect anomalies and sudden changes, adjust thresholds dynamically, and predict service degradation

References:

- List of stats functions:
<https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/CommonStatsFunctions>
- Download the Splunk Machine Learning Toolkit (MLTK):
<https://splunkbase.splunk.com/app/2890/>
- Download Python for Scientific Computing add-on for your OS:
<https://splunkbase.splunk.com/app/2882>

Use standard deviation to find anomalies

Apply standard deviation and variance to transactions like sessions and users.

This is very useful when looking for bots or looking at trends across segments of users

username ↕	avg(clicks_per_minute) ↕	stdev(clicks_per_minute) ↕
fsamuels	1.9473684210526316	0.229415
fholmesfc	1.0909090909090908	0.30151
mfoxch	1.0909090909090908	0.30151
tscottj4	1.0909090909090908	0.30151
areyese7	1.1	0.316227
astewartp6	1.1	0.316227
darmstrongr8	1.1	0.316227

Example: Detect bots in web logs (very low standard deviation, constant clicks/minute)

```
sourcetype=access_combined
| bin _time span=1m
| stats count as clicks_per_minute by _time username
| stats avg(clicks_per_minute) stdev(clicks_per_minute) by username
| where 'avg(clicks_per_minute)' > 1
| sort + "stdev(clicks_per_minute)"
```

Peer-based adaptive thresholds with eventstats

Use eventstats to build adaptive thresholds based on peer groups (services, application tags, business units), and and track service health over time.

Example: Find all transactions with very slow durations, with variable speed impact by service

```
sourcetype=access_combined
| stats earliest(_time) as _time range(_time) as duration by JSESSIONID uri_path
| eventstats avg(duration) stdev(duration) by uri_path
| where duration > 'avg(duration)' + 2*'stdev(duration)'
```

JSESSIONID ↕	uri_path ↕	_time ↕	duration ↕
SD10ASL10FF1ADFF2	/cart.do	2024-01-17 19:34:57	2731
SD10ASL10FF3ADFF1	/product.screen	2024-01-17 19:34:44	3309
SD10ASL3FF10ADFF7	/cart.do	2024-01-17 19:35:05	2850
SD10ASL4FF3ADFF5	/cart.do	2024-01-17 19:32:00	3114
SD10ASL5FF10ADFF5	/cart.do	2024-01-17 19:30:45	2906
SD10BSL10FF5ADFF3	/product.screen	2024-01-17 19:33:23	3353
SD10BSL3FF3ADFF6	/cart.do	2024-01-17 19:31:28	3146
SD10BSL4FF3ADFF5	/product.screen	2024-01-17 19:32:15	3151

Historical adaptive thresholds with streamstats

Use streamstats to build historical adaptive thresholds and track service health over time.

Find anomalous transactions which are impacting service health (compute rolling avg and stdev of duration by service, find durations which are 2 stdevs slower than historical avg)

```
sourcetype=access_combined
| stats earliest(_time) as _time range(_time) as duration by JSESSIONID uri_path
| streamstats avg(duration) stdev(duration) by uri_path
| where duration > 'avg(duration)' + 2*'stdev(duration)'
```

JSESSIONID ↕	uri_path ↕	_time ↕	duration ↕
SD10BSL10FF5ADFF3	/product.screen	2024-01-17 19:33:23	3353
SD10CSL5FF3ADFF4	/product.screen	2024-01-17 19:33:15	2782
SD10SL5FF10ADFF4	/product.screen	2024-01-17 19:34:57	2831
SD10SL5FF6ADFF5	/product.screen	2024-01-17 19:37:09	2827
SD10SL6FF4ADFF6	/product.screen	2024-01-17 19:35:05	3029
SD10SL6FF9ADFF9	/cart.do	2024-01-17 19:33:15	3509
SD10SL9FF10ADFF7	/product.screen	2024-01-17 19:31:39	3600
SD1SAL3FF4ADFF9	/category.screen	2024-01-17 19:31:18	3576

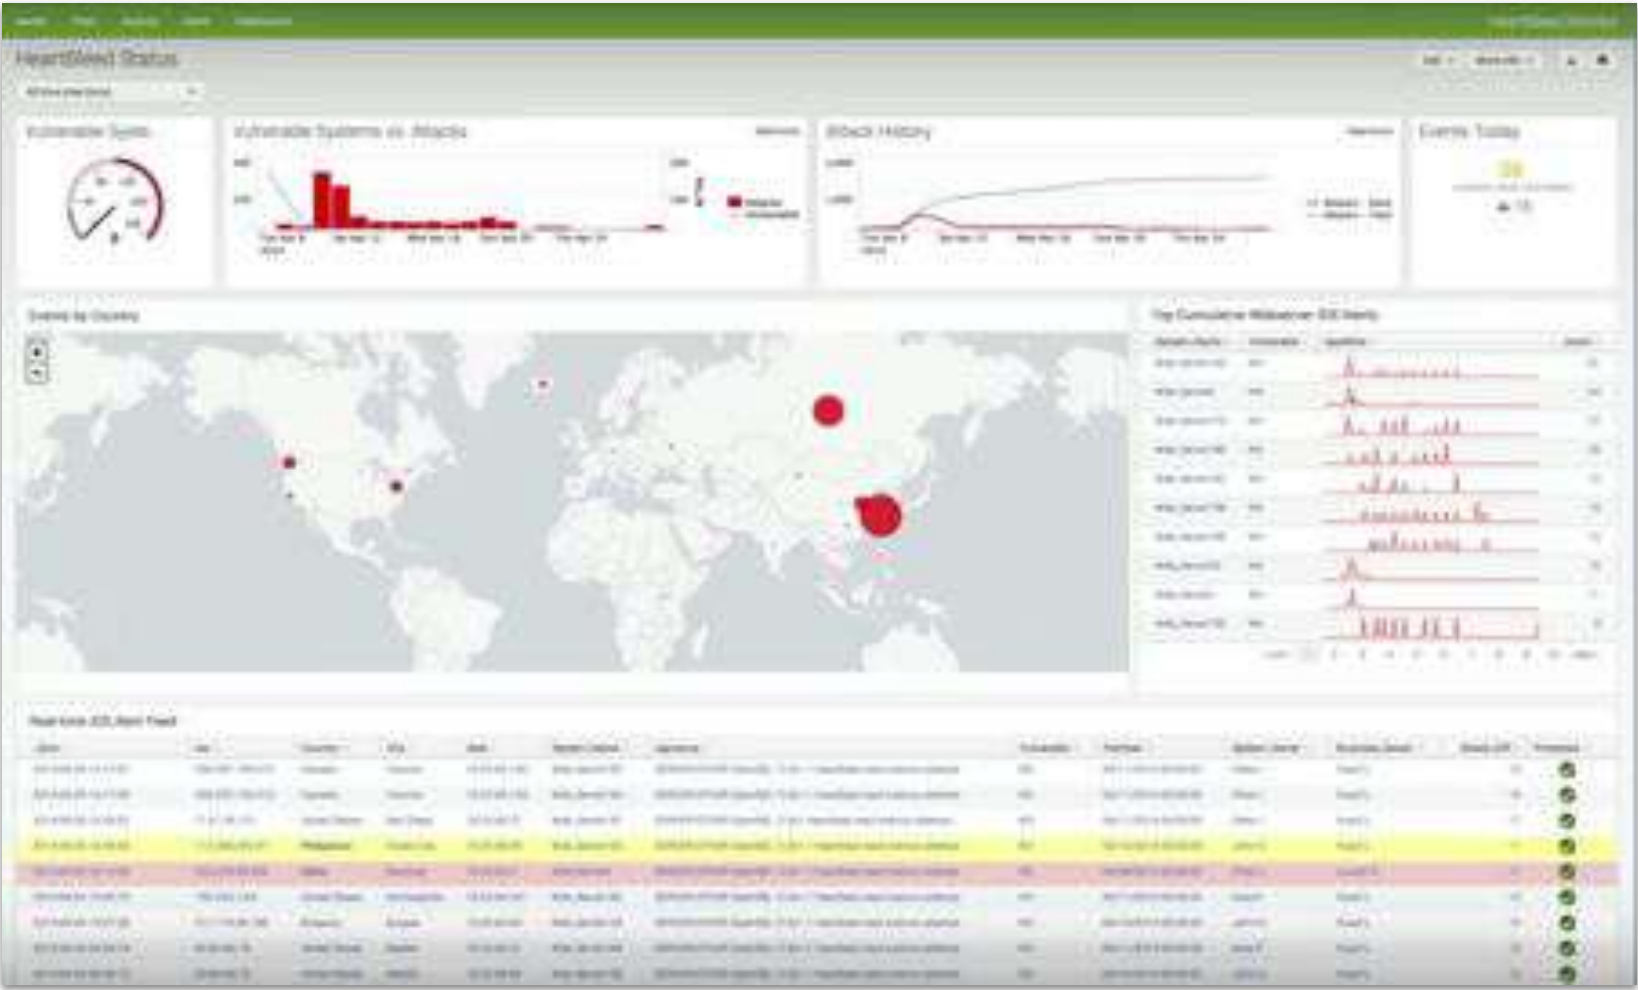
Advanced Dashboarding

splunk>



Examples of Good Dashboards

NASDAQ's Heartbleed Status dashboard for CISO/CIO:

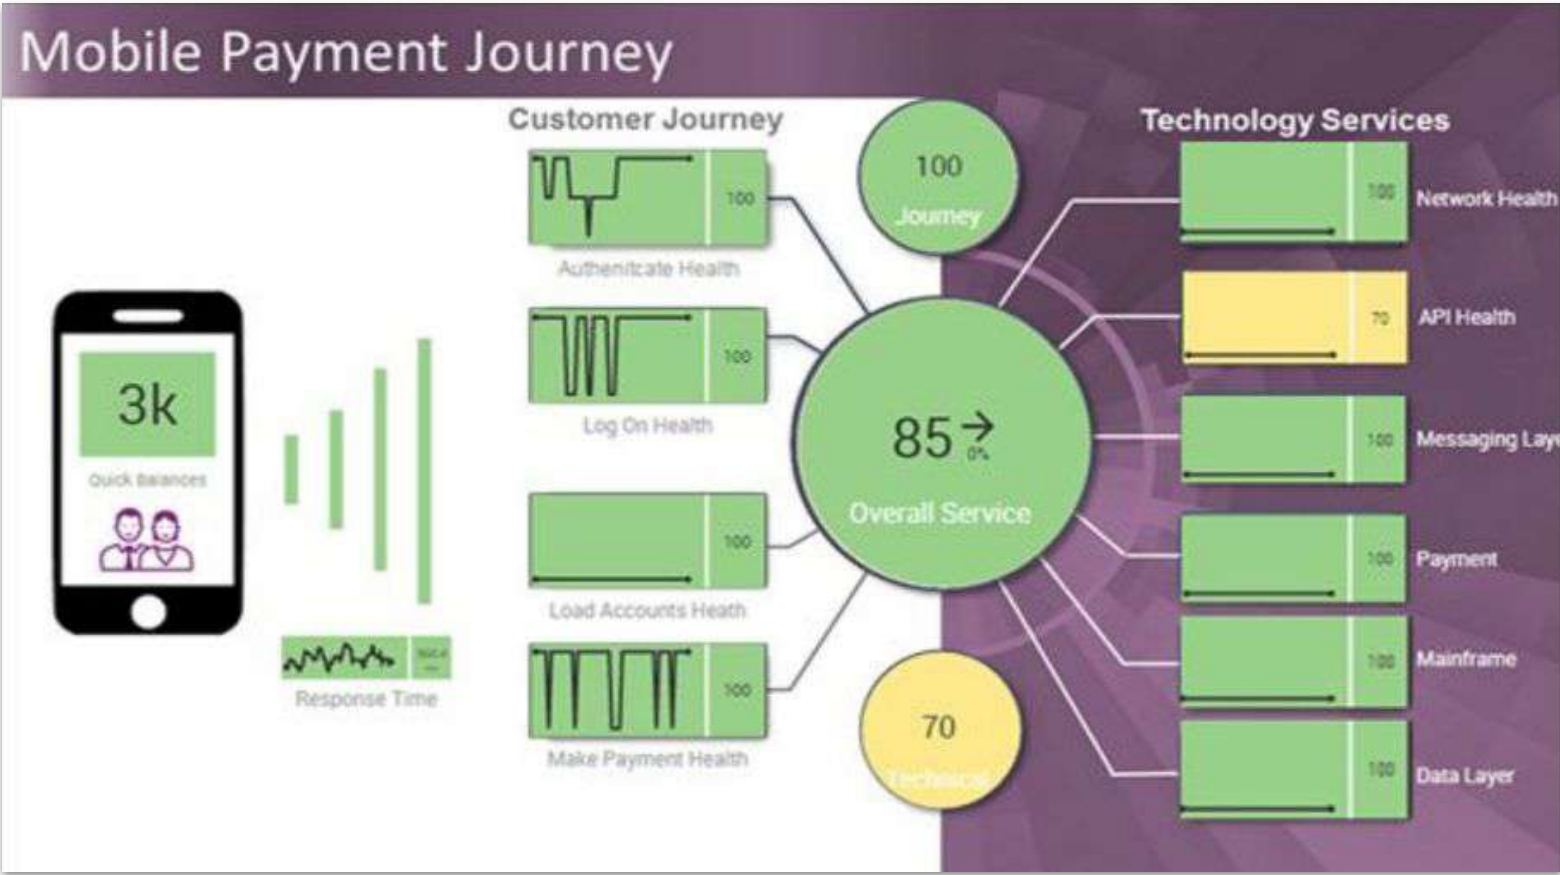


Dashboard showing patient diagnostics:



Examples of Good Dashboards (cont.)

Allied Irish Bank’s Mobile Payment Journey:



API Gateway Performance and Health:





Dashboarding Best Practices

Understand Requirements

- **Document** stakeholders' needs and context
- **Align** your strategy & metrics to the Business Goals
- **Be prepared** with the key projects, the most vital ones needing approval, and strategy in simple numbers. E.g., “if we go SaaS, reduce costs by \$5 million”
- **Don't** build your dashboard in a vacuum

Design Dashboard

- **Define** required charts and metrics
- **Show** Single Values 100 and Trending Charts  before showing Raw Data 
- **Build** a simple prototype/minimum viable product, then iterate with stakeholder.
- **Create** a dashboard in the Splunk app for the population that will view it (e.g., Trading dashboard in Trading team's app workspace)

Tell a Story

- **Focus** on what your audience NEEDS to see, not what's easy to calculate
- **Use** Real Life examples and stories that are contextually relevant to the business
- **Provide** to stats and competing organizations is helpful.

Dashboarding Discussion & Homework

- What is the best dashboard you've seen?
- Why was it so good?
- Who was the stakeholder?
- What value did it give them?

Form Inputs and Tokens

Make dashboards easy to use with form inputs

E.g., split by time, service/application, user, host

Add a form input

- For time input: adjust the time picker on each panel
- For other inputs: use token syntax in the panels:

search host=\$host\$ username=\$username\$

Best practice:

- Use Static Options for small number of choices, and ALL:
 - Name: ALL
 - Value: *
- Use Dynamic Search to populate data-driven set of choices:
 - | **tstats** **count** **where** sourcetype=access_combined **by** host | **sort** host
 - sourcetype=access_combined | **stats** **count** **by** username | **sort** username

The screenshot shows the 'Slowest Transactions Navigator' interface. At the top, there are two filter inputs: 'Hosts' with a dropdown menu showing 'webserver-01' and a close button 'X', and 'Username' with a dropdown menu showing 'fsamuels' and a close button 'X'. To the right of these is a 'Hide Filters' link. Below the filters is a section titled 'Top Slowest Transactions' containing a table. The table has six columns: 'JSESSIONID', 'duration', 'username', 'uri_path', 'bytes', and 'clientip'. The first row of data shows 'SD3SL1FF7ADFF9' for JSESSIONID, '811.160461' for duration, 'fsamuels' for username, '/store.php' for uri_path, and '10645', '10841', '11146' for bytes, and '175.45.177.7' for clientip.

JSESSIONID	duration	username	uri_path	bytes	clientip
SD3SL1FF7ADFF9	811.160461	fsamuels	/store.php	10645 10841 11146	175.45.177.7

The screenshot shows the 'Static Options' and 'Dynamic Options' configuration interface. Under 'Static Options', there is a table with 'Name' and 'Value' columns. A single row is visible with 'Name' as 'ALL' and 'Value' as '*'. There is a '+ Add New' link below this table. Under 'Dynamic Options', there is a 'Content Type' dropdown menu with a search icon and a 'Search String' input field. The search string is 'sourcetype=access_combined | stats count by username | sort username'.

Name	Value
ALL	*

+ Add New

Content Type: [Search Icon] [Dropdown Icon]

Search String: sourcetype=access_combined | stats count by username | sort username

Post Process Search

Build “base searches” (with id tag) on your dashboard which pull the data, then “post-process searches” (with base tag) which process from same base dataset.

Best practice: Place values and charts above data and use shared time picker

Step 1) start with a base search generating transaction:

```
<search id="all_web_sessions">
  <query>sourcetype=access* | stats earliest(_time) as _time values(*) as * range(_time)
as duration by JSESSIONID uri_path</query>
  <earliest>$time_picker.earliest$</earliest>
  <latest>$time_picker.latest$</latest>
</search>
```

Step 2) Use post-process searches to compute # sessions & avg duration by service:

```
<search base="all_web_sessions">
  <query>| stats count as num_sessions by uri_path</query>
</search>
```

```
<search base="all_web_sessions">
  <query>| timechart avg(duration) by uri_path</query>
</search>
```

Dashboard Studio

Deliver high-value dashboards & reports to business & operational teams with Dashboard Studio

Example Hub has lots of examples for you to use and customize

eCommerce Monitoring & Performance



VPN Health by Region



Geography-Based Analysis

Geography is important because all data takes place in time AND space

Example use cases:

- DDOS attack from Nation State
- Fast-moving logons or badge swipes
- Find out where web-based transactions are coming from
 - why do we have latency for some transactions?
- Is our sales/marketing strategy in agreement with geography-based segments and outcomes?



Best practices:

- Calibrate geo views to your audience. Show a GOOD map with good info to make decisions
- Don't use "pew pew" laser maps because they're easy, make sure it's valuable
- Use accurate geo info. You can buy updated MaxMind geo database in Splunk, or enrich lat/lon from your own DB

View geographic data with cluster maps

Use cluster maps for localized information (specific latitude/longitude)

Prepare data with iplocation+geostats command

Find all activity by specific locations:

```
sourcetype=access_combined  
| iplocation clientip  
| geostats count by status
```



Learn more about Cluster Maps:

<https://docs.splunk.com/Documentation/Splunk/latest/Viz/MarkerMap>

Regex/rer commands

Extract fields at search time with `rer` & `regex` commands. Develop regex and pass back to Splunk admin team.

Best practice: Use a website like regex101.com and Google search to practice regex. E.g., “How do I build regex for: everything up to the first semicolon?”

Example: Splunk admin should already extract timestamp

1573704878 200 - User fsamuels failed login

```
| rer field="_raw" "(?<status>\d*) \- User (?<username>\S*)"
```

See this example on regex101.com: <https://regex101.com/r/BrqKls/1>

Regex command:

<https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Regex>

Rer command:

<https://docs.splunk.com/Documentation/Splunk/latest/SearchReference/Rer>

Schema-on-read and Field Extraction



Schema-on-read vs Schema-on-write

Splunk has many options for storing and searching data

Schema-on-read / search-time extractions:

- We only write `_raw` and metadata to disk.
- Extract new fields when you run a search.
- This makes ingest very fast, search very flexible.
- Good for asking questions you didn't know in advance (i.e., during security investigations!!!)

Schema-on-write / index-time extractions:

- We write additional field information to disk at ingest-time (e.g., hash, username, source/dest IPs).
- This adds time to ingest and uses extra disk space, but enables fast, optimized searches.
- Good for asking questions you know you need to ask (i.e., for security monitoring).

Data model acceleration:

- We can write additional field information after the fact.
- Acceleration searches run in background on regular basis.

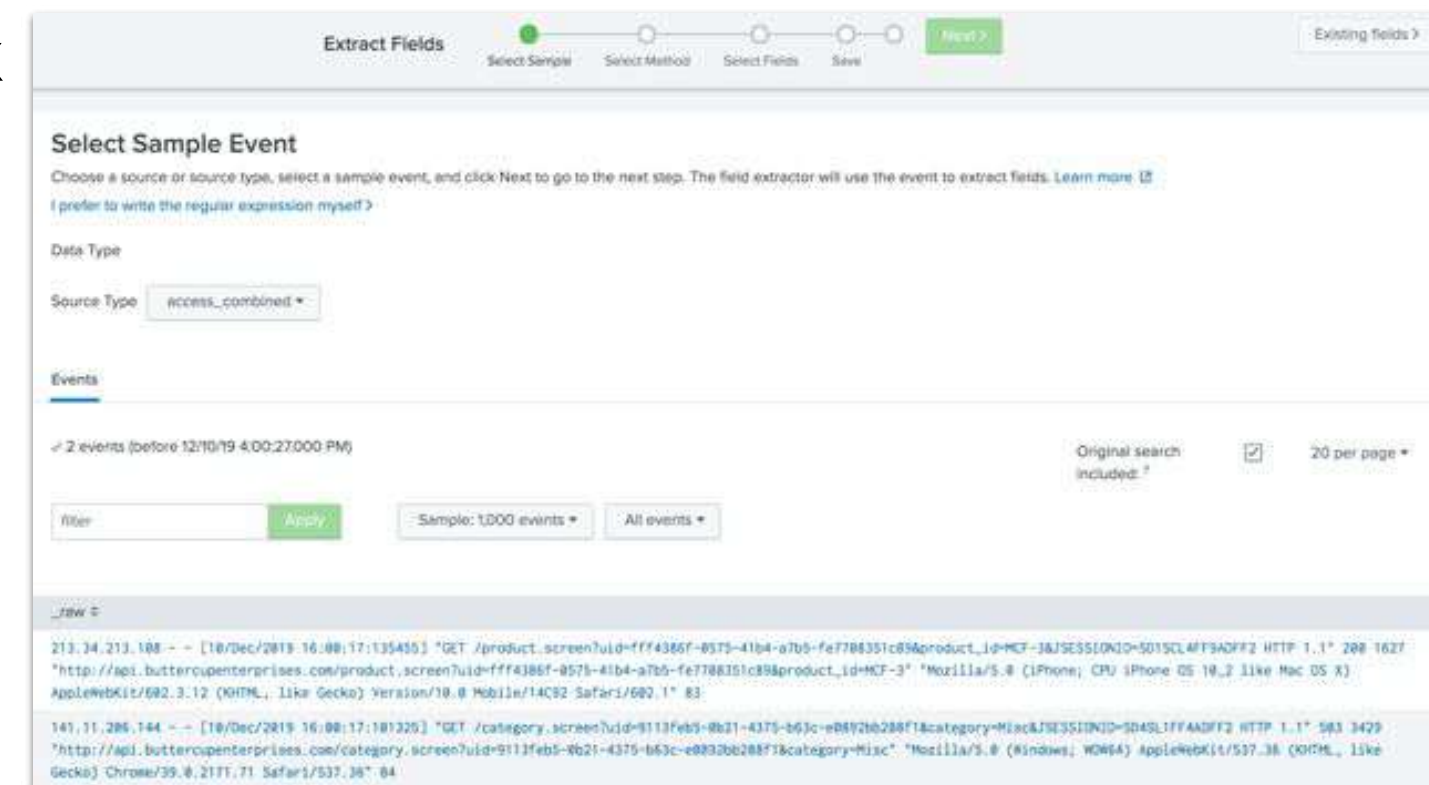
Interactive Field Extractor

Extract fields with UI-based Interactive Field Extractor - it creates regex for you!

Best practice: Extract ALL fields you need using single regex

Bad practice: A customer who used 72 different field extractions for 4 different types of data grouped into one sourcetype.

- **How to fix:** Separate into 4 different sourcetypes, each with 1-3 field extractions covering the ~18 fields on each sourcetype.



Work with Splunk admin teams on field extractions, but the Interactive Field Extractor is good for a quick one-off

Eventtypes, tags and macros

We can “hide” field extractions under eventtypes and tags. Eventtype=myfavoritedata => (sourcetype=X type=foo) OR (sourcetype=Y style=bar)

E.g. tag=network as part of Common Information Model hides LOTS of contributing information to make it easy for you to search for network data

Macro: store reusable search strings. E.g., wrap string “lookup assets.csv host” as macro ``asset_enrich``:

```
sourcetype=X | `asset_enrich`  
=> sourcetype=X | lookup assets.csv host
```

Save lots of time if you’re reusing search strings. You can pass variables (``mymacro(X)`` => something)

Best practice: Work with your Splunk admin team

Common Information Model

Use field aliases and tags to harmonize data across different types

E.g. field name “src” to appear on all network data, and sometimes represents “source”, “source_ip”, “starting_ip”, etc. This way you can search on src=1.2.3.4 across all data.

Talk to your Splunk admin team. If working with common data types, they’re probably already aligned to CIM. Security teams should definitely align to CIM.

Learn more about the Common Information Model:
<https://docs.splunk.com/Documentation/CIM/latest/User/Overview>

Validate CIM compliance:
https://github.com/hire-vladimir/SA-cim_vladiator

Data model	File name
Alerts	Alerts.json
Application State	Application_State.json
Authentication	Authentication.json
Certificates	Certificates.json
Change	Change.json
Change Analysis	Change_Analysis.json
CIM Validation (S.o.S)	Splunk_CIM_Validation.json
Databases	Databases.json
Data Loss Prevention	DLP.json
Email	Email.json

Thank you

