Splunk4Ninjas - Machine Learning

Hands on Introduction to the Splunk Machine Learning Toolkit (MLTK)

splunk>

Forward-Looking Statements

This presentation may contain forward-looking statements regarding future events, plans or the expected financial performance of our company, including our expectations regarding our products, technology, strategy, customers, markets, acquisitions and investments. These statements reflect management's current expectations, estimates and assumptions based on the information currently available to us. These forward-looking statements are not guarantees of future performance and involve significant risks, uncertainties and other factors that may cause our actual results, performance or achievements to be materially different from results, performance or achievements expressed or implied by the forward-looking statements contained in this presentation.

For additional information about factors that could cause actual results to differ materially from those described in the forward-looking statements made in this presentation, please refer to our periodic reports and other filings with the SEC, including the risk factors identified in our most recent quarterly reports on Form 10-Q and annual reports on Form 10-K, copies of which may be obtained by visiting the Splunk Investor Relations website at www.investors.splunk.com or the SEC's website at www.sec.gov. The forward-looking statements made in this presentation are made as of the time and date of this presentation. If reviewed after the initial presentation, even if made available by us, on our website or otherwise, it may not contain current or accurate information. We disclaim any obligation to update or revise any forward-looking statement based on new information, future events or otherwise, except as required by applicable law.

In addition, any information about our roadmap outlines our general product direction and is subject to change at any time without notice. It is for informational purposes only and shall not be incorporated into any contract or other commitment. We undertake no obligation either to develop the features or functionalities described, in beta or in preview (used interchangeably), or to include any such feature or functionality in a future release.

splunk>

Agenda

- Welcome/Introduction
- Intro to Machine Learning at Splunk
- Demo of Machine Learning Toolkit (MLTK) with Q&A
- Intro to the Trackday Dataset
- Four Different Challenges (~30 mins each)

  Challenge 1 - Explore the track_day.csv Dataset

  Challenge 2 - Detect Numeric Outliers

  Challenge 3 - Supervised Learning: Predict Categorical Fields

  Challenge 4 - Unsupervised Learning: Clustering
- Wrap Up, Discussion and Feedback

splunk>

Disclaimer

What this session is not about and what it is about

- **NO** replacement for a PhD in machine learning, data science or AI
- **NO** replacement for Splunk's Education class for Data Science
- **NO** comprehensive lecture about all possible concepts and algorithms in ML … but,

- **YES** first introduction into Machine Learning @ Splunk
- **YES** getting to know of Splunk's Machine Learning Toolkit
- **YES** guided hands-on challenges to explore a few typical ML tasks

Machine Learning Tour

splunk>

Splunk Customers Want Answers from their Data

Anomaly detection

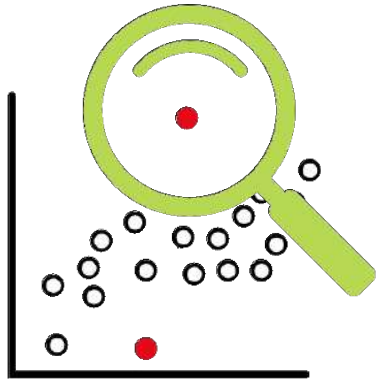

- Deviation from past behavior
- Deviation from peers
- (aka Multivariate AD or Cohesive AD)
- Unusual change in features
- **ITSI MAD**

Predictive Analytics

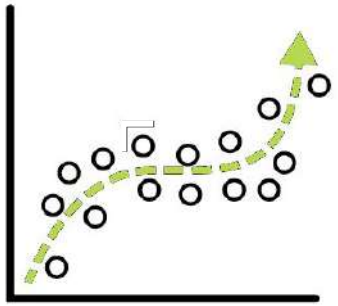

- Predict Service Health Score/Churn
- Predicting Events
- Trend Forecasting
- Detecting influencing entities
- Early warning of failure

Clustering

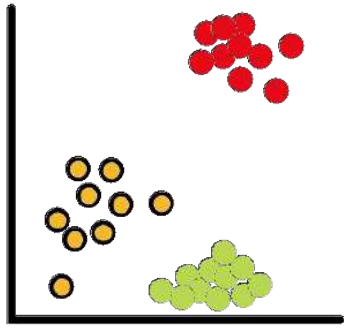

- Identify peer groups
- Event Correlation
- Reduce alert noise
- Behavioral Analytics
- **ITSI Event Analytics**

splunk>

Types of Machine Learning

Supervised Learning (labeled data)

- Regression
- Classification

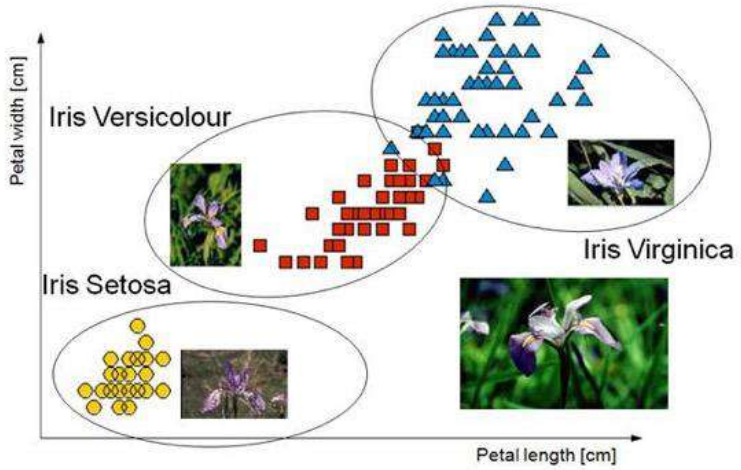

Unsupervised Learning (unlabeled data)

- Clustering
- Anomaly Detection

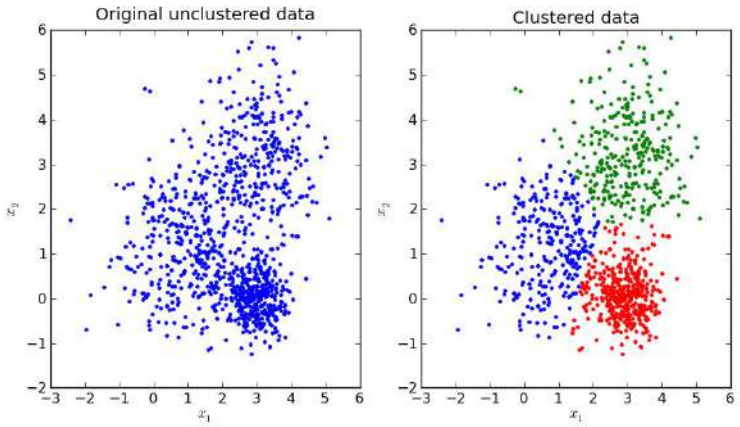

Mixed Models (with reinforcement or feedback)

- Human in the Loop
- Autonomous Systems

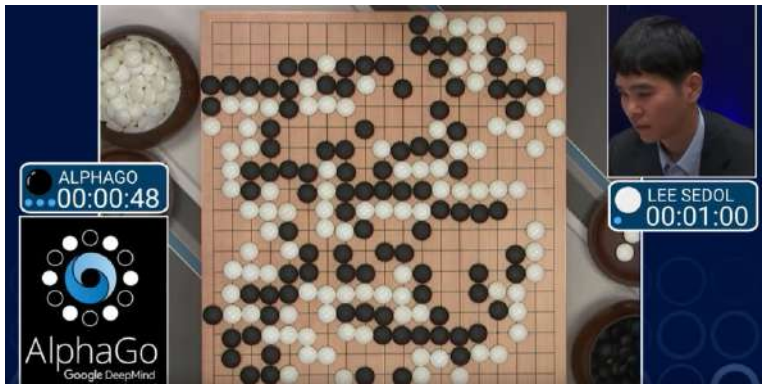

splunk>

Overview of Machine Learning at Splunk

Skill Areas for Machine Learning @ Splunk

Premium solutions provide out of the box ML capabilities.

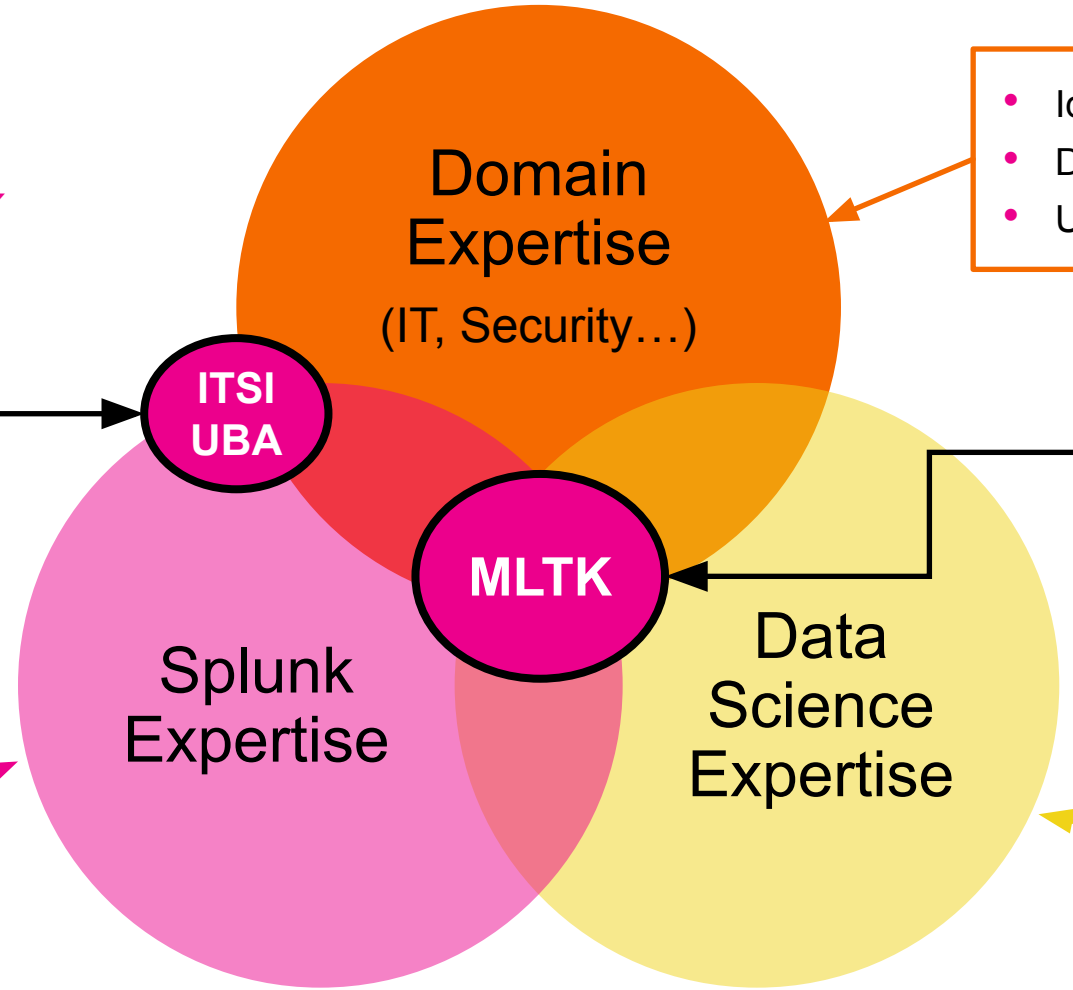

- Identify use cases
- Drive decisions
- Understanding of business

Splunk ML Toolkit facilitates and simplifies via examples & guidance

- Searching
- Reporting
- Alerting
- Workflow

- Statistics/math background
- Algorithm selection
- Model building

splunk>

What Data Scientists Really Do

Data Preparation accounts for about 80% of the work of data scientists

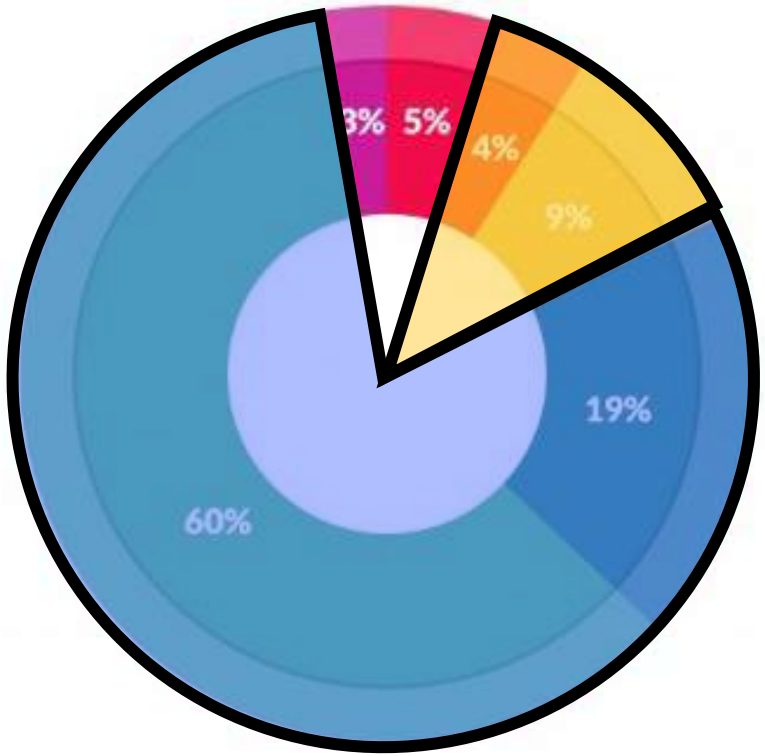

What data scientists spend the most time doing

- *Building training sets: 3%*
- *Cleaning and organizing data: 60%*
- *Collecting data sets; 19%*
- *Mining data for patterns: 9%*
- *Refining algorithms: 4%*
- *Other: 5%*

"Cleaning Big Data: Most Time-Consuming, Least Enjoyable Data Science Task, Survey Says", Forbes Mar 23, 2016

Custom ML with the Splunk Platform

Ecosystem — Splunk's App Ecosystem contains 1000's of free add-ons for getting data in, applying structure and visualizing your data giving you faster time to value.

MLTK — The Machine Learning Toolkit delivers new SPL commands, custom visualizations, assistants, and examples to explore a variety of ml concepts.

Splunk — Splunk Enterprise is the mission-critical platform for indexing, searching, analyzing, alerting and visualizing machine data.

Operationalized Data Science Pipeline

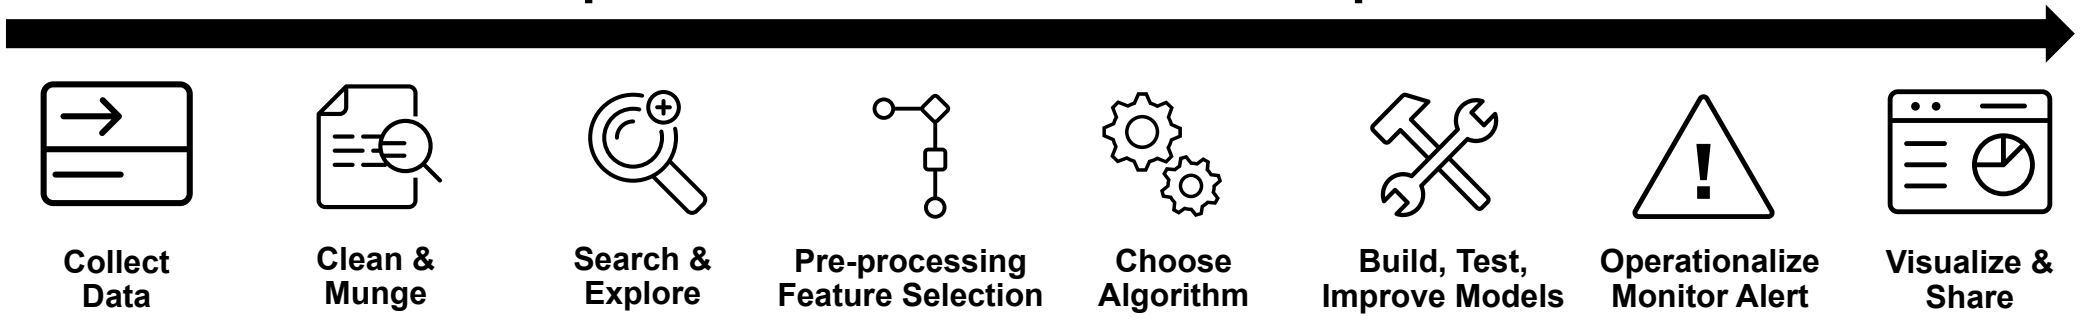

Collect Data	Clean & Munge	Search & Explore	Pre-processing Feature Selection	Choose Algorithm	Build, Test, Improve Models	Operationalize Monitor Alert	Visualize & Share
Ecosystem	Ecosystem	Splunk	MLTK	MLTK	MLTK	MLTK	Ecosystem
Splunk	Splunk		Splunk	Splunk	Splunk	Splunk	Splunk

splunk> Platform for Operational Intelligence

Continuous Data Ingest at Scale

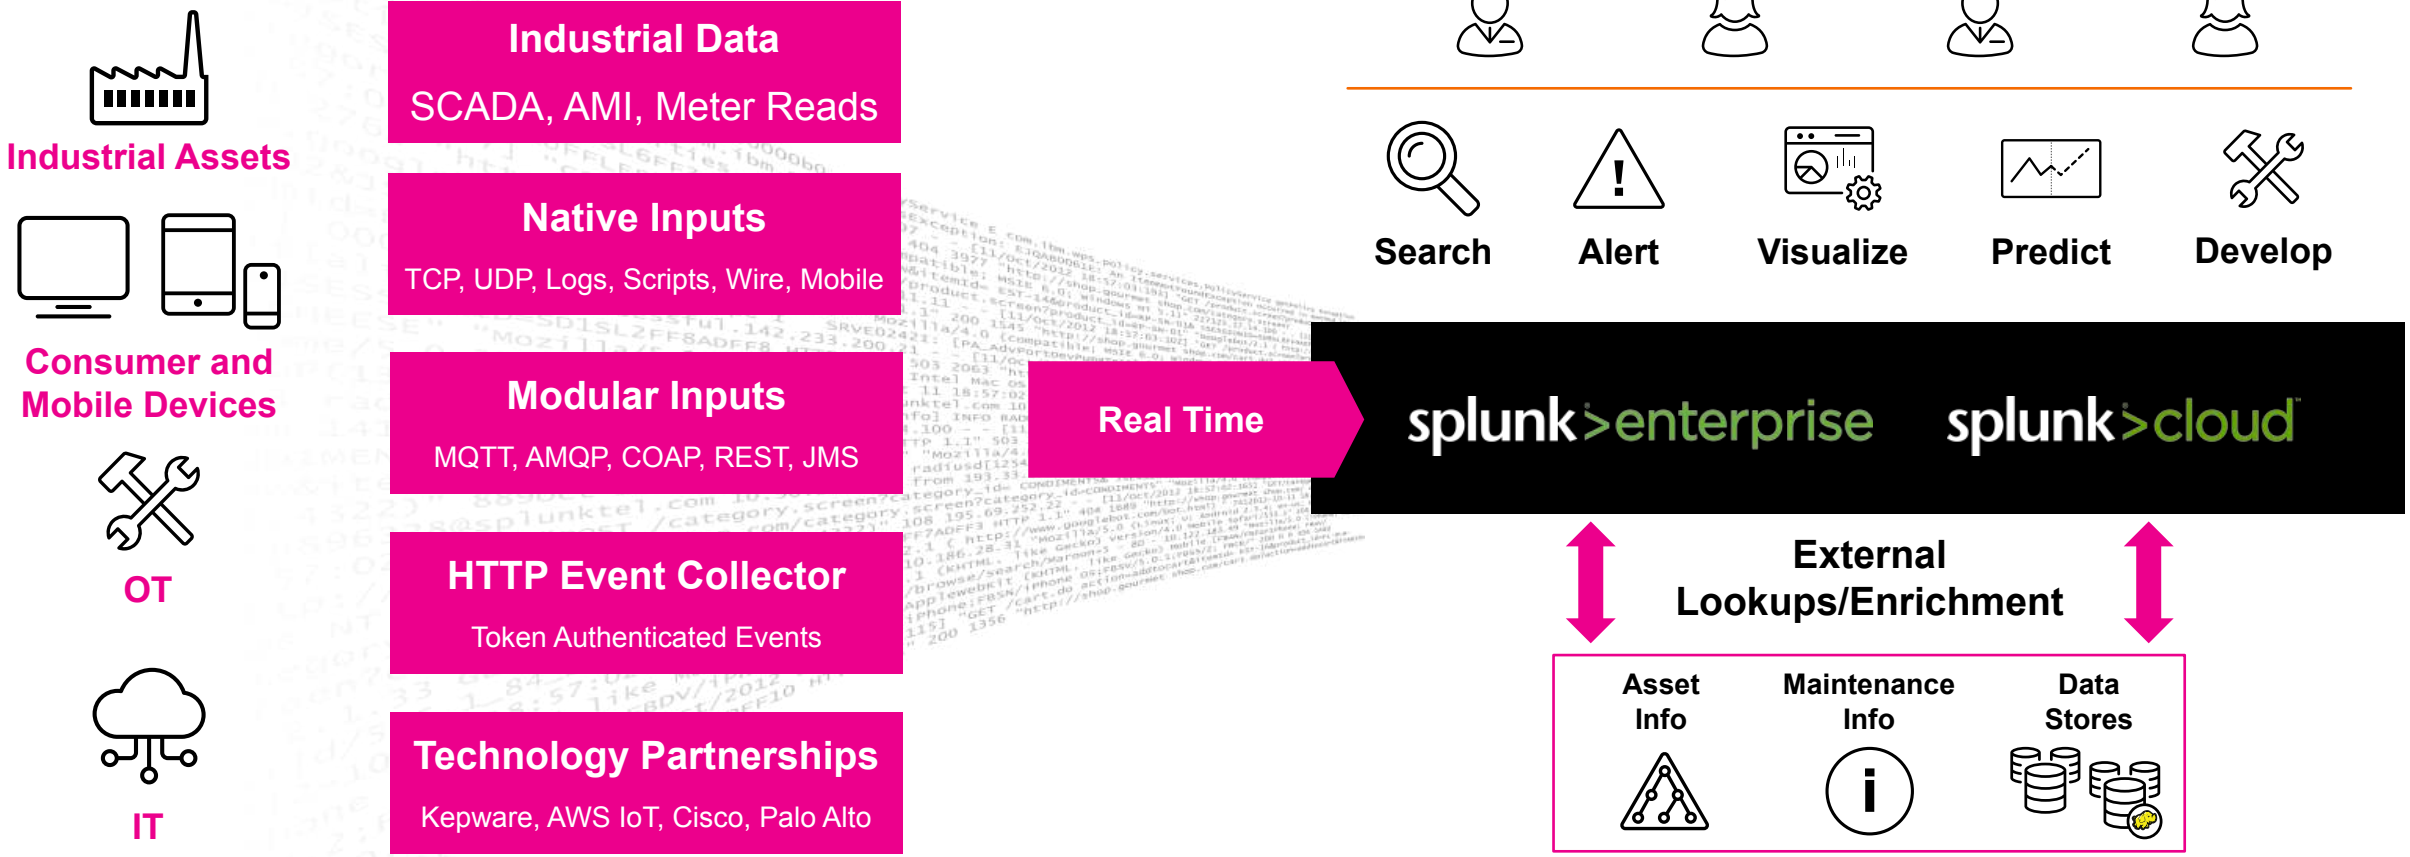

splunk>

Sense and Respond

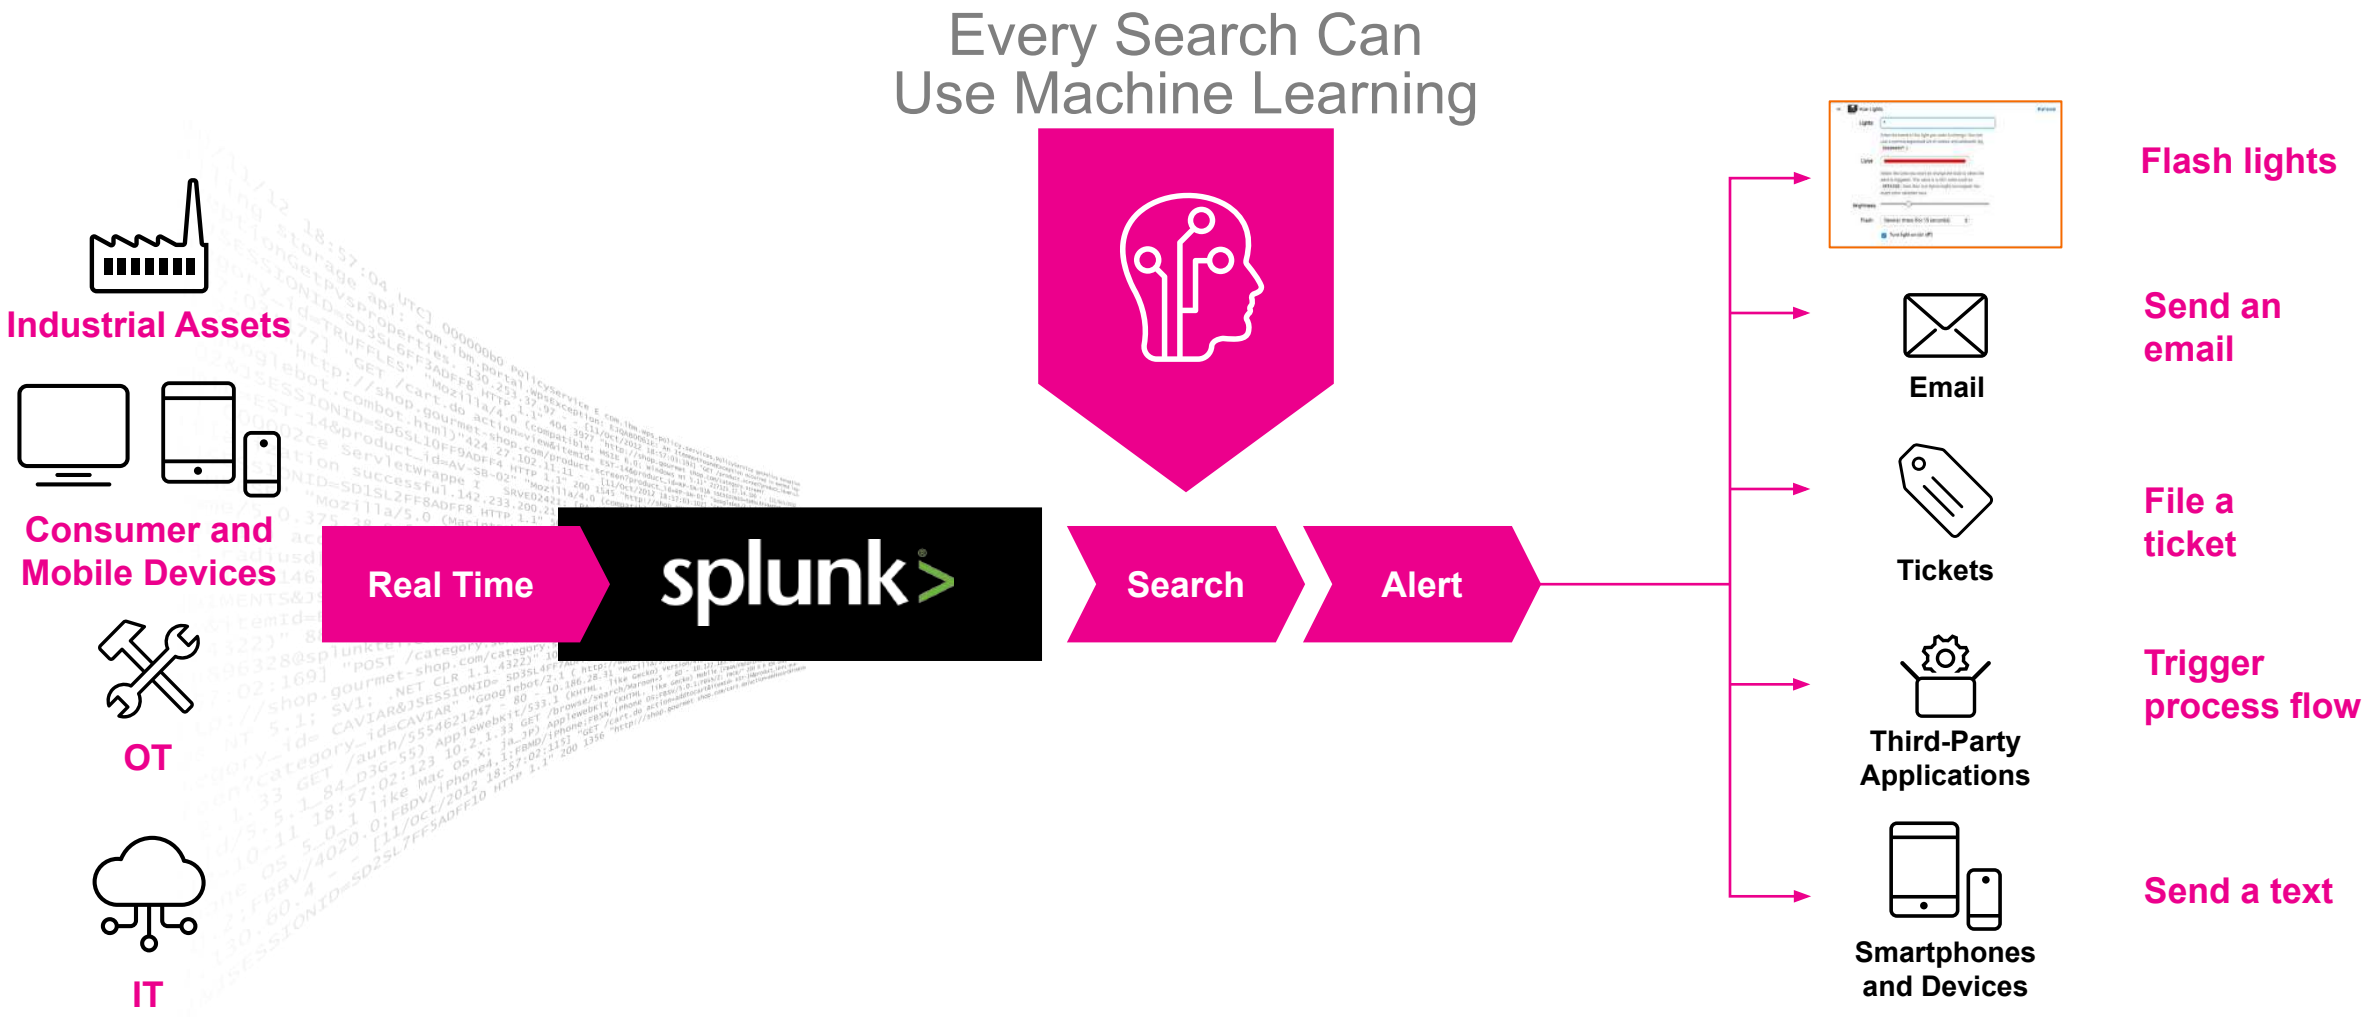

splunk>

MLTK + Python for Scientific Computing

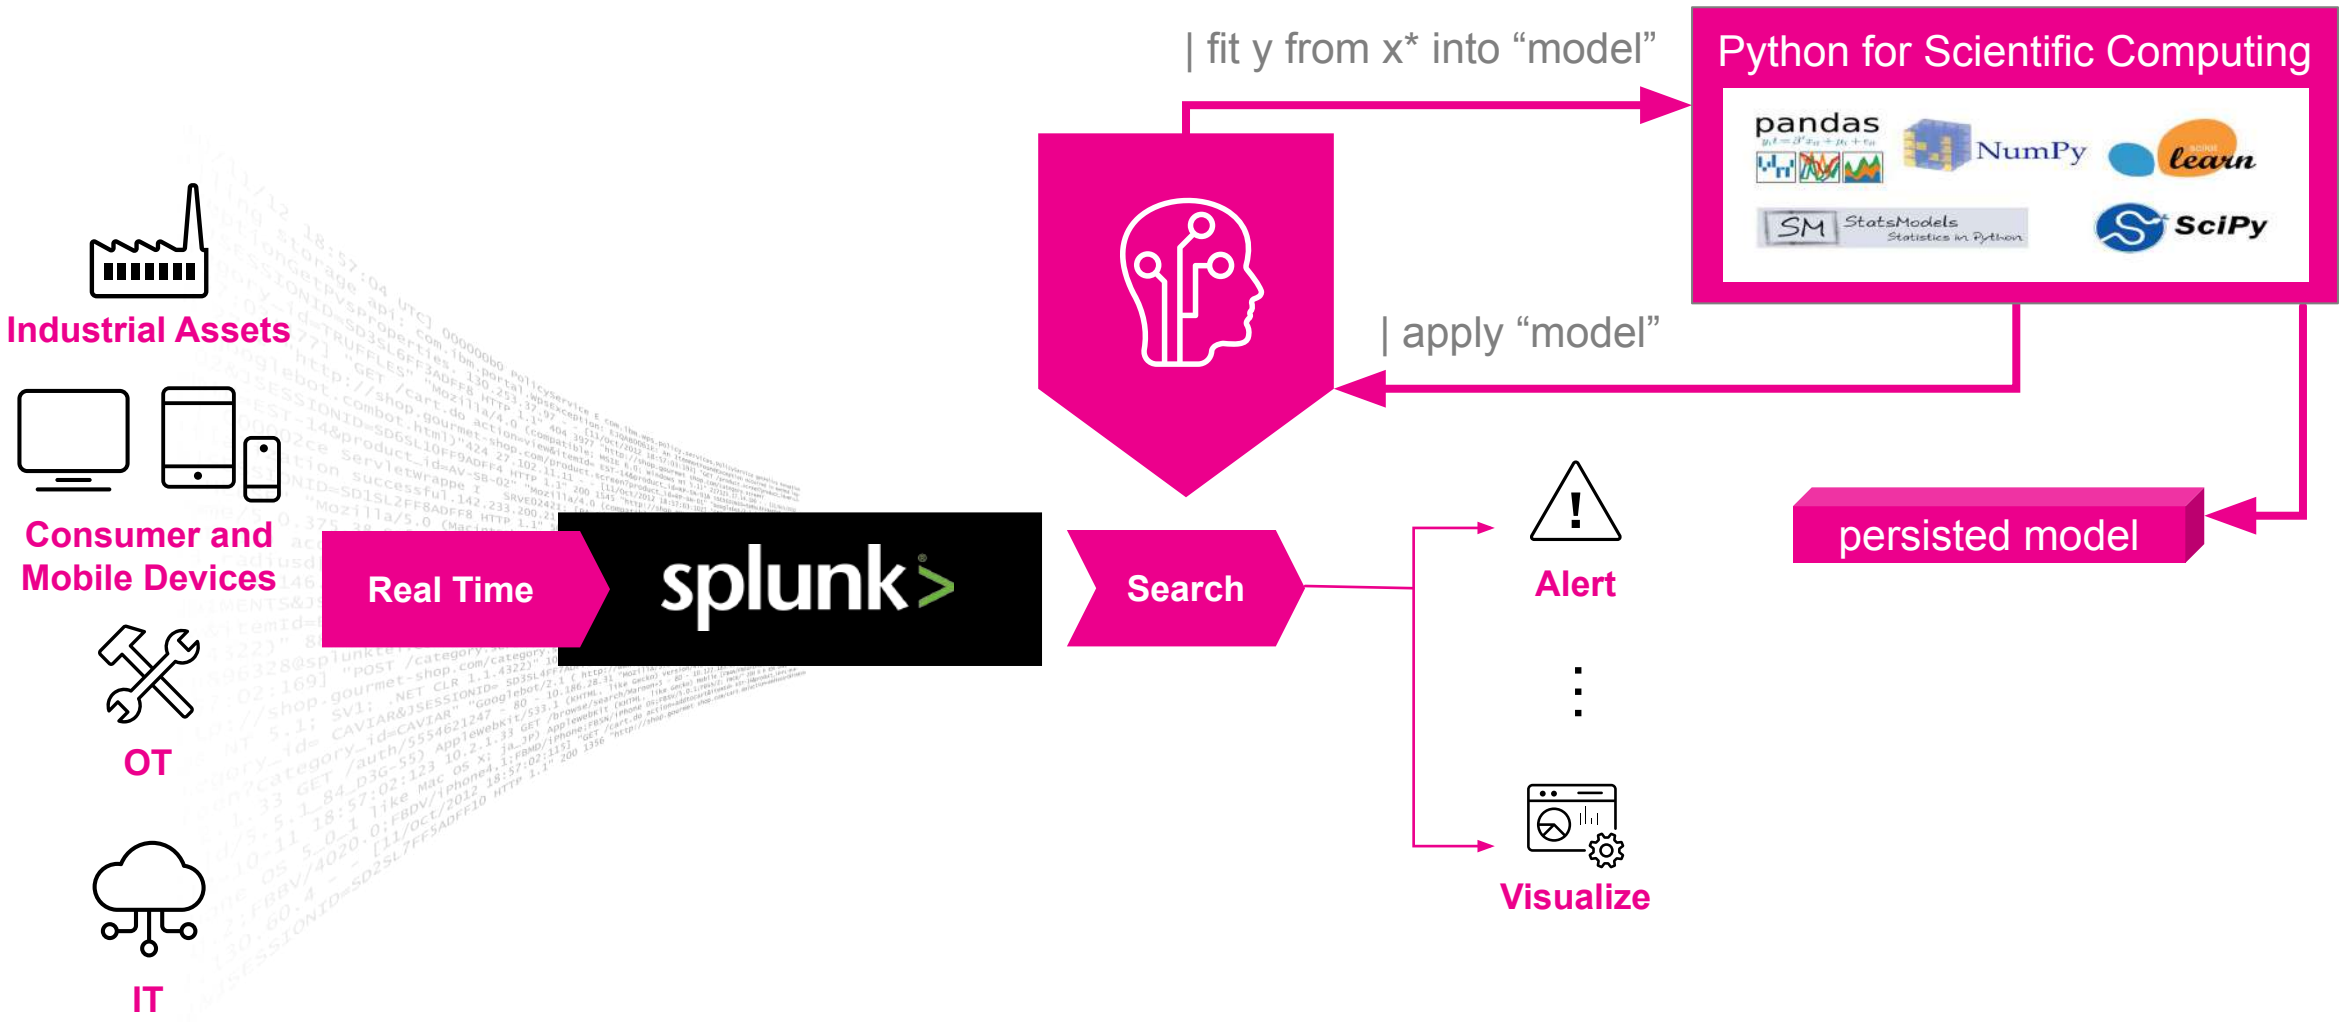

Splunk Machine Learning Toolkit (MLTK)

Extends Splunk platform functions and provides a guided modeling environment

Built for the Citizen Data Scientist

- **Experiments and Assistants**: Guided model building, testing, and deployment for common objectives
- **Algorithms**: 80+ standard algorithms (supervised & unsupervised)

Extensible to operationalize any use case

- **Python for Scientific Computing Library**: Access to 300+ open source algorithms
- **Deep Learning Toolkit** : Supports NN and GPU accelerated machine learning
- **ML-SPL API**: Import any open-source or proprietary algorithm

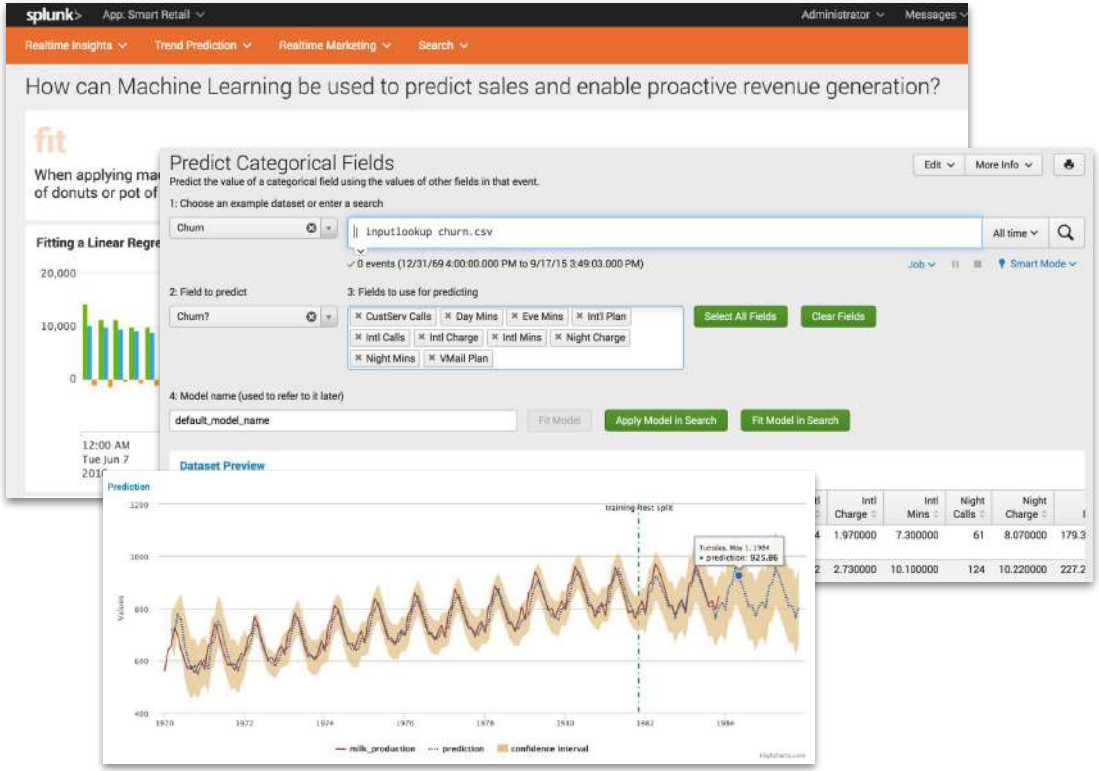

splunk>

Example: MLTK powered DGA App for Splunk

Detect Malicious Domain Names using Machine Learning

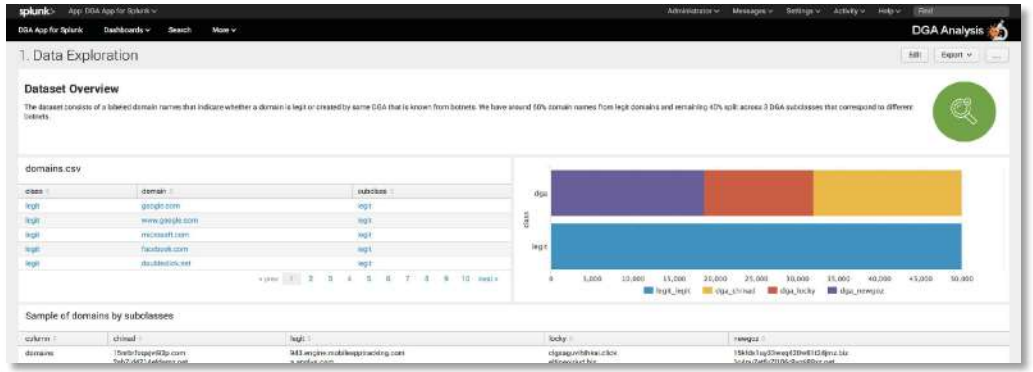

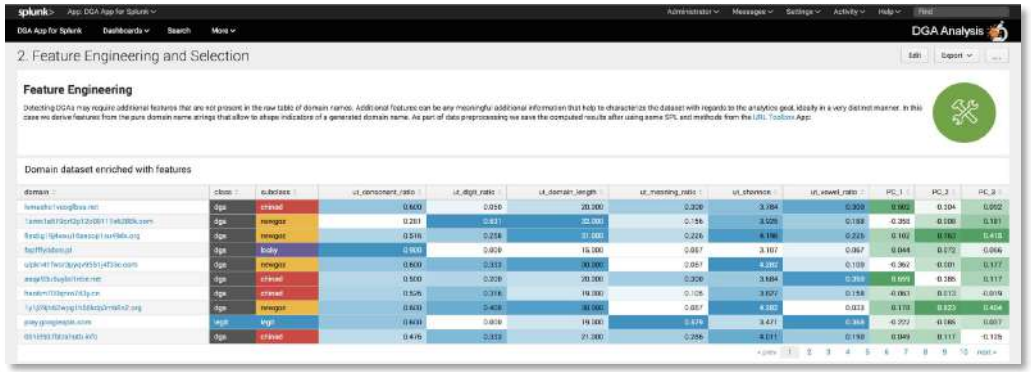

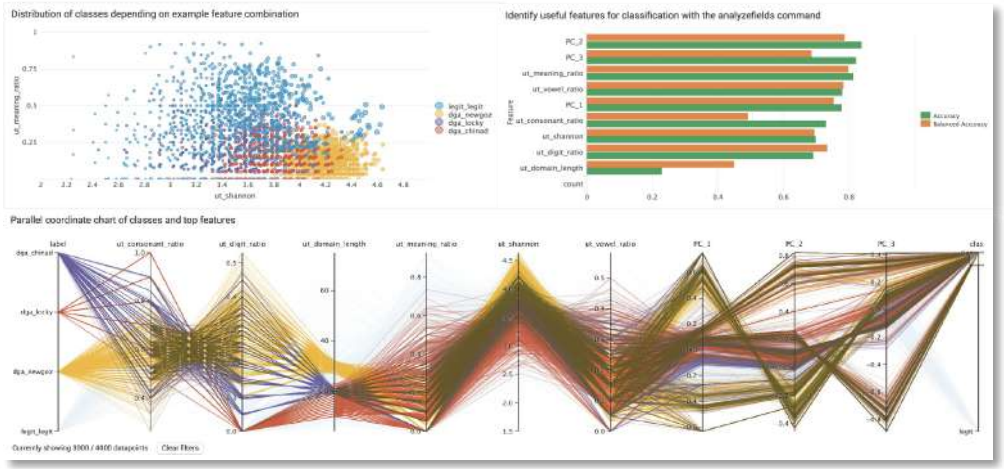

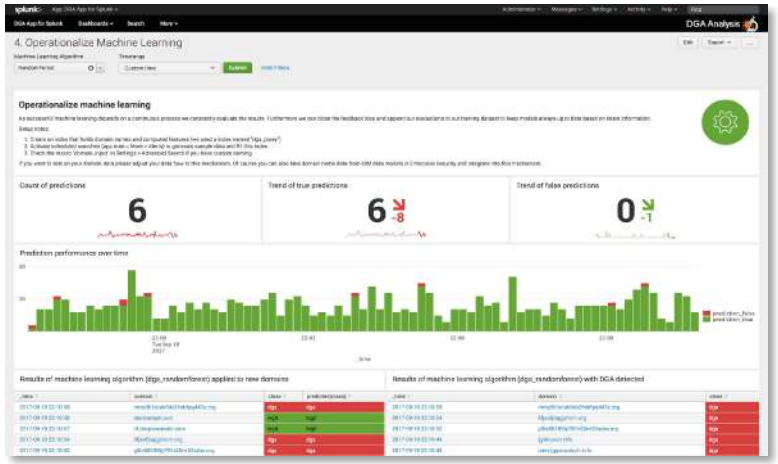

splunk>

Overview of ML including DL at Splunk

(not covered in this workshop)

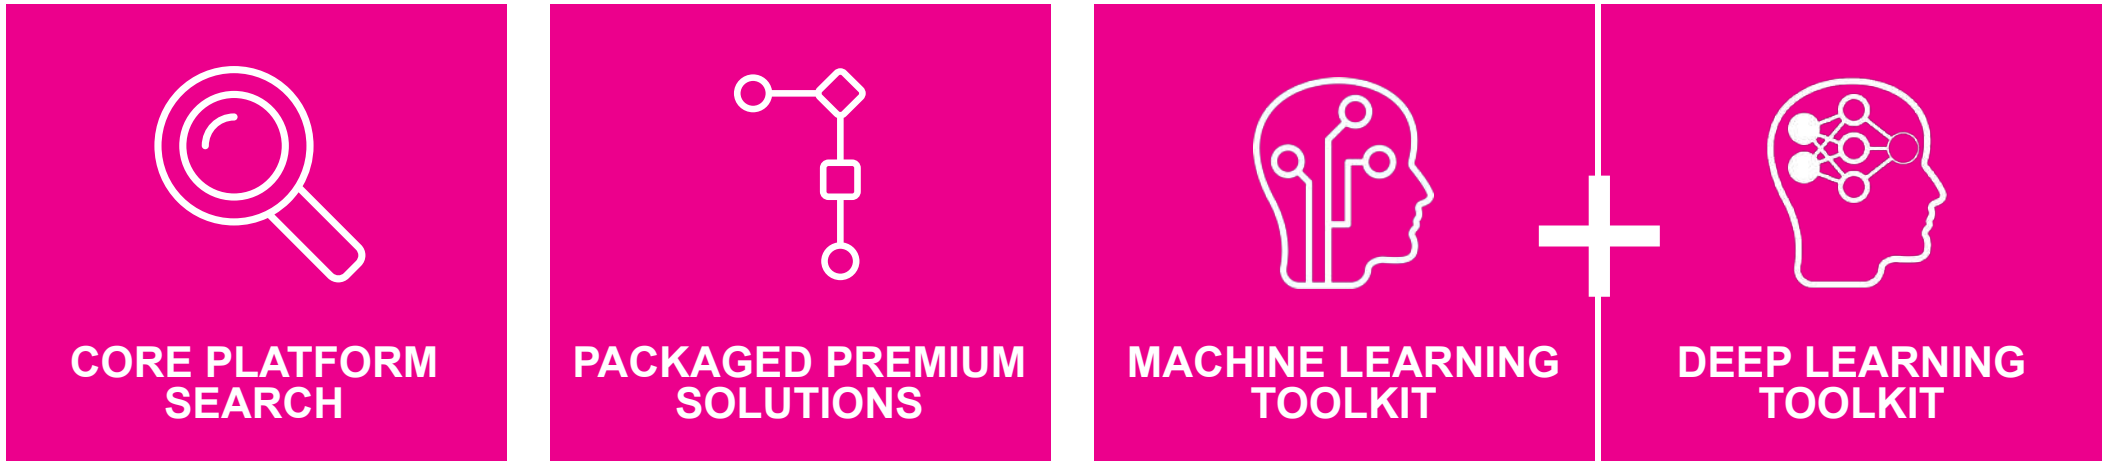

DLTK for Splunk

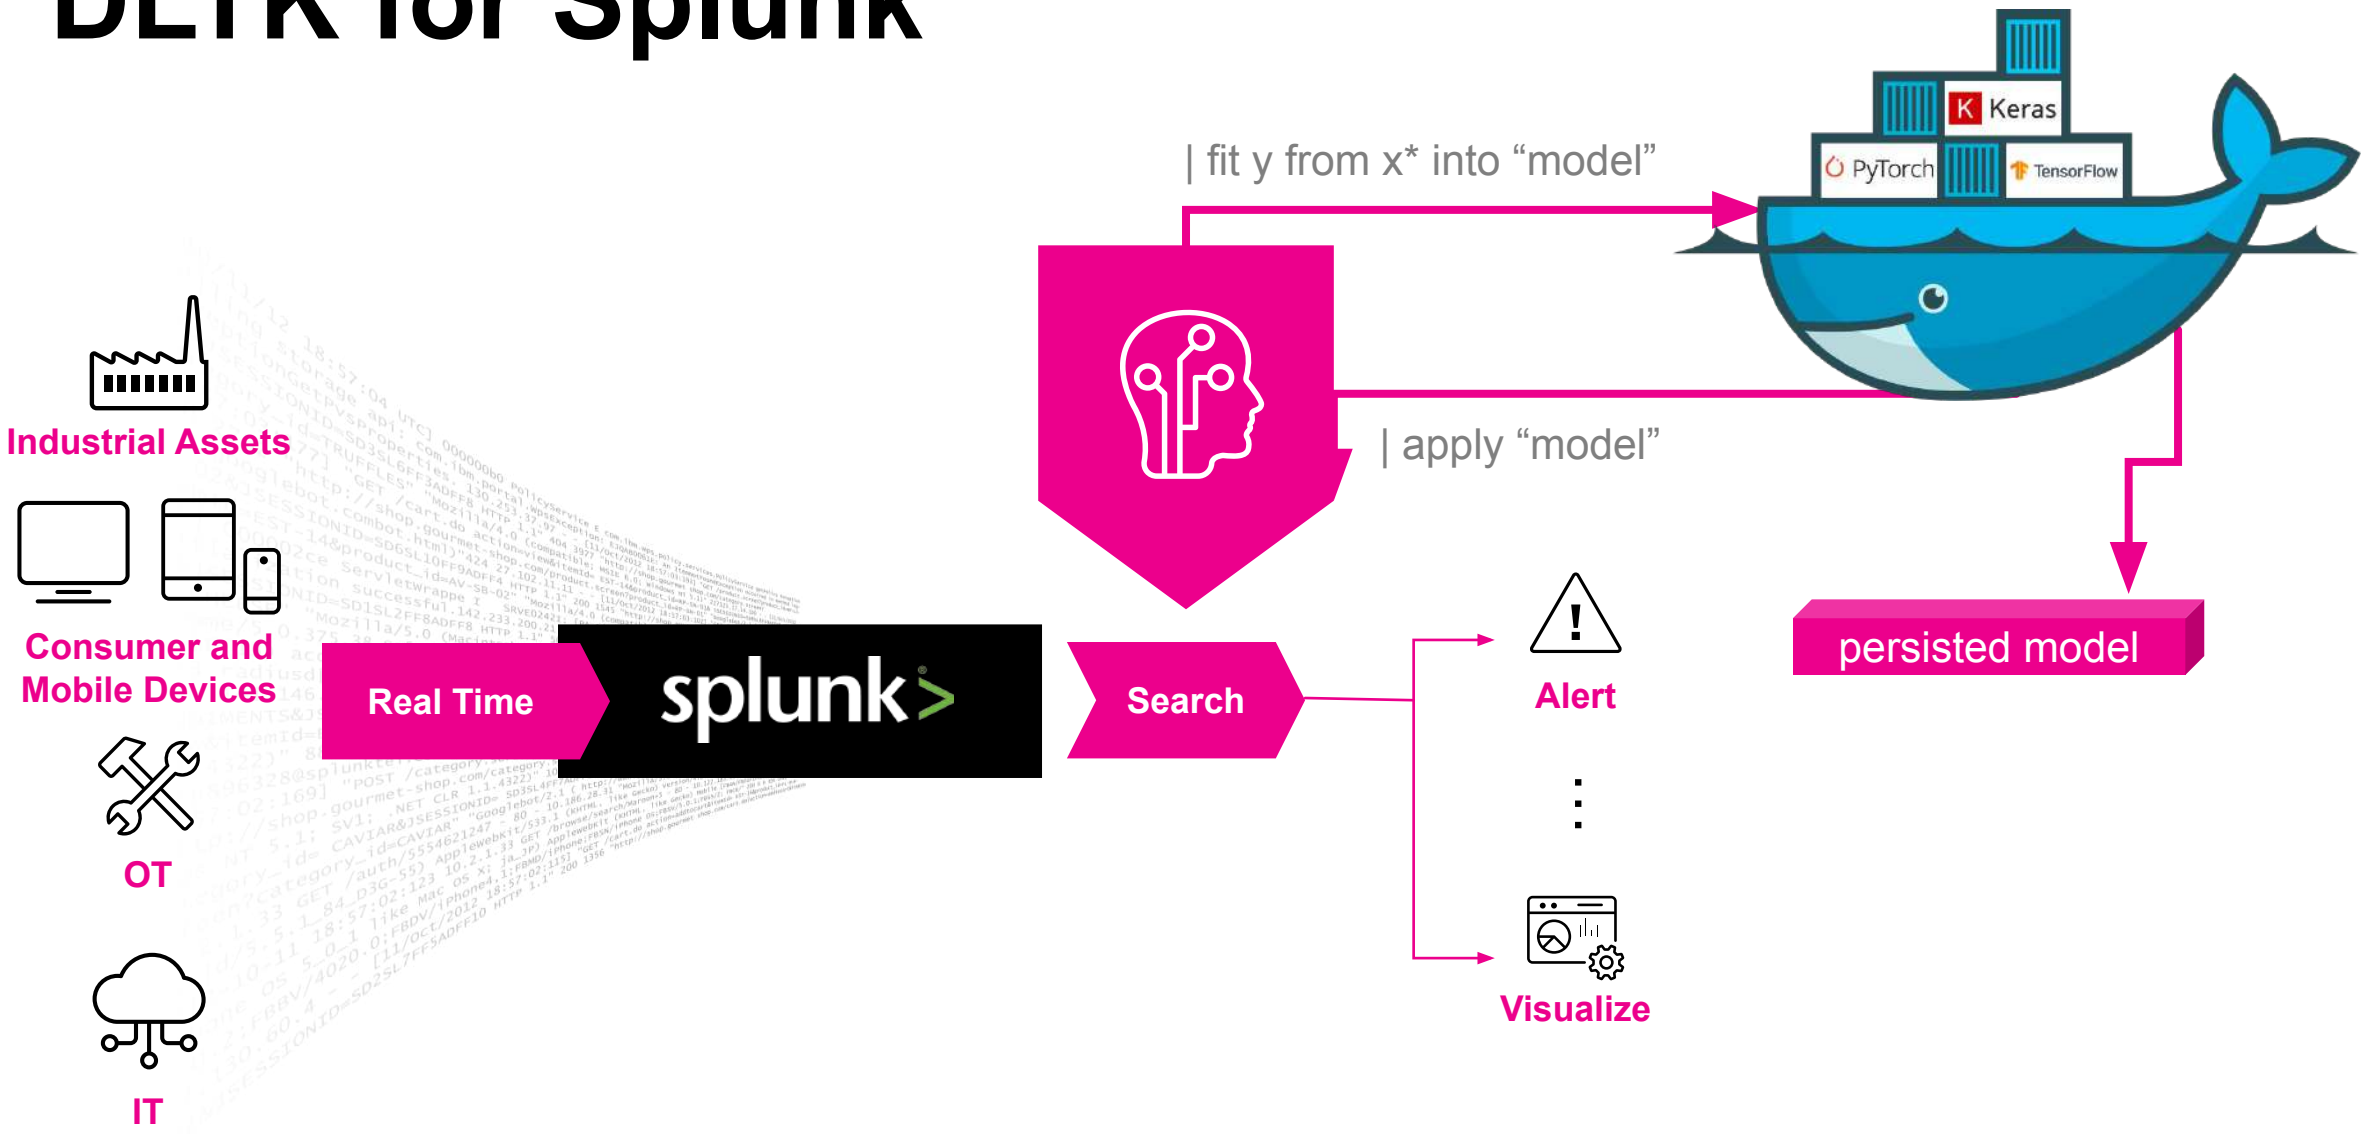

Hyatt ensured every customer experience with hotel wifi across the globe

1. All the data showing every customer sign in to the Hotel Wifi programs along with other relevant data (provider and lookup file of provider info, hotel ID and lookup file of hotel info, indexes for real time data of customers, local and global holidays).
2. A workflow (dashboard(s)) to show customers at every hotel signing in, normalize _time so a hotel in one time zone at 8 am to can be compared to another hotel at 8am.
3. Forecasting the likely wifi logins based on each property, day of week, and local and global holiday out two days into the future to show our expectations to executives.
4. Detecting meaningful anomalies as real time data comes in and is compared to the forecast . I must be able to insert business rules into the anomalies based on analyst feedback in a quick and nimble way. I should update my learning every night.

splunk>

Demo: Machine Learning Toolkit

splunk>

Hands-on Challenges

splunk>

REGISTER 5 MINS

Enroll in Today's Workshop

Tasks

1. Get a splunk.com account if you don't have one yet: https://splk.it/SignUp
2. Enroll in the Splunk Show workshop event: https://show.splunk.com/event/<eventID>
3. Download a copy of today's slide deck: https://splk.it/S4N-ML-Attendee

Goal

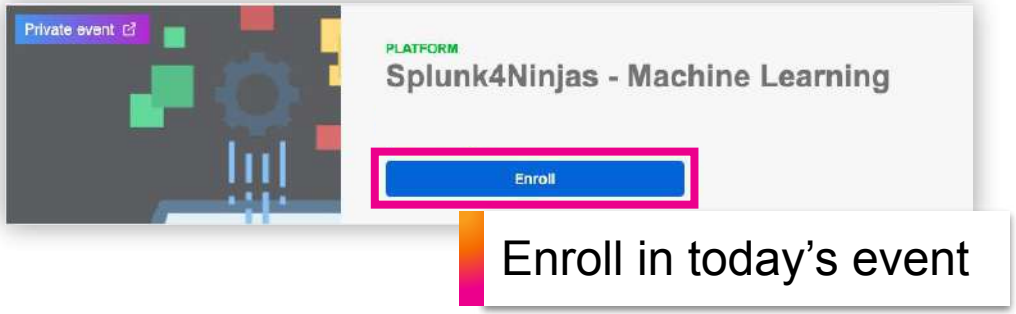

splunk>

Fun Facts about the Track Day dataset

A popular private event of racing and sportscar affine Splunkers in the early days.

Simple concept

Go on a race track, have fun and collect some car data to get insights about driving behavior etc.

A subset of this data is available in MLTK!

Image Source: https://www.youtube.com/watch?v=meBjl-ay9-U

Today's Challenges

We are going to create four dashboards:

1. **Explore the Dataset:** Create a sample dataset and explore it using different types of visualizations such as SPL

2. **Detect Numeric Outliers:** Explore the MLTK showcase and adapt it to start a new experiment with your own dataset

3. **Use a Classification Model:** Create a classification model and use it to predict vehicle types from your sensor data

4. **Use a Clustering Model:** Create a clustering model and and use it to analyze your dataset

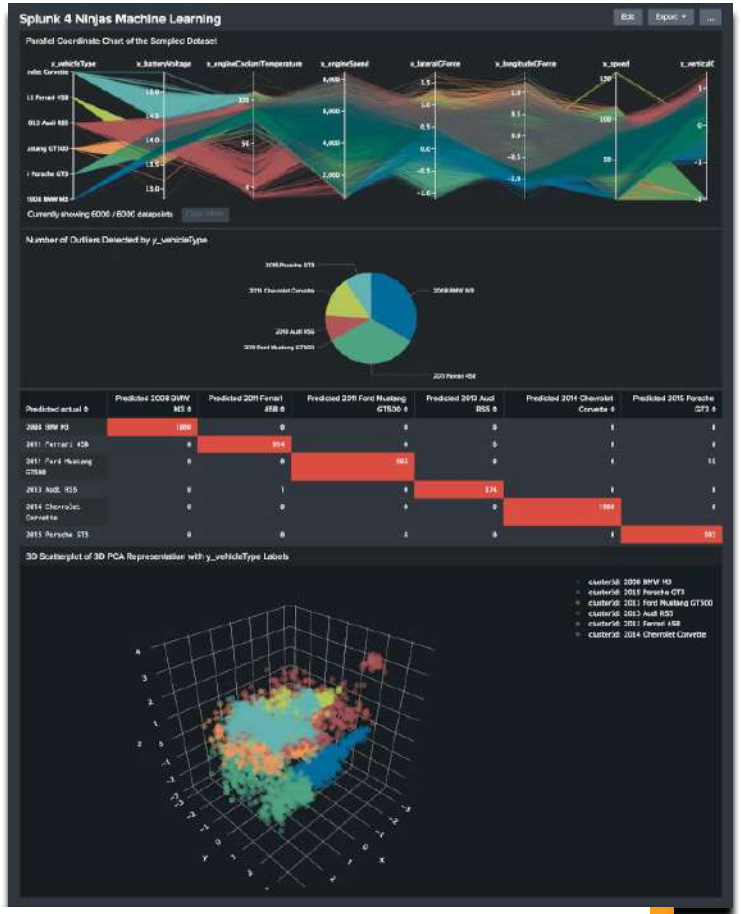

We're aiming for a dashboard like this!

splunk>

Workshop Goals

- Getting to know Splunk in the context of Machine Learning
- Prepare and analyze a dataset and summarize results on 4 dashboards


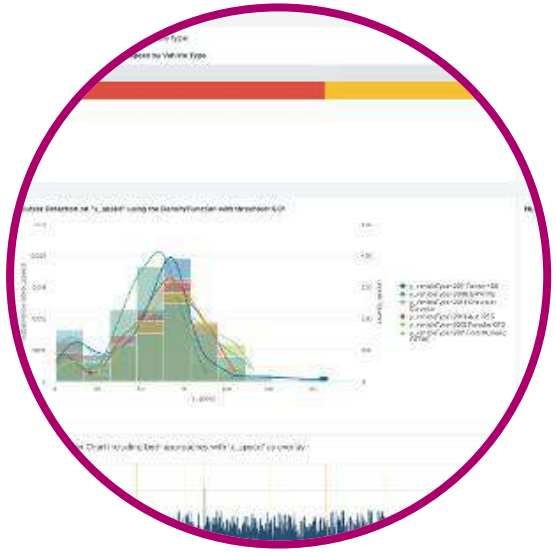

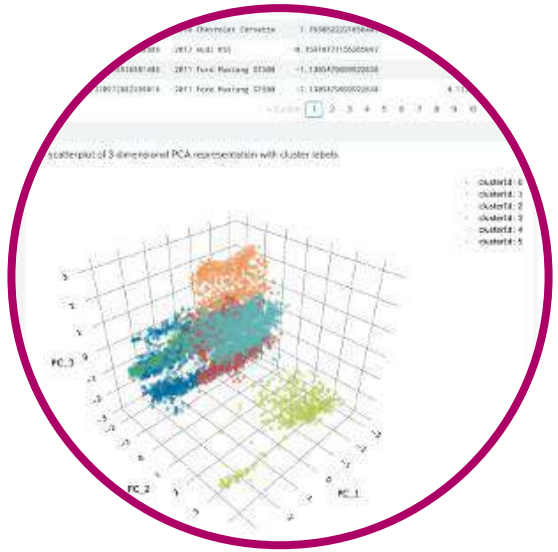

splunk>

Login to Splunk

Locate your instance URL and credentials in the Splunk Show event https://show.splunk.com

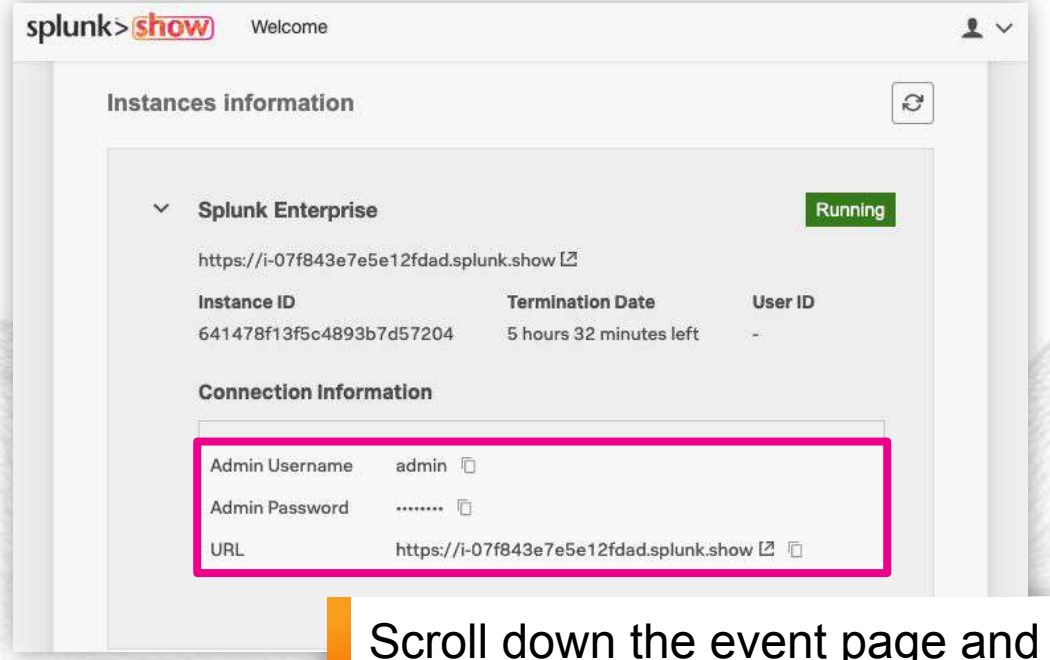

Scroll down the event page and expand the **Splunk Enterprise** section to view your login details

Log in to your Splunk instance

Login using the credentials from Splunk Show

splunk>

Create a Sample Dataset

1

1. Access the Splunk Machine Learning Toolkit

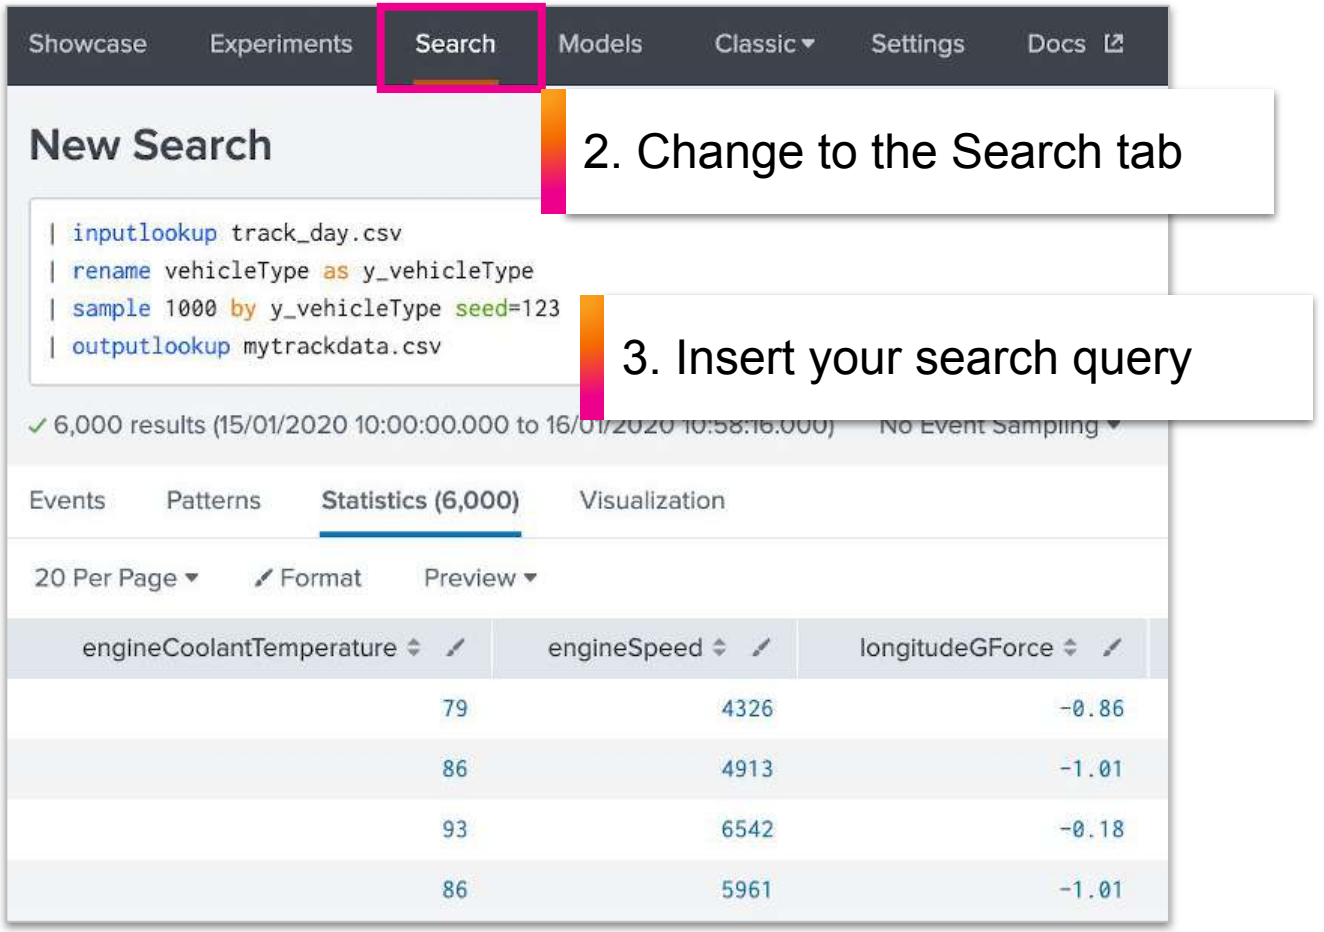

2. Change to the Search tab

3. Insert your search query

What's the benefit of renaming variables?

splunk>

1 Use Fieldsummary to Explore your Dataset

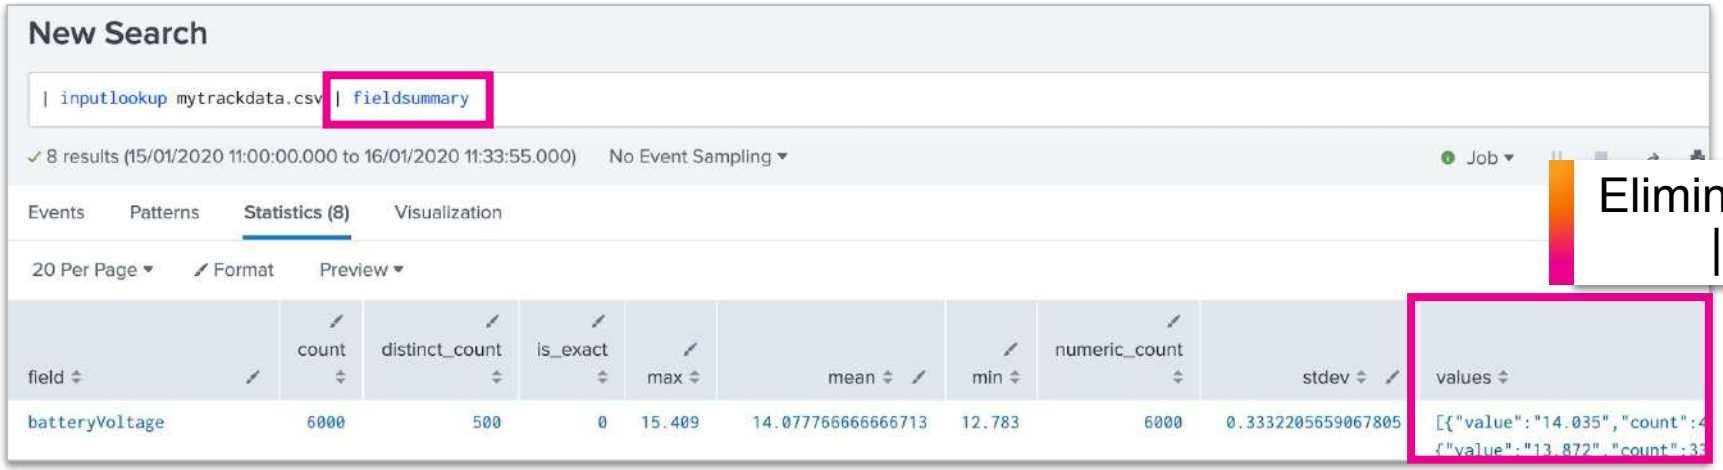

Eliminate unwanted Fields with
`| fields - values`

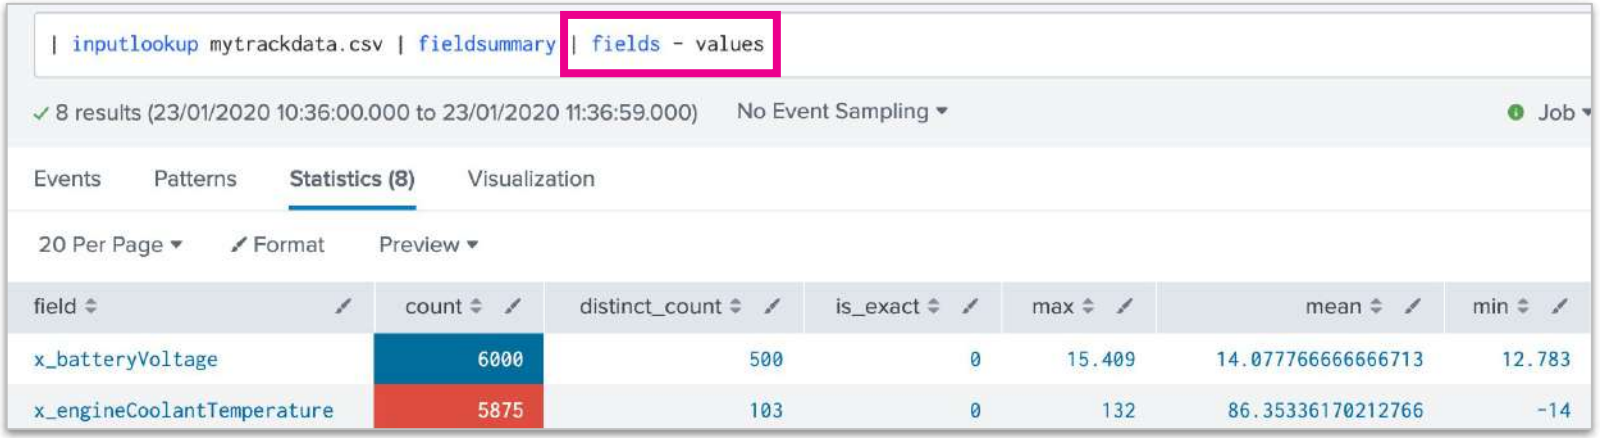

What's going on with the engine coolant temperature?

splunk>

1 Explore your Dataset with Visualizations

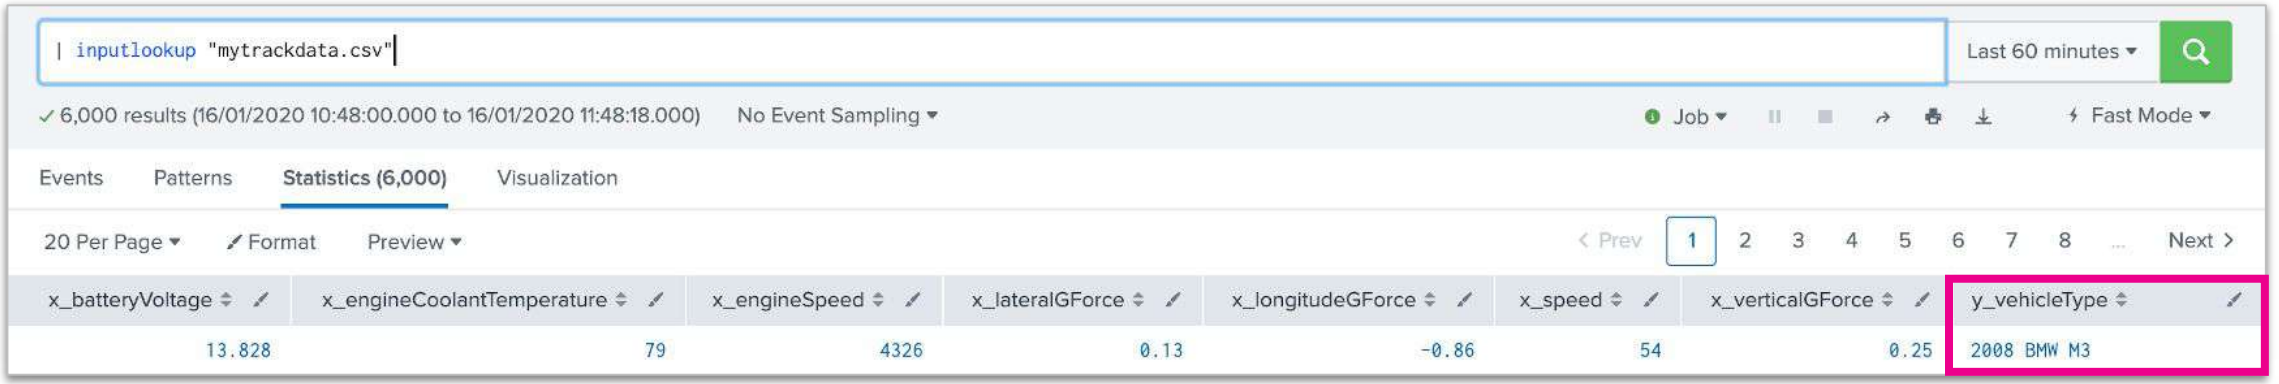

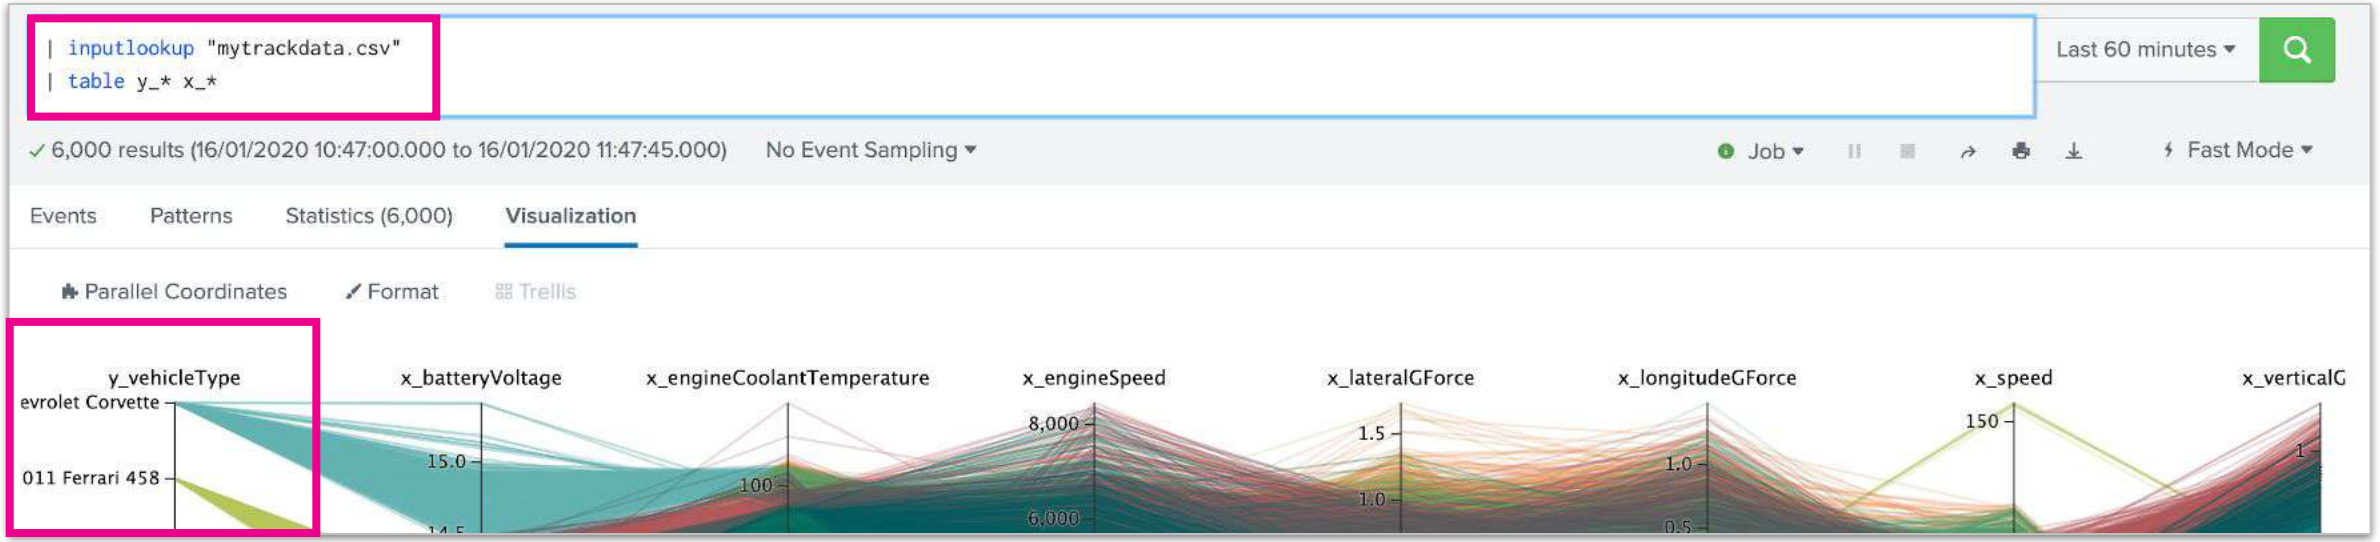

splunk>

1 Using Splunk's Histogram Macro

Check Macro in settings

OR

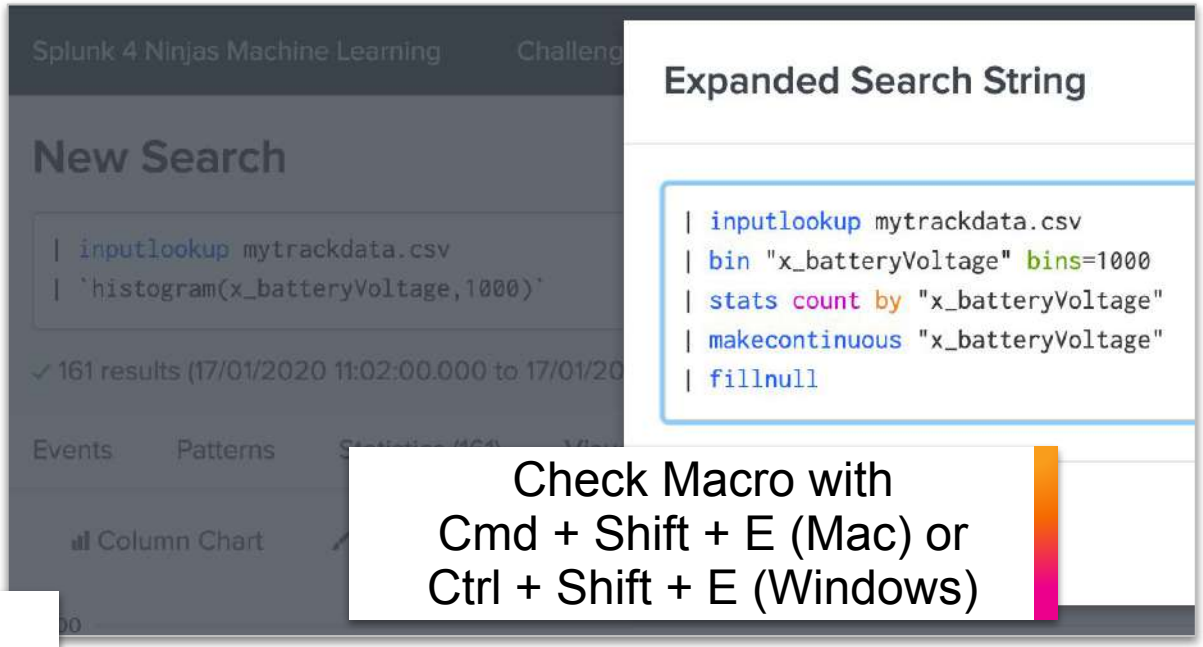

Check Macro with
Cmd + Shift + E (Mac) or
Ctrl + Shift + E (Windows)

splunk>

1 Adjusting the Histogram Macro

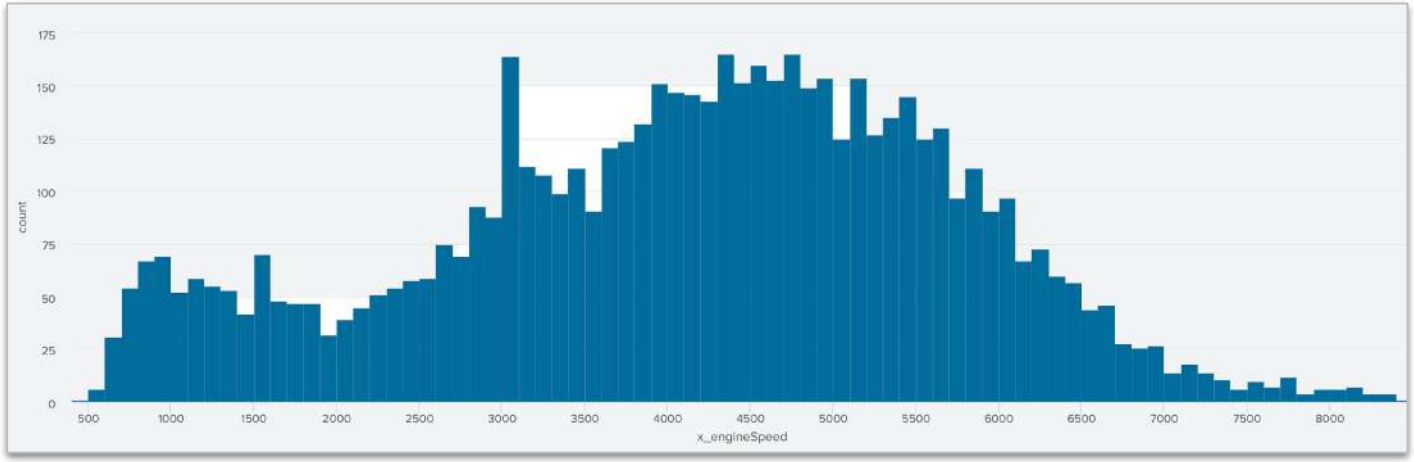

How can we get from the top to the bottom histogram?

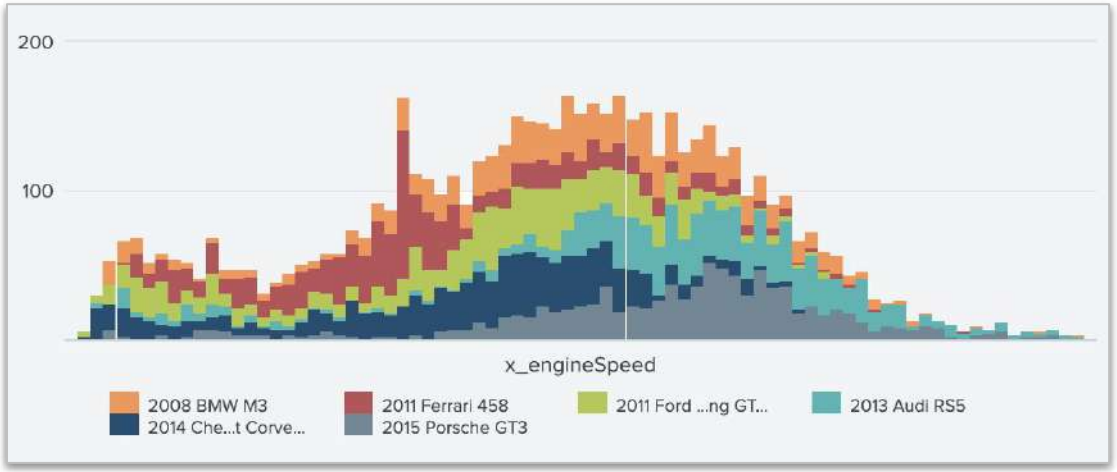

splunk>

1 Adjust the Macro to Split by Vehicle Type

```
| chart count by y_vehicleType
```

vehicleType	count
Ferrari	641
Audi	42
BMW	51
Chevrolet	44
Ford	95

```
| chart count over x_engineSpeed by y_vehicleType
```

batteryVoltage	Ferrari	Audi	BMW	Chevrolet	Ford
13	0	0	0	0	1
14	0	0	0	1	1
15	1	0	1	0	0
16	1	1	1	0	0
17	1	0	0	0	1

1

Working with the Boxplot Macro

How can this query be improved?

```
| inputlookup "mytrackdata.csv"
| `boxplot`
```

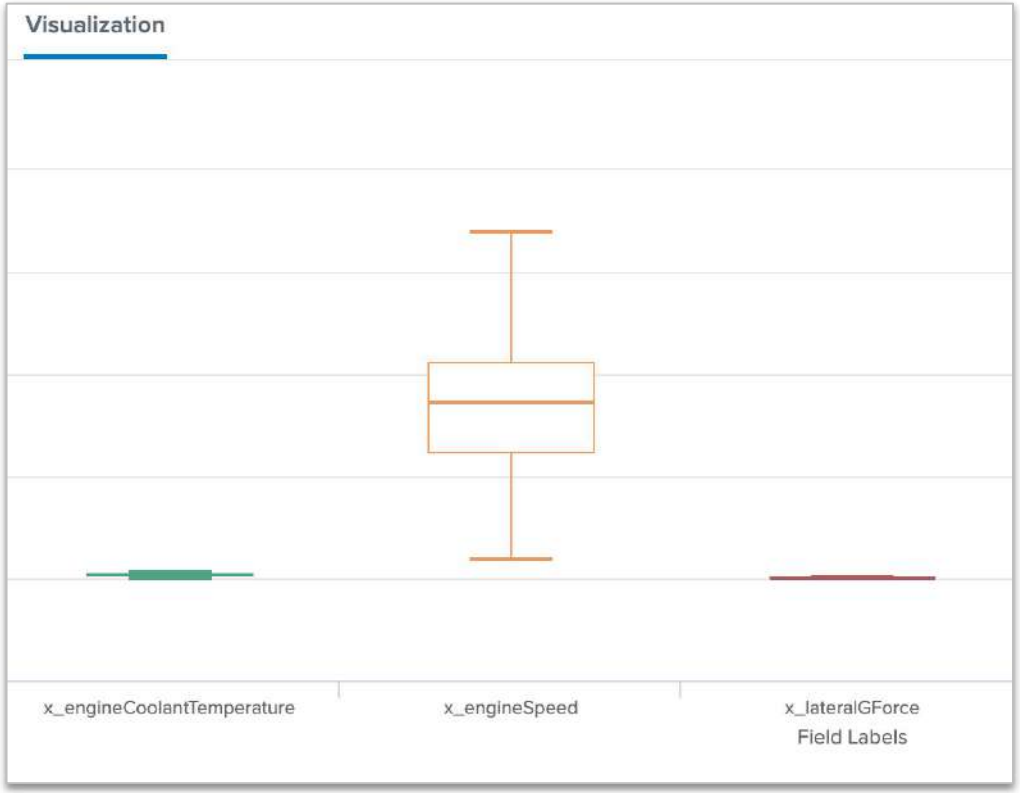

Hints:

Scale numeric values using the `fit` command with the StandardScaler

splunk>

1 Explore the Dataset with Box Plots

```
| inputlookup "mytrackdata.csv"
| fit StandardScaler x_*
| table SS_*
| `boxplot`
```

- Standardized data fields have a mean of 0 and a standard deviation of 1
- The box plots are less stretched and can be analyzed more easily

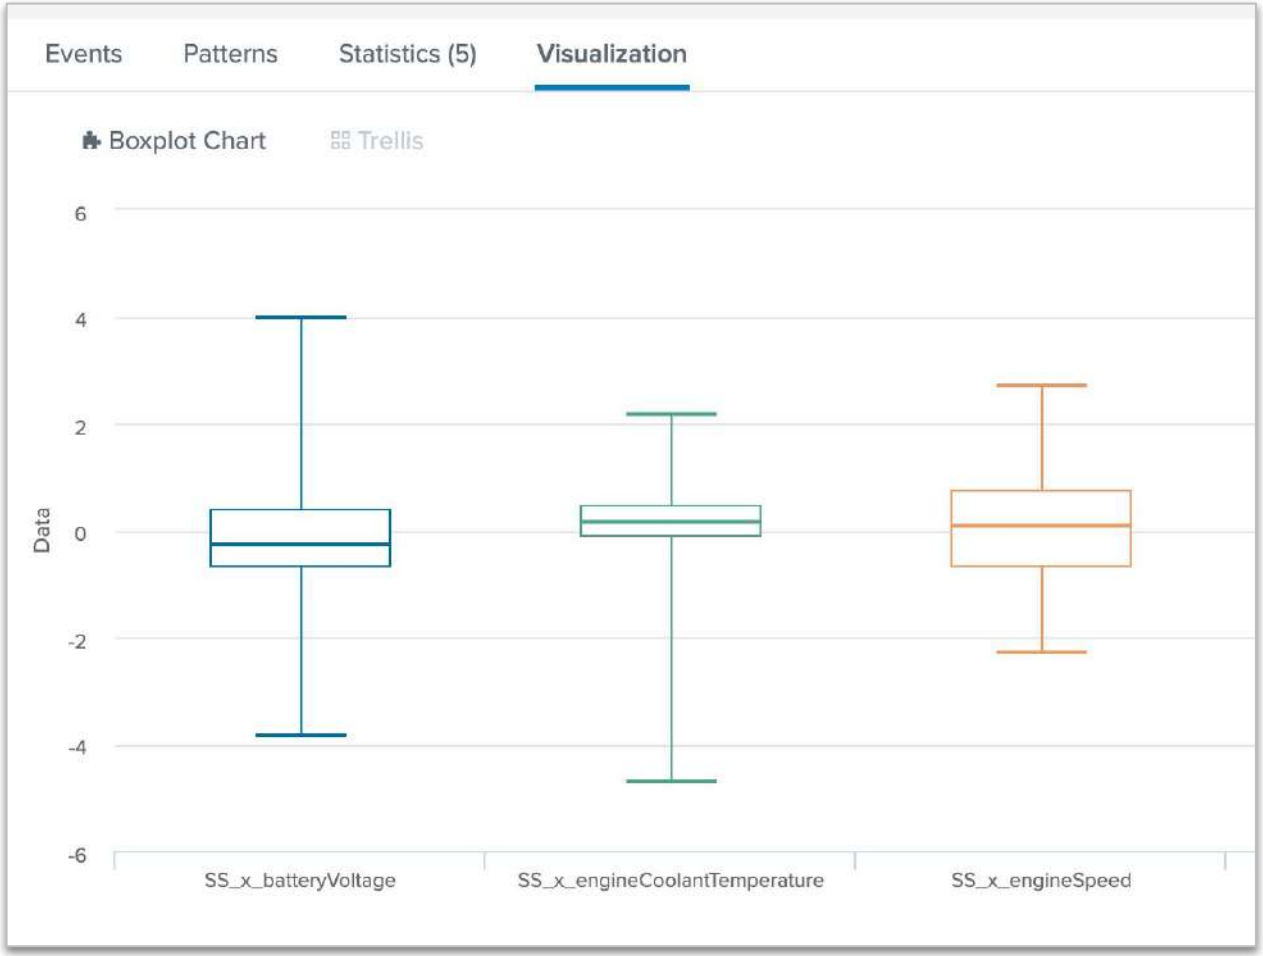

splunk>

2

Detect Numeric Outliers:

Explore the MLTK showcase and adapt it to start a new experiment with your own dataset

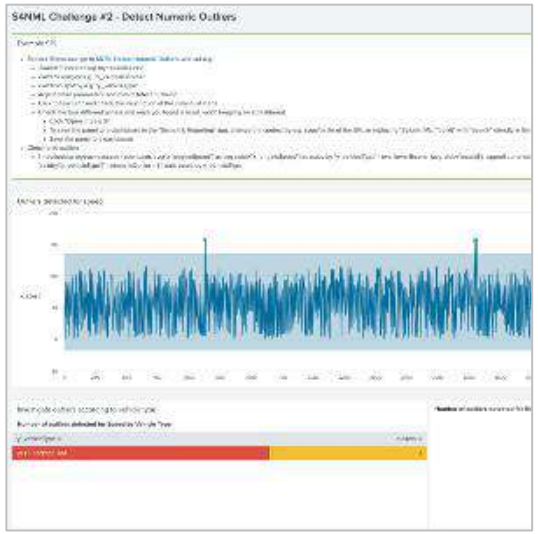

Explore the Outlier Detection Showcases

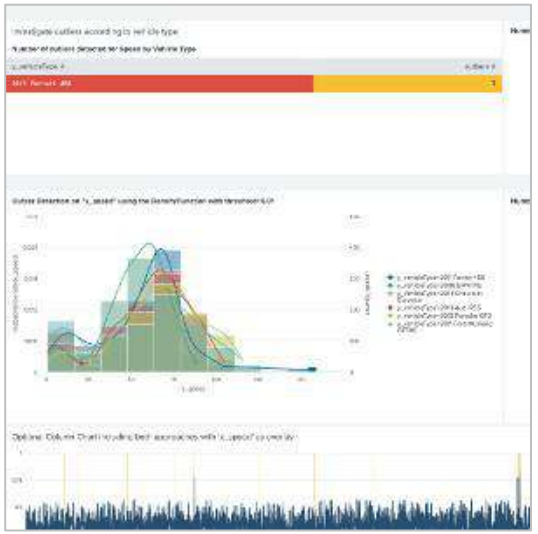

Start your own Outlier Detection Experiment

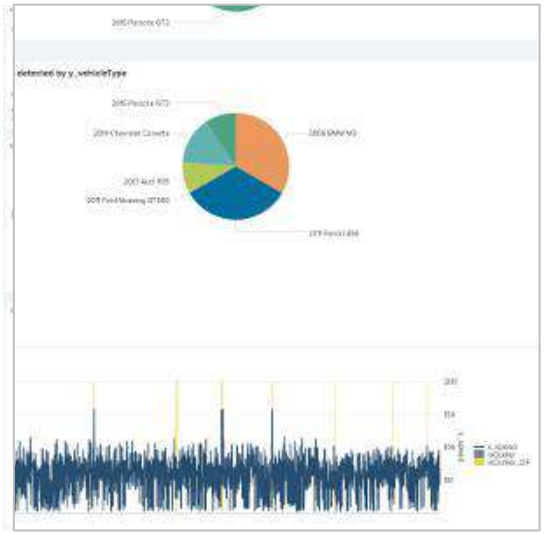

Optionally try to compare different outlier detection approaches

splunk>

2 Explore the Outlier Detection Showcases

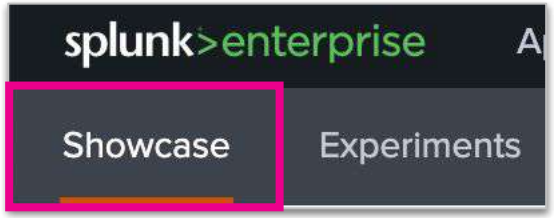

- Switch to the Showcase tab of the MLTK and explore the assistant to detect outliers in server response time
- We are now going to use statistics to detect the outliers

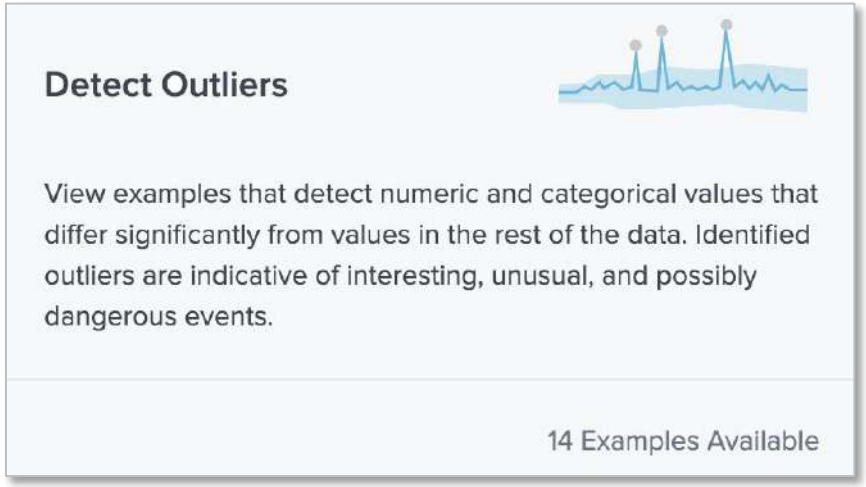

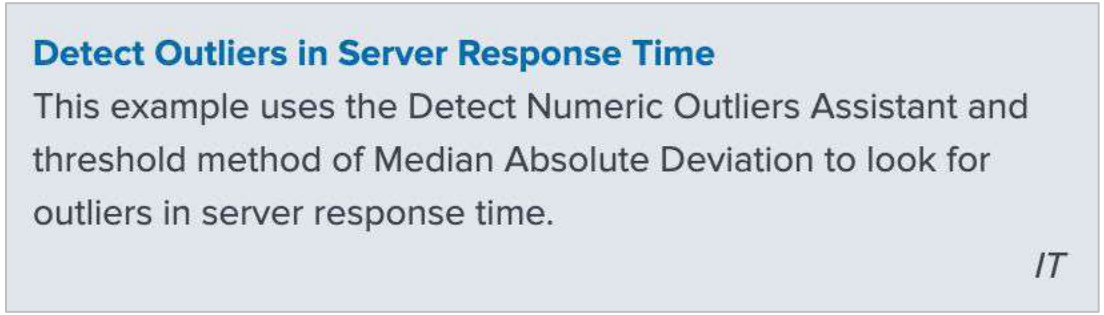

splunk>

2 Explore the Outlier Detection Showcases

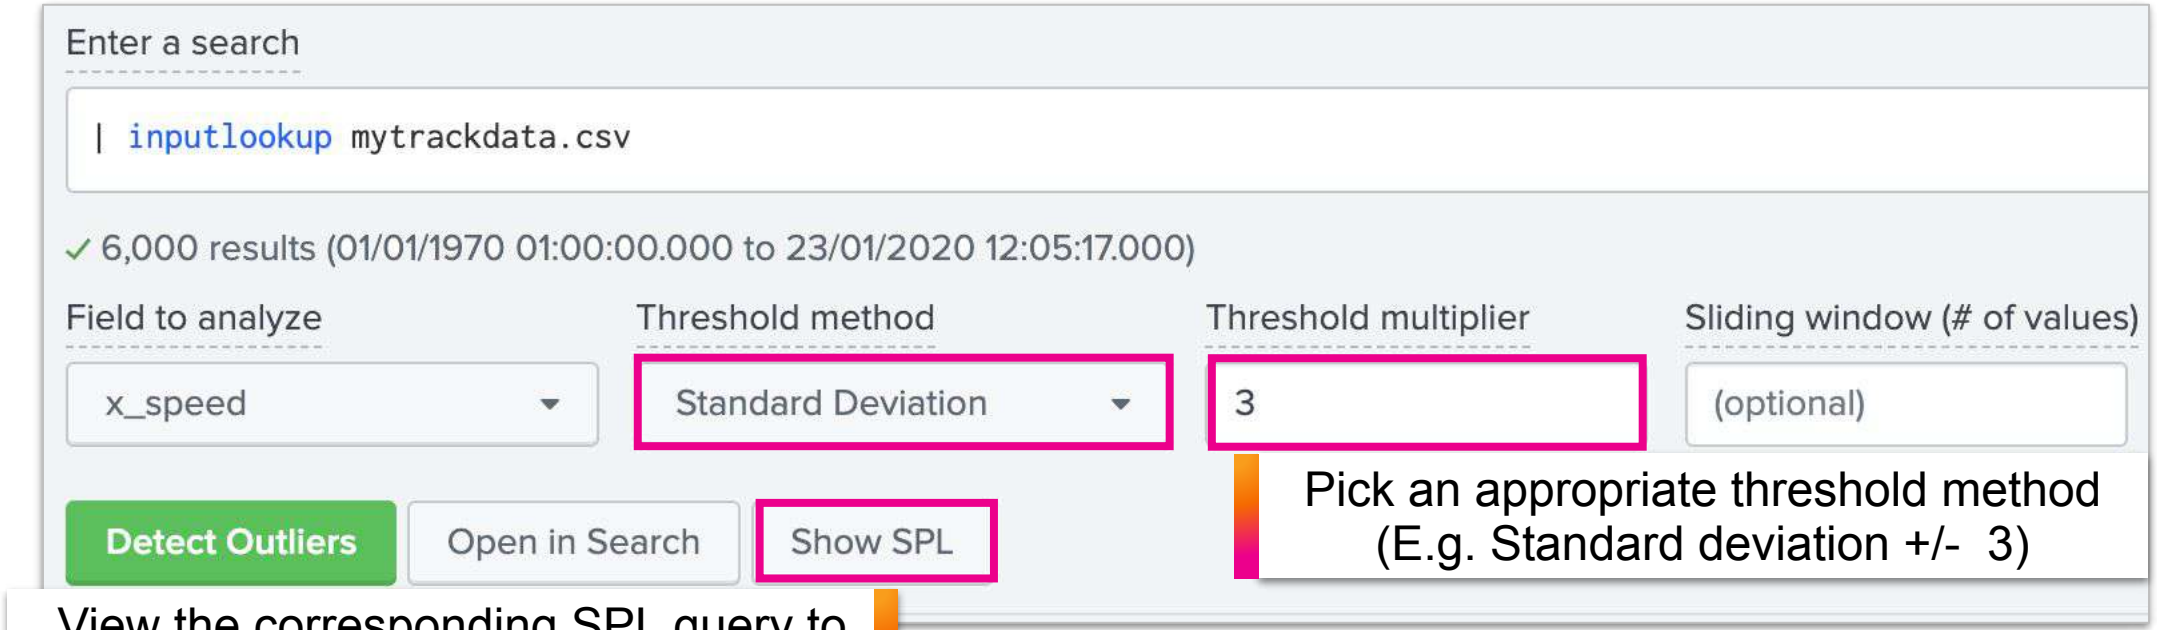

splunk>

2 Detecting Outliers with the Density Function

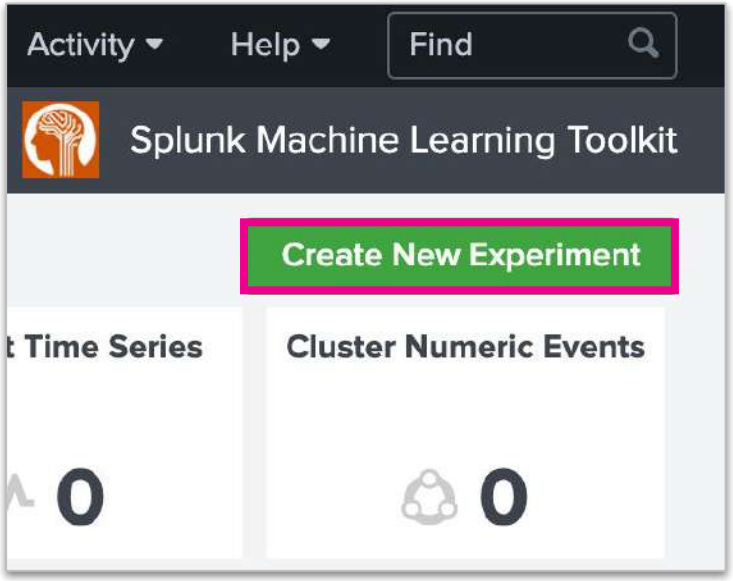

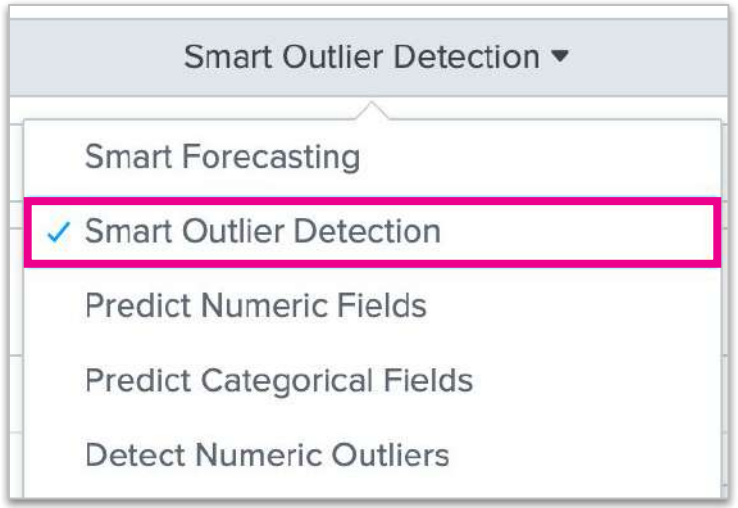

- Switch to the Experiments tab of the MLTK and create a new experiment
- Instead of an approach based on statistics
  we are now going to use the density function to detect outliers

2 Create Your Own Smart Outlier Experiment

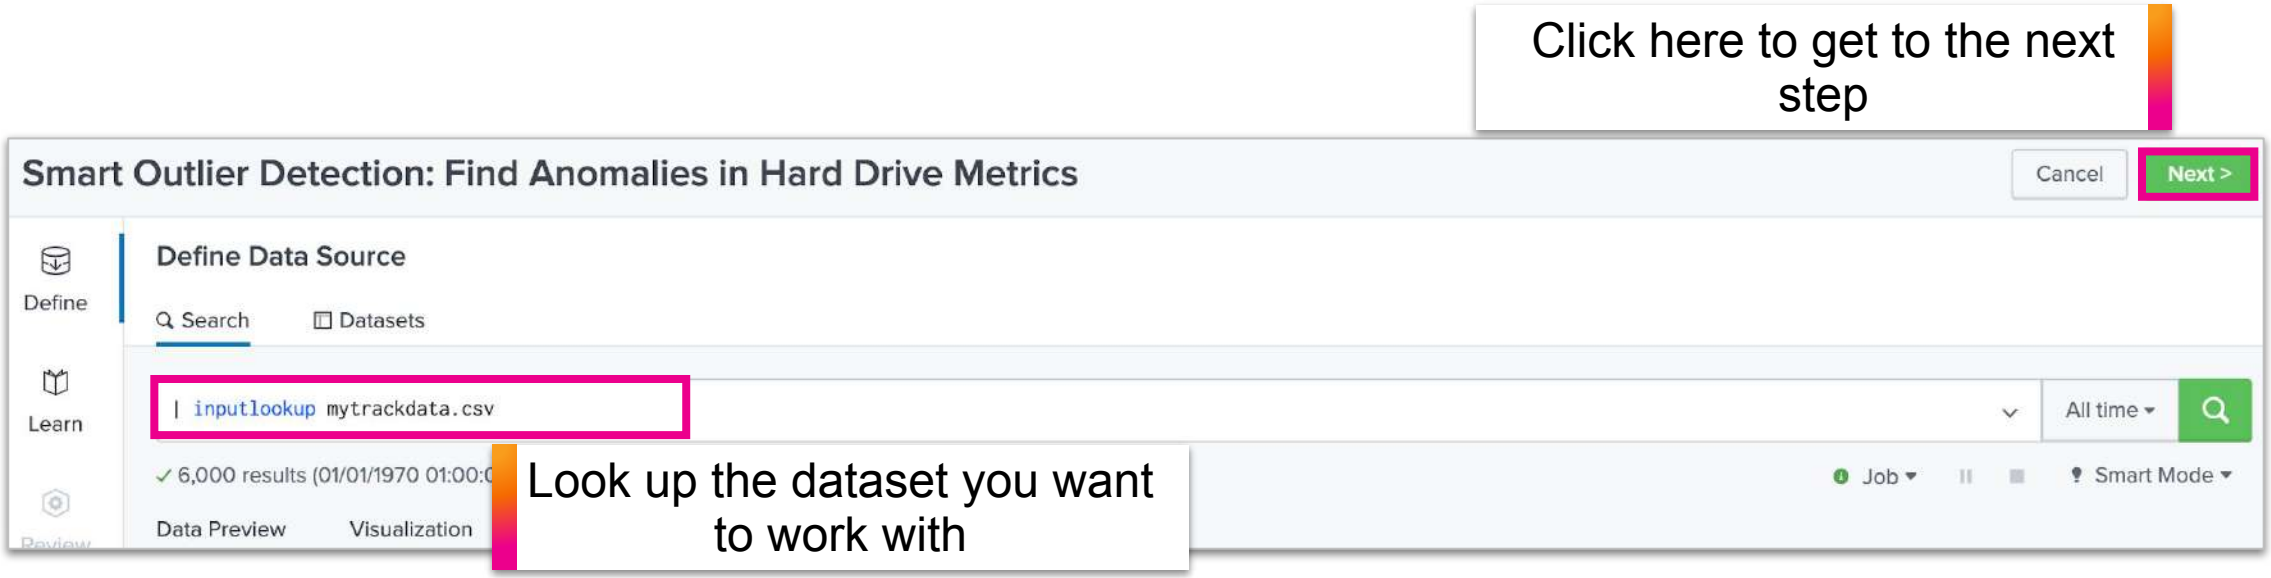

splunk>

3

SPL for MLTK: The `fit` and `apply` Commands

Examples:

```
<your search> | fit <model name>
<your search> | apply <model name>
```

```
| inputlookup mytrackdata.csv
| apply car_outlier_df_speed
```

- The `fit` command produces a machine learning model based on the behaviour of a set of events. It applies the model to the current search results in the search pipeline
- The `apply` command applies the machine learning model that was learned using the fit command

3

SPL for MLTK: The `fit` and `apply` Commands

Examples:

```
<your search> | fit StandardScaler <fields> into <model name>

<your search> | apply <model name> | `<macro name>`

<your search> | fit SVM “X X X" from “XXX" “XXX" kfold_cv=3
```

Check out the confusion matrix and classification statistics macros!

- The StandardScaler algorithm uses the scikit-learn StandardScaler algorithm to standardize data fields
- Splunk’s MLTK allows you to cross-validate your models right from the search queries that train them. Simply specify the number of cross-validation folds you want by setting the `fit` command’s parameter `kfold_cv`

3

Use a Classification Model:

Create a classification model and use it to predict vehicle types from your sensor data

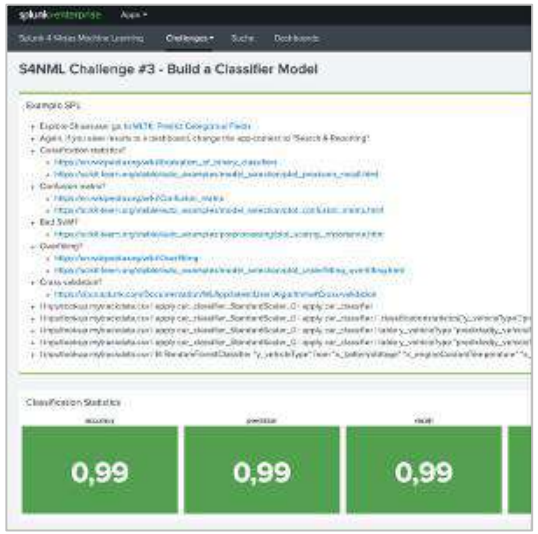

Explore the Classification Assistant

Put your Algorithm into Practice

Optionally find a way to deal with model overfitting

splunk>

3 Explore the Classification Assistant

Prediction Results

y_vehicleType	predicted(y_vehicleType)
2015 Porsche GT3	2011 Ford Mustang GT500
2011 Ferrari 458	2011 Ford Mustang GT500
2014 Chevrolet Corvette	2014 Chevrolet Corvette
2011 Ferrari 458	2011 Ford Mustang GT500
2011 Ford Mustang GT500	2011 Ford Mustang GT500
2015 Porsche GT3	2011 Ford Mustang GT500
2008 BMW M3	2011 Ford Mustang GT500
2011 Ford Mustang GT500	2011 Ford Mustang GT500
2011 Ferrari 458	2011 Ferrari 458
2011 Ferrari 458	2011 Ferrari 458

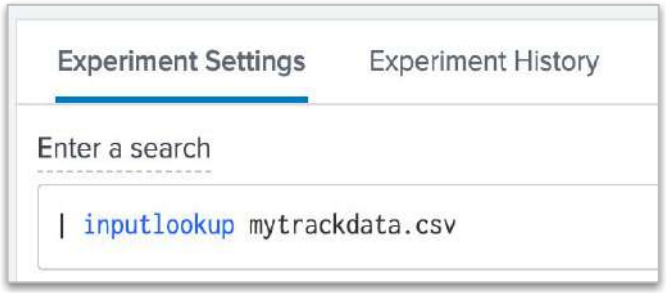

Why is SVM doing so bad?

splunk>

3

Save your Classification Model

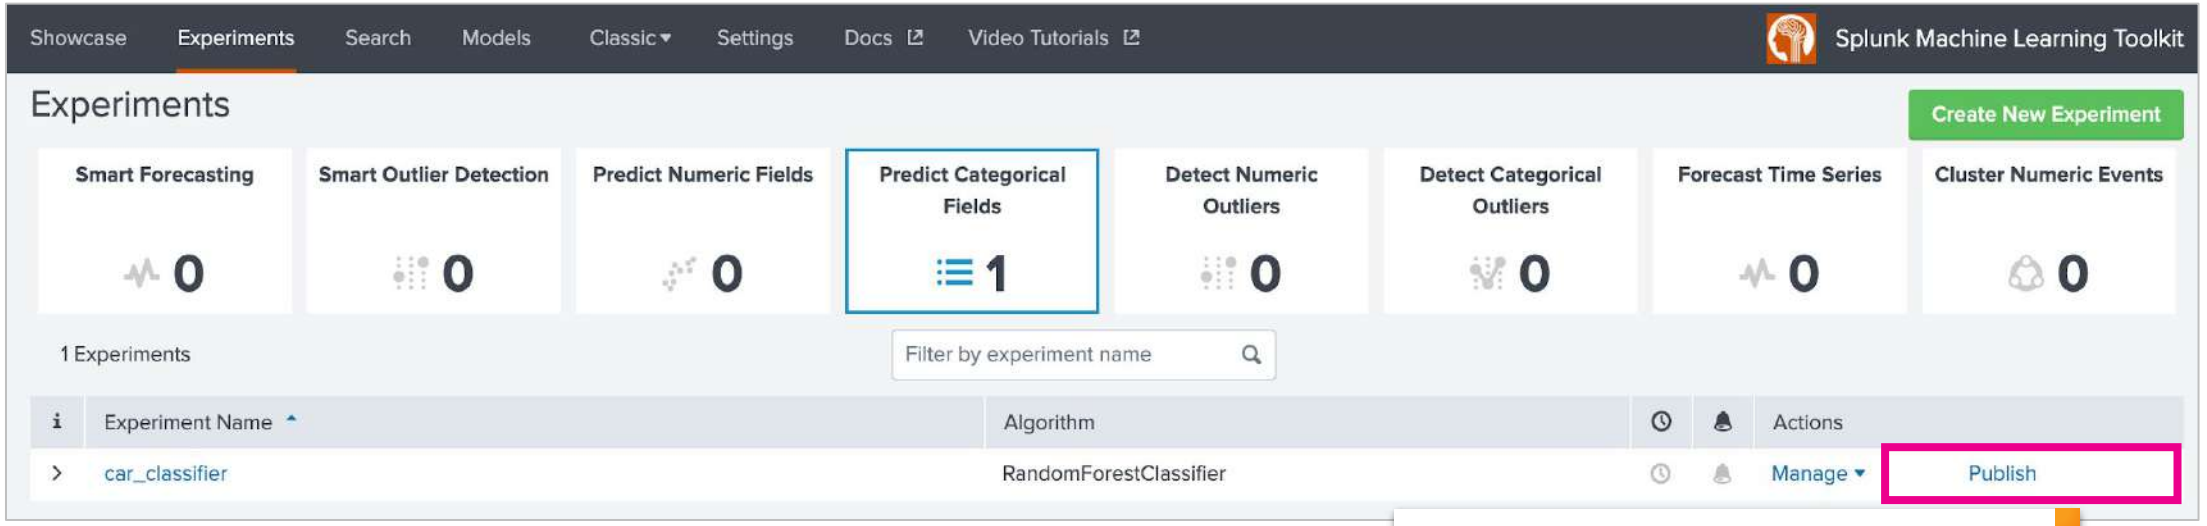

Publish your model in the app of your choice

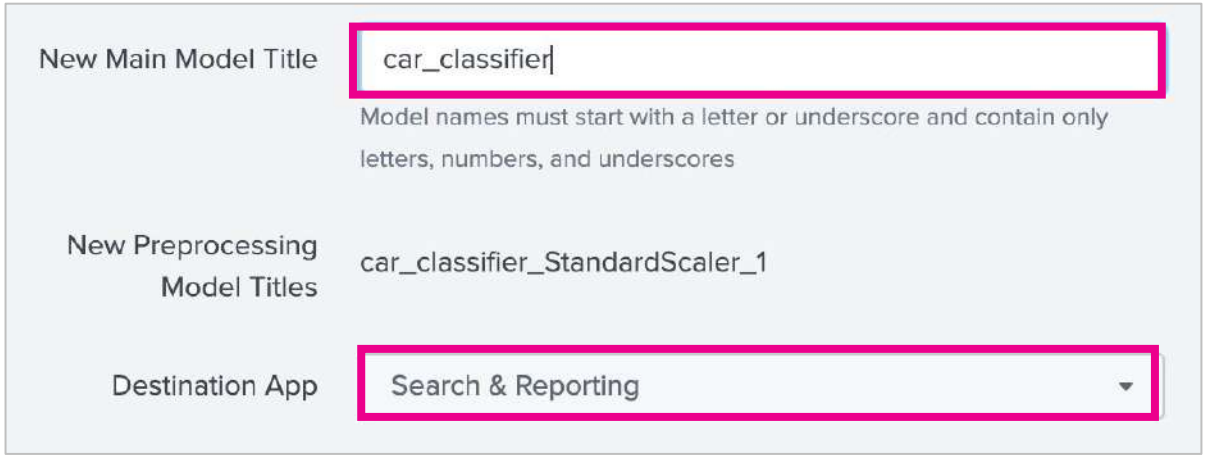

splunk>

3

Apply your Classification Model

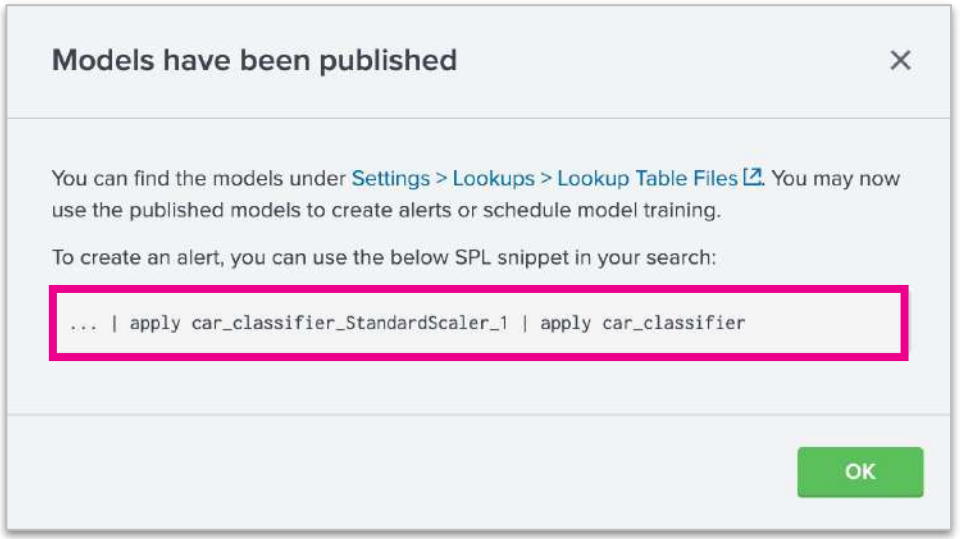


```
| inputlookup mytrackdata.csv
| apply car_classifier_StandardScaler_0
| apply car_classifier
| table y_vehicleType "predicted(y_vehicleType)" *
|  `confusionmatrix("y_vehicleType","predicted(y_vehicleType)")`
```

3

Which Car Gets Classified Worst?

How can you find out where your model is off?

Further analysis of wrong classifications

Predicted actual	Predicted 2011 Ferrari 458	Predicted 2011 Ford Mustang GT500	Predicted 2013 Audi RS5	Predicted 2015 Porsche GT3
2011 Ferrari 458	0	0	6	0
2011 Ford Mustang GT500	0	0	0	15
2013 Audi RS5	1	0	0	0
2015 Porsche GT3	0	8	0	0

splunk>

3

Use a Clustering Model:

Create a clustering model and use it to analyze your dataset

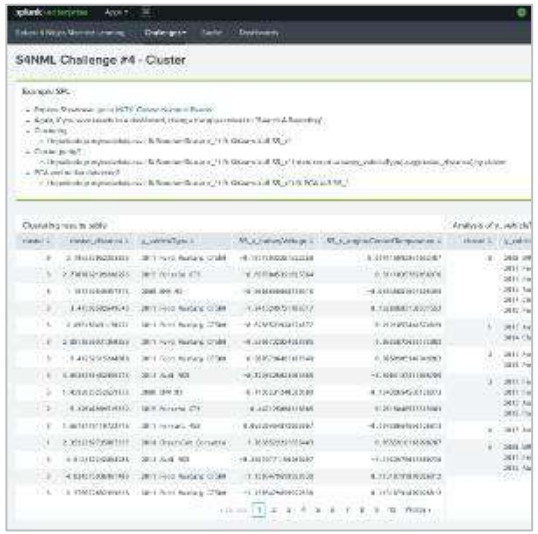

Explore the Clustering Assistant

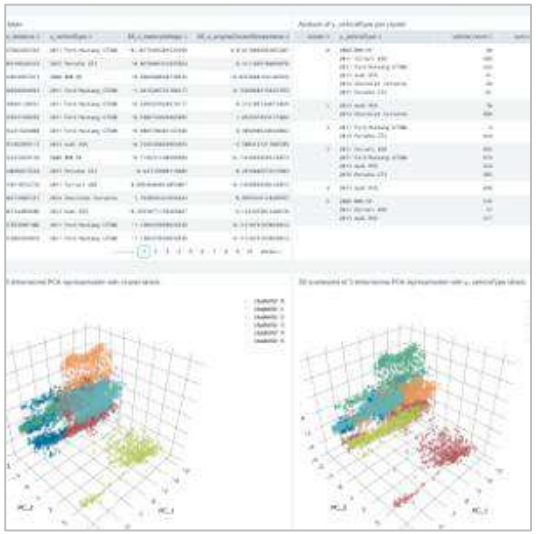

Cluster Analysis of the mytrackdata-Dataset

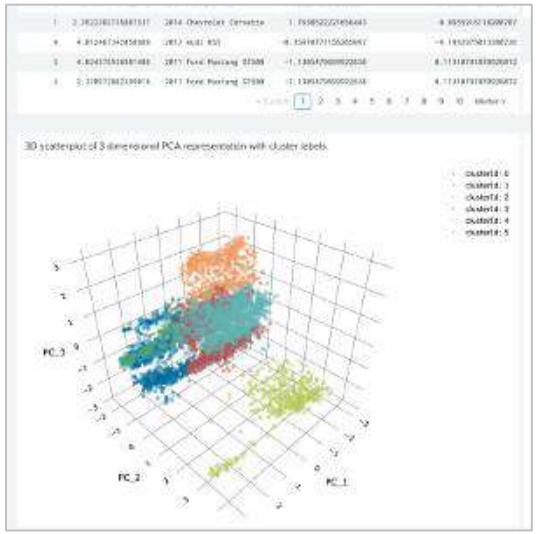

Optionally try and detect outliers

splunk>

4 Explore the Cluster Showcases

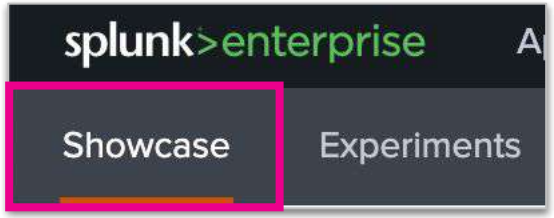

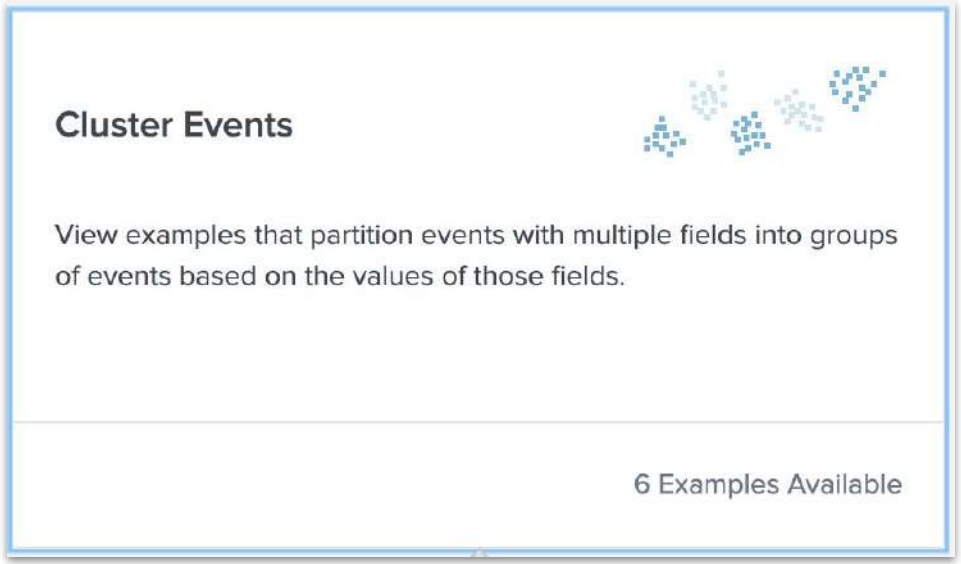

- Switch to the Showcase tab of the MLTK and explore the assistant to identify clusters of events

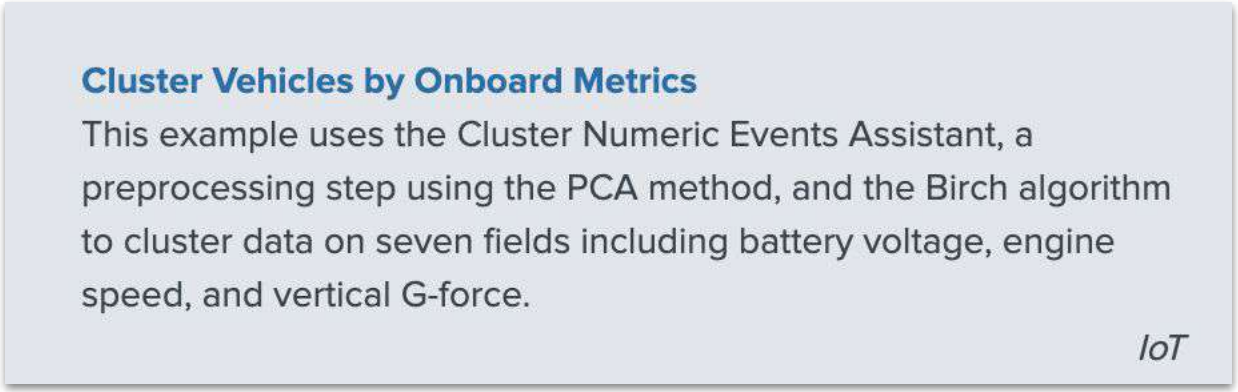

splunk>

4

The MLTK Comes with Many Different Algorithms

Example:

```
<your search> | fit PCA k=<int> <fields>
```

- Factor analysis with an algorithm such as PCA can reduce the number of variables one must deal with
- The k parameter specifies the number of features to be extracted from the data

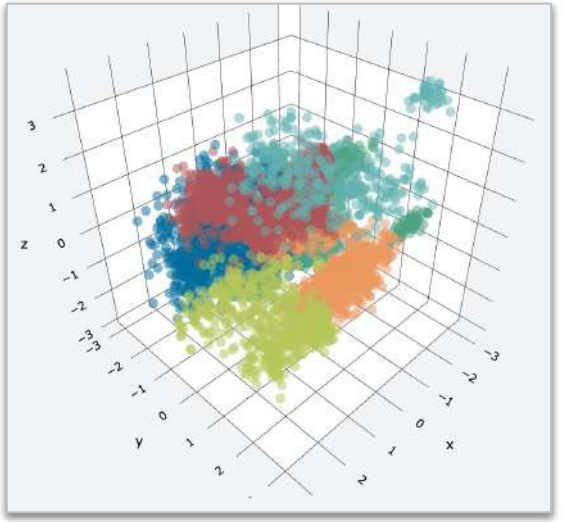

- clusterId: 2.0
- clusterId: 0.0
- clusterId: 5.0
- clusterId: 1.0
- clusterId: 3.0
- clusterId: 4.0
- clusterId: null

Why is there a cluster with "clusterId: null" ?

4

The MLTK Comes with Many Different Algorithms

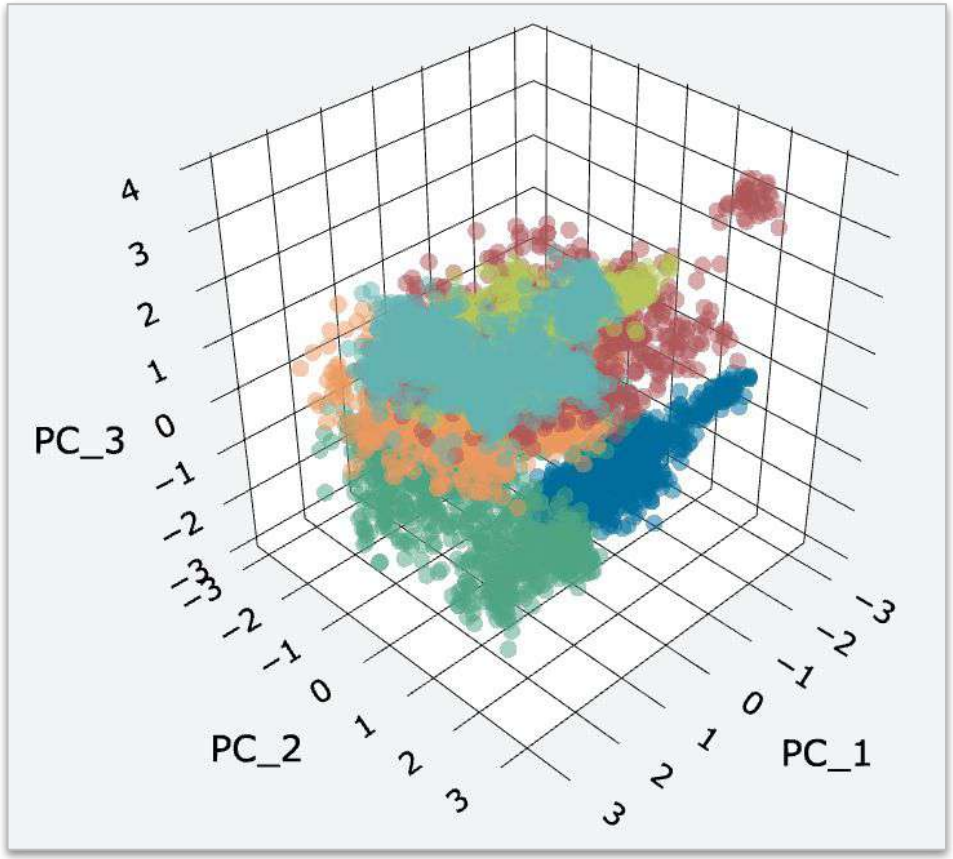

We have missing values in "x_engineCoolantTemperature" that we didn't fix/impute in mytrackdata.csv

- clusterId: 2008 BMW M3
- clusterId: 2015 Porsche GT3
- clusterId: 2011 Ford Mustang GT500
- clusterId: 2013 Audi RS5
- clusterId: 2011 Ferrari 458
- clusterId: 2014 Chevrolet Corvette

splunk>

Wrap Up

splunk>

Wrap Up

- **Feedback:** How was your experience, what worked well, what did not?

- **Discussion Brainstorming:** How could you transfer the topics learned today to other use cases or departments?

- **You want to learn more about Splunk's Machine Learning?**
  - Check our latest Splunk Blogs around Machine Learning
  - Watch videos from Splunk Machine Learning YouTube Channel
  - Take the Splunk Education Class for Data Science and Advanced Analytics
  - Learn more about Splunk's Machine Learning Advisory Program

Thank You!

Additional Information

Login:

Username: **admin**

Password: **<See Splunk Show>**

Challenge Solution Examples:

We created a dashboard for each challenge with example solutions in the hidden app “Splunk 4 Ninjas Machine Learning”. Use this app for preparation, debriefing after the challenges or as assistance for unexperienced attendees.

- https://{your-host}/en-GB/app/s4n_ml/splunk_4_ninjas_ml
  or click the button next to “Splunk 4 Ninjas Machine Learning” on top of the Home dashboard